

Optimization for Deep Learning

Lecture 1-2: Introduction

Kun Yuan

Peking University

Main contents in the class

- **Part I: Fundamental algorithms for optimization**

Gradient descent; projected gradient descent; proximal gradient descent; Nesterov acceleration; quasi-Newton algorithms; zeroth-order methods

- **Part II: Fundamental algorithms for deep learning**

Stochastic gradient descent (SGD); SGD stability; momentum SGD; adaptive SGD; variance reduction

- **Part III: Advanced algorithms for deep learning**


Mixed precision training; gradient clipping; adversarial learning; multi-task learning; meta learning; bilevel optimization

- **Part IV: Distributed algorithm for deep learning**

Communication compression; federated learning; decentralized learning; asynchronous SGD; Byzantine learning;

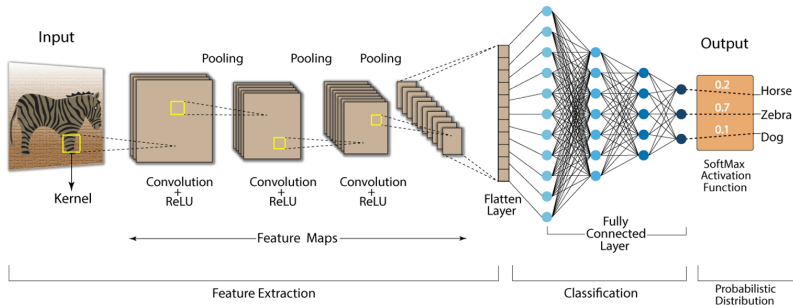
Deep netrual network (DNN)

- DNN is widely used in almost all AI applications
- A typical DNN model includes a **feature extractor** and a **classifier**
- Well-trained DNN can make precise predictions



A practical DNN example¹

Convolution Neural Network (CNN)




¹Source: analyticsvidhya.com

DNN model

- We model DNN as $h(x; a) : \mathbb{R}^d \rightarrow \mathbb{R}^c$
 - $x \in \mathbb{R}^d$ is the DNN model parameter to be trained
 - a is a random input data sample
 - c is the number of classes
- Given the model parameter x , DNN outputs prediction scores \hat{b} for input ξ_i




a



$$\hat{b} = h(x; a) \in \mathbb{R}^c$$

DNN model: a trivial example

- Given model parameter $x = [W; h]$, and a linear model $h(x; a) = Wa + h$,
- An illustration of the trivial DNN model and its output is as follows²



²Source: <https://cs231n.github.io/linear-classify/>

How to train a DNN model?

- Given good model x , DNN $h(x; a)$ can make precise predictions
- But how to train/achieve the model parameter x ?
- Given a dataset $\{a_i, b_i\}_{i=1}^m$ where b_i is the ground-truth label for data a_i
- Define $L(\hat{b}_i, b_i) = L(h(x; a_i), b_i)$ as a loss function to measure the difference/mismatch between predictions and ground-truth labels
- DNN training is to find a model parameter x such that the mismatch (between pred and real) are minimized across the entire dataset:

$$x^* = \arg \min_{x \in \mathbb{R}^d} \left\{ \frac{1}{m} \sum_{i=1}^m L(h(x; a_i), b_i) \right\}$$


DNN model is notoriously difficult to train

- DNN model $L(h(x; a), b)$ is highly **non-convex**, and probably non-smooth

$$h(x; a) = \psi(\cdots \psi(W_2 \cdot \psi(W_1 a + h_1) + h_2) \cdots)$$


$$L(\hat{b}; b) = \frac{1}{2} \|b - \hat{b}\|^2 \text{ or } - \sum_i b_i \log(\hat{b}_i) \text{ or others}$$

where $x = \{W_i, h_i\}$ and $\psi(\cdot)$ is a non-linear activation function



DNN model is notoriously difficult to train

- Cannot find global minima; trapped into local minima and saddle points
- The dimension of model parameter $x = \{W_i, h_i\}$ (or model size) is huge³



³Image source: neowin.net

DNN model is notoriously difficult to train

- Cannot find global minima; trapped into local minima and saddle points
- The dimension of model parameter $x = \{W_i, h_i\}$ (or model size) is huge
- The size of the dataset $\{a_i, b_i\}_{i=1}^m$ is huge

DNN Trainig = Non-convexity training + Huge dimension + Huge dataset

- Our lectures will focus on algorithms to train DNN

Stochastic optimization: problem

- Consider the stochastic optimization problem:

$$\min_{x \in \mathbb{R}^d} f(x) = \mathbb{E}_{\xi \sim D}[F(x; \xi)]$$

- ξ is a random variable indicating data samples
 - D is the data distribution; unknown in advance
 - $F(x; \xi)$ is differentiable in terms of x
- Many applications in signal processing and machine learning

Example: deep neural network

- Recall the DNN training problem

$$\min_{x \in \mathbb{R}^d} \quad \frac{1}{m} \sum_{i=1}^m L(h(x; a_i), b_i)$$

which is a finite-sum problem

- Suppose we have infinite data (a, b) following distribution D , the above problem becomes

$$\min_{x \in \mathbb{R}^d} \quad f(x) = \mathbb{E}_{(a,b) \sim D} L(h(x; a), b)$$

Stochastic gradient descent: algorithm

- Recall the problem

$$\min_{x \in \mathbb{R}^d} f(x) = \mathbb{E}_{\xi \sim D}[F(x; \xi)]$$

- Closed-form of $f(x)$ is unknown; gradient descent is not applicable
- Stochastic gradient descent (SGD):

$$x_{k+1} = x_k - \gamma \nabla F(x_k; \xi_k)$$

where ξ_k is a data realization sampled at iteration k .

- Our lecture will explore the following questions
 - Can SGD converge to the desired solution and how fast?
 - How does stochastic noise influence the convergence rate?
 - How to tune learning rate?

SGD Stability

- SGD was first proposed to relieve the computational overhead suffered in GD
- Surprisingly, it is often observed to generalize better than GD
- SGD first escapes from the GD solution and then converges to a better solution that can generalize better⁴





Figure 6: Fast escape phenomenon in fitting corrupted FashionMNIST.

⁴Figure is from (Wu et al., 2018)

SGD Stability

- SGD favors flat solution⁵



- Our lecture will explore the following questions:
 - Why does SGD favor flat solution?
 - Why does flat solution generalize better?
 - Why does SGD outperform GD in generalization?

⁵Figure is from Lei Wu's talk

Momentum SGD

- Momentum SGD can be rewritten into ($m_0 = 0$)

$$m_k = \beta m_{k-1} + \nabla F(x_k; \xi_k)$$

$$x_{k+1} = x_k - \gamma m_k$$

- Equivalent to

$$x_{k+1} = x_k - \gamma \nabla F(x_k; \xi_k) + \underbrace{\beta (x_k - x_{k-1})}_{\text{momentum}}$$

- Our lecture will explore the following questions:
 - Can momentum SGD converge to the desired solution and how fast?
 - Is momentum SGD theoretically faster than vanilla SGD?

Adaptive gradient method (AdaGrad)

- SGD performance is very sensitive to the learning rate. A good learning rate schedule can significantly speed up the convergence
- Adaptive stochastic gradient method

$$g_k = \nabla F(x_k; \xi_k)$$

$$s_k = s_{k-1} + g_k \odot g_k$$

$$x_{k+1} = x_k - \frac{\gamma}{\sqrt{s_k} + \epsilon} \odot g_k$$

where $1/\sqrt{s_k} = \text{col}\{1/\sqrt{s_{k,1}}, \dots, 1/\sqrt{s_{k,d}}\} \in \mathbb{R}^d$ is an element-wise operation, s_0 is initialized as 0, and a small ϵ is added for safe-guard.

- Learning rate is automatically tuned by s_k

Adaptive gradient method (AdaGrad)

- AdaGrad falls into preconditioned SGD
- If we let $P_k = \text{diag}\{\frac{1}{\sqrt{s_{k,1}}+\epsilon}, \dots, \frac{1}{\sqrt{s_{k,d}}+\epsilon}\} \in \mathbb{R}^{d \times d}$, AdaGrad becomes

$$x_{k+1} = x_k - \gamma P_k g_k$$

where P_k is a time-varying preconditioning matrix.

- AdaGrad imposes smaller learning rates for notable gradient directions
- AdaGrad imposes larger learning rates for insignificant gradient directions

Adaptive gradient method (AdaGrad)

AdaGrad alleviates the “Zig-Zag” phenomenon




Figure: SGD converges slow for ill-conditioned problem

Adaptive gradient method (AdaGrad)

AdaGrad alleviates the “Zig-Zag” phenomenon




Figure: AdaGrad has alleviated “Zig-Zag” phenomenon

Adaptive gradient method (AdaGrad)

The learning rate in AdaGrad is adaptive; **no need to tune**.




Figure: AdaGrad automatically adapts to problem structure⁶.

⁶These examples are from <https://conferences.mpi-inf.mpg.de/adfocs/material/alina/adaptive-L1.pdf>

More adaptive algorithms

- RMSProp, ADAM, AdaBound, etc.
- Animation of different adaptive SGD: <https://imgur.com/a/Hqolp>
- Our lecture will explore the following questions
 - How to prove convergence of these adaptive algorithms?
 - Can the adaptability bring theoretical benefits?

Finite-sum minimization

- In deep learning, we typically have a finite dataset
- If ξ takes samples uniformly from a finite dataset $\{\xi_1, \dots, \xi_m\}$, then

$$\mathbb{E}[F(x; \xi)] = \frac{1}{m} \sum_{i=1}^m F(x; \xi_i) = \frac{1}{m} \sum_{i=1}^m F_i(x)$$

where $F_i(x) := F(x; \xi_i)$. Number m is typically very large.

- The stochastic optimization problem becomes

$$\min_{x \in \mathbb{R}^d} f(x), \quad \text{where} \quad f(x) = \frac{1}{m} \sum_{i=1}^m F_i(x),$$

which is called **finite-sum minimization** problem.

Solving finite-sum minimization with SGD

- Since finite-sum minimization is a special example of stochastic optimization, we can solve it with SGD
- For any $k = 0, 1, \dots$, the SGD recursion is as follows

Pick $i_k \in [m]$ uniformly randomly;

$$x_{k+1} = x_k - \gamma \nabla F_{i_k}(x_k)$$

- Unbiased stochastic gradient: $\mathbb{E}[\nabla F_{i_k}(x)] = \frac{1}{m} \sum_{i=1}^m \nabla F_i(x) = \nabla f(x)$
- Convergence property follows SGD

Stochastic variance-reduced gradient

- In SGD, stochastic gradient $\nabla F_{i_k}(x)$ has constant variance:

$$\mathbb{E}[\nabla F_{i_k}(x)] = \nabla f(x), \quad \mathbb{E}\|\nabla F_{i_k}(x) - \nabla f(x)\|^2 \leq \sigma.$$

- The constant variance is the root reason why SGD is slow
- For infinite-data scenario, we can do little to reduce gradient variance
- But for finite-sum scenario, we can construct variance-reduced stochastic gradient with diminishing variance, and hence significantly improve the convergence rate, i.e., $\mathbb{E}\|\nabla F_{i_k}(x) - \nabla f(x)\|^2 \rightarrow 0$.
- We leave it to the main lecture

Numerical performance

Stochastic variance-reduced methods are as cheap to update as **SGD**, and also have a fast exponential convergence like full gradient descent.




Figure: SGD, SVRG, and SAGA; figure is from (Gower et al., 2020)

Summary

- DNN training can be formulated into a non-convex stochastic optimization problem
- We introduced various fundamental algorithms to train deep neural network: SGD, momentum SGD, adaptive SGD, and variance-reduction

References I

- L. Wu, C. Ma *et al.*, “How sgd selects the global minima in over-parameterized learning: A dynamical stability perspective,” *Advances in Neural Information Processing Systems*, vol. 31, 2018.
- R. M. Gower, M. Schmidt, F. Bach, and P. Richtárik, “Variance-reduced methods for machine learning,” *Proceedings of the IEEE*, vol. 108, no. 11, pp. 1968–1983, 2020.