

# Optimization for Deep Learning

## Lecture 1-3: Introduction

**Kun Yuan**

Peking University

# Main contents in the class

- **Part I: Fundamental algorithms for optimization**

Gradient descent; projected gradient descent; proximal gradient descent; Nesterov acceleration; quasi-Newton algorithms; zeroth-order methods

- **Part II: Fundamental algorithms for deep learning**

Stochastic gradient descent (SGD); SGD stability; momentum SGD; adaptive SGD; variance reduction

- **Part III: Advanced algorithms for deep learning**

Mixed precision training; gradient clipping; adversarial learning; multi-task learning; meta learning; bilevel optimization

- **Part IV: Distributed algorithm for deep learning**

Communication compression; federated learning; decentralized learning; asynchronous SGD; Byzantine learning;

## Mixed precision training

- Efficient DNN training relies on using lower precision data types
- Float16 matrix multiplication is  $16\times$  faster than float32 in A100
- Using lower precision can save memory
- Training large DNN is infeasible without using mixed precision

## Mixed precision training

- Mixed precision training:

$$x_{k+1} = x_k - \gamma Q(\nabla F(x_k, \xi_k))$$

where  $Q(\cdot)$  is a quantization operator using lower precision

- Our lecture will explore the following questions:
  - Can the algorithm converge to the desired solution?
  - How does  $Q(\cdot)$  influence the convergence rate?
  - How to design  $Q(\cdot)$ ?
  - How to use mixed precision for momentum SGD or adaptive SGD?


# Gradient clipping: recurrent neural network (RNN)

- RNN has the following recursion:

$$h_t = f(x_t, h_{t-1}; \theta_h)$$

$$\hat{y}_t = g(h_t; \theta_o)$$

where  $\theta_h$  and  $\theta_o$  are the parameters of  $f(\cdot)$  and  $g(\cdot)$ , respectively, and  $h_0$  can be initialized to arbitrary values.



## Gradient clipping: vanishing and exploding gradients

- The gradient in linear RNN is:

$$\frac{\partial F}{\partial W_x} = \frac{1}{T} \sum_{t=1}^T \sum_{i=1}^t (W_h^\top)^{t-i} W_o^\top \frac{\partial L(\hat{y}_t, y_t)}{\partial \hat{y}_t} x_i^\top \in \mathbb{R}^{n \times d}.$$

- $(W_h^\top)^t$  will cause a significant numerical issue in  $\partial F / \partial W_x$
- If the largest magnitude of the eigenvalue is less than 1, i.e.,  $|\lambda(W_h^\top)| < 1$ , it holds that  $(W_h^\top)^{t-i} \rightarrow 0$  as  $t$  (or  $T$ ) gets large; **Gradient vanishing!**
- If the largest magnitude of the eigenvalue is greater than 1, i.e.,  $|\lambda(W_h^\top)| > 1$ , it holds that  $(W_h^\top)^t \rightarrow +\infty$  as  $t$  (or  $T$ ) gets large; **Gradient exploding!**
- Activation functions may also amplify gradient vanishing and exploding

## Gradient clipping: algorithm

- Consider the following non-convex optimization problem

$$\min_{x \in \mathbb{R}^d} f(x)$$

- The gradient clipping algorithm iterate as follows

$$x_t = x_t - \gamma g_t \quad \text{where} \quad g_t = \text{clip}(\nabla f(x_t), c)$$

for some positive constant  $c > 0$ .

- The clipping operator is defined as

$$\begin{aligned} \text{clip}(u, c) &= \min\left\{1, \frac{c}{\|u\|}\right\} u \quad \forall u \in \mathbb{R}^d \\ &= \begin{cases} u & \text{if } \|u\| \leq c \\ \frac{c}{\|u\|} u & \text{if } \|u\| > c \end{cases} \end{aligned}$$

where  $\|\cdot\|$  is an  $\ell_2$ -norm.

## Gradient clipping: resolving gradient exploding

- Clipping operator does not change the gradient direction; just scales gradient.
- Clipping operator squeezes large gradient when  $\|\nabla f(x)\| > c$ , but does nothing to small gradient.
- After clipping, it is guaranteed that  $\|u\| \leq c$  for any  $u \in \mathbb{R}^d$ .
- Our lectures will explore the following questions:
  - Can gradient clipping overcome gradient exploding?
  - Under what conditions can gradient clipping perform better than GD?



## Adversarial learning: DNN is fragile to adversarial attacks




Figure: A demonstration of the adversarial example [Goodfellow et.al., 2015].

## Adversarial learning: Adversarial stop sign

### Adversarial Examples



Clean Stop Sign



Real-world Stop Sign  
in Berkeley



Adversarial Example



Adversarial Example



"Stop sign"

"Stop sign"

"Speed limit sign 45km/h"

"Speed limit sign 45km/h"

## Adversarial learning: Adversarial T-shirt



# Adversarial learning: Adversarial attacks in NLP

Targeted caption



## Original top caption

A man holding a tennis racquet on a tennis court

## Adversarial top caption

A woman brushing her teeth in a bathroom

Keywords



## Original top caption

A cake that is sitting on a table

## Adversarial top caption

A **dog** and a **cat** are playing with a **Frisbee**

## Adversarial keywords:

"dog", "cat" and "Frisbee"

## Adversarial learning: Adversarial attacks in NLP

<b>Original Input</b>	Connoisseurs of Chinese film will be pleased to discover that Tian's meticulous talent has not withered during his enforced hiatus.	Prediction: <u>Positive (77%)</u>
<b>Adversarial example [Visually similar]</b>	<u>Aonnoisseurs</u> of Chinese film will be pleased to discover that Tian's meticulous talent has not withered during his enforced hiatus.	Prediction: <u>Negative (52%)</u>
<b>Adversarial example [Semantically similar]</b>	Connoisseurs of Chinese <u>footage</u> will be pleased to discover that Tian's meticulous talent has not withered during his enforced hiatus.	Prediction: <u>Negative (54%)</u>

## Adversarial learning: Construct adversarial examples

- An adversarial example is a perturbation  $\eta$  to maximize misclassification
- Given an input pair  $(\xi, y)$ , its adversarial example  $\eta \in \mathbb{R}^d$  is defined as

$$\eta \in \arg \max_{\eta: \|\eta\| \leq \epsilon} L(h(x^*, \xi + \eta), y)$$

where  $x^*$  is the optimal DNN model.

- How to solve the above problem? We leave it to the main lecture

## Adversarial learning: problem formulation

- Adversarial machine learning is to make models robust to attacks

$$\min_{x \in \mathbb{R}^d} \frac{1}{m} \sum_{i=1}^m f_i(x) \quad \text{where} \quad f_i(x) = \max_{\eta: \|\eta\|_{\infty} \leq \epsilon} L(h(x; \xi_i + \eta), y_i)$$

- We maximize  $\eta$  to construct adversarial examples but minimize  $x$  to construct robust machine learning models; **minimax** optimization!
- How to solve the above problem? We leave it to the main lecture

# Meta learning

- Deep learning algorithms have too many hyper-parameters
  - learning rate, momentum coefficient, network architecture, etc.
- Manually tuning the hyper-parameters is neither efficient nor effective
- Can we tune the hyper-parameter automatically? Yes! **Meta learning!**



# Meta learning

Task 1



⋮

Task N



train task 1 by  $\theta_{k+1}^{(1)} = \theta_k^{(1)} - \gamma g(\theta_k^{(1)})$   
to get a good classifier  $\theta_\star^{(1)}$

⋮

train task N by  $\theta_{k+1}^{(N)} = \theta_k^{(N)} - \gamma g(\theta_k^{(N)})$   
to get a good classifier  $\theta_\star^{(N)}$

Learn a good learning  
rate  $\gamma_\star$  that performs  
well in tall tasks

New task



train new task by  $\theta_{k+1} = \theta_k - \gamma_\star g(\theta_k)$

use  $\gamma_\star$

# Meta learning

- Suppose we have a collection of  $M$  tasks  $\{\mathcal{T}_i\}_{i=1}^M$ . Each task is associated with a dataset pair  $(\mathcal{D}_i^{\text{tr}}, \mathcal{D}_i^{\text{test}})$ .
- Let  $\phi$  be the hyper-parameter to learn, which is common to all tasks.
- Let  $\theta_i$  be the model for task  $i$ . Given an specific algorithm  $\mathcal{Alg}$  (such as SGD), the hyper-parameter  $\phi$  (such as the learning rate), and the training dataset  $\mathcal{D}_i^{\text{tr}}$ ,  $\theta_i$  can be learned by

$$\theta_i = \mathcal{Alg}(\phi, \mathcal{D}_i^{\text{tr}}) = \arg \min_{\theta \in \mathbb{R}^d} \left\{ \mathbb{E}_{\xi_i \sim \mathcal{D}_i^{\text{tr}}} [F(\theta; \phi, \xi_i)] \right\}$$

- The hyper-parameter  $\phi$  can be learned by the meta-learning problem

$$\phi^* = \arg \min_{\phi \in \mathbb{R}^s} \left\{ \frac{1}{M} \sum_{i=1}^M L(\theta_i, \mathcal{D}_i^{\text{test}}) \right\} \text{ where } \theta_i = \mathcal{Alg}(\phi, \mathcal{D}_i^{\text{tr}})$$

# Meta learning


- Meta learning is essentially a bi-level optimization problem:

$$\min_{x \in \mathbb{R}^p} \Phi(x) := f(x, y^*(x)), \quad \text{where} \quad y^*(x) = \arg \min_{y \in \mathbb{R}^q} \{g(x, y)\}$$

- How to solve bilevel optimization? We leave it to the main lecture

# Multi-task learning

- Multi-task learning solves many tasks simultaneously with one neural network
  - Dataset of each task can be shared with each other
  - More efficient than train each individual network independently
- One body with multiple heads



# Multi-task learning

- Multi-task learning can be formulated as

$$\min_{u, \{v_i\}} \frac{1}{M} \sum_{i=1}^M f(u, v_i) \quad \text{where} \quad f(u, v_i) = \mathbb{E}[F(u, v_i, \xi_i)]$$

where  $u$  is the shared model while  $v_i$  is the specific model for task  $i$ .

- How to solve the above problem? We leave it to the main lecture

## More advanced deep learning problems

- Large-batch learning
- Self-supervised learning
- Contrastive learning
- Multi-model learning

We will discuss them in lectures if time allows

## Summary

- Many deep learning tasks can be formulated into more advanced optimization problems
- We previewed mixed precision learning, gradient clipping, adversarial learning, meta learning, and multi-task learning in this lecture