

Parameters, Memories, and Computations in Transformers

Kun Yuan

Center for Machine Learning Research @ Peking University



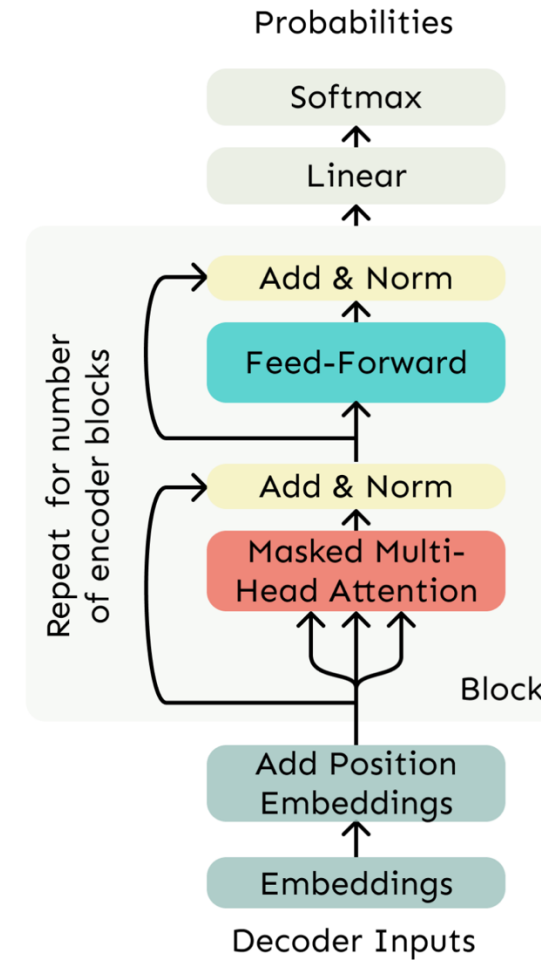
PART 01

Settings and Basics



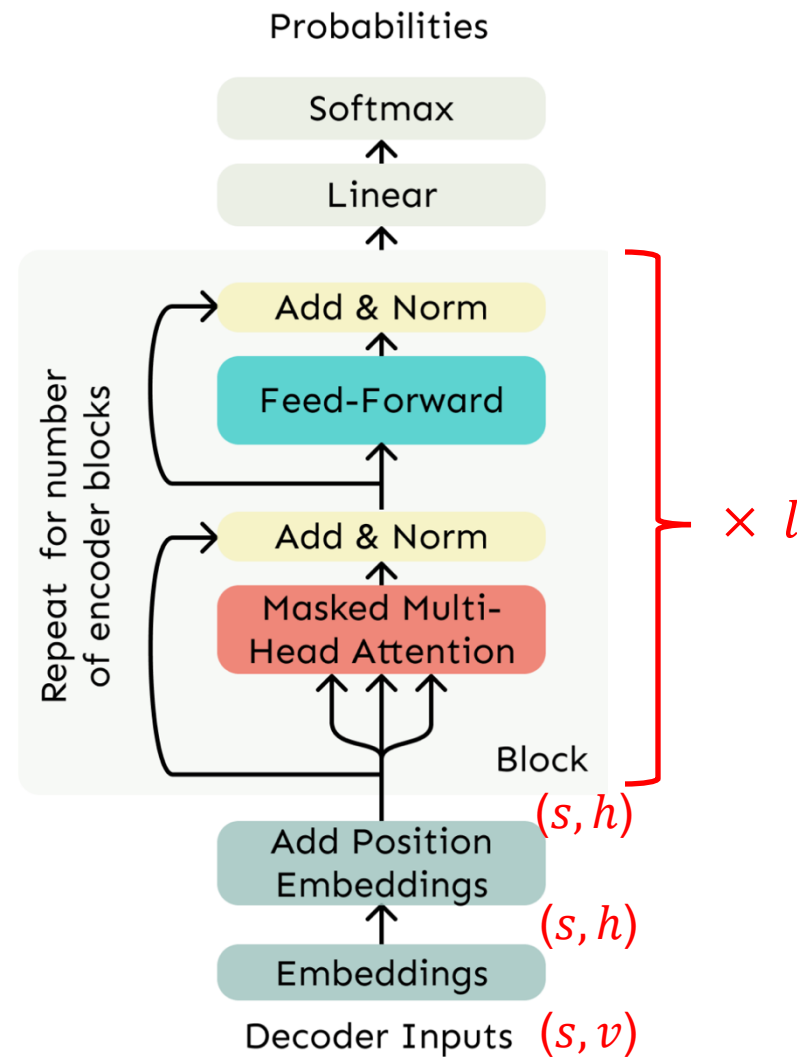
Decoder-only Transformer

- GPT is based on the **decoder-only** transformer
- We will analyze the parameters, memories, and computation costs for decoder-only transformer




Notations

- Number of the transformer layers: l
- Sequence length: s
- Vocabulary size: v
- Embedding representation dims: h



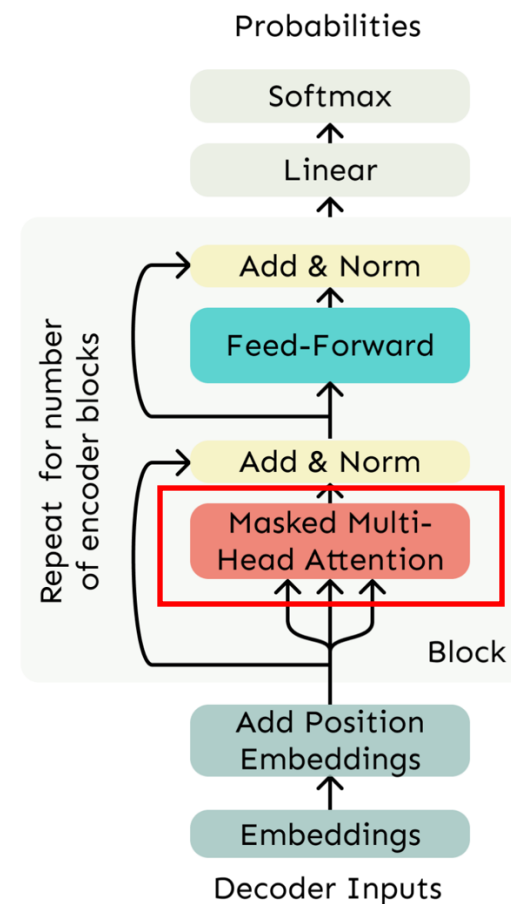

Embedding



Self-attention


- Number of heads: a

- Dims of each W_{Q_i} , W_{K_i} and W_{V_i} : $h \times \frac{h}{a}$




Self-attention

- Number of heads: a
- Dims of each W_{Q_i} , W_{K_i} and W_{V_i} : $h \times \frac{h}{a}$



Self-attention

- Number of heads: a
- Dims of each W_{Q_i} , W_{K_i} and W_{V_i} : $h \times \frac{h}{a}$




- Number of heads: a
- Dims of each W_{Q_i} , W_{K_i} and W_{V_i} : $h \times \frac{h}{a}$

$$\text{softmax}\left(\frac{Q_i K_i^T}{\sqrt{h/a}}\right) V_i = \text{softmax} \left[\begin{array}{|c|} \hline \text{Blue Grid (5x2)} \\ \hline \end{array} \times \begin{array}{|c|} \hline \text{Yellow Grid (2x5)} \\ \hline \end{array} \right] \times \begin{array}{|c|} \hline \text{Orange Grid (5x2)} \\ \hline \end{array} = \begin{array}{|c|} \hline \text{Light Blue Grid (5x2)} \\ \hline \end{array} \begin{matrix} h/a \\ s \\ Z_i \end{matrix}$$

One-head attention


Multi-head attentions

- Number of heads: a
- Dims of each W_{Q_i} , W_{K_i} and W_{V_i} : $h \times \frac{h}{a}$



Multi-head attentions

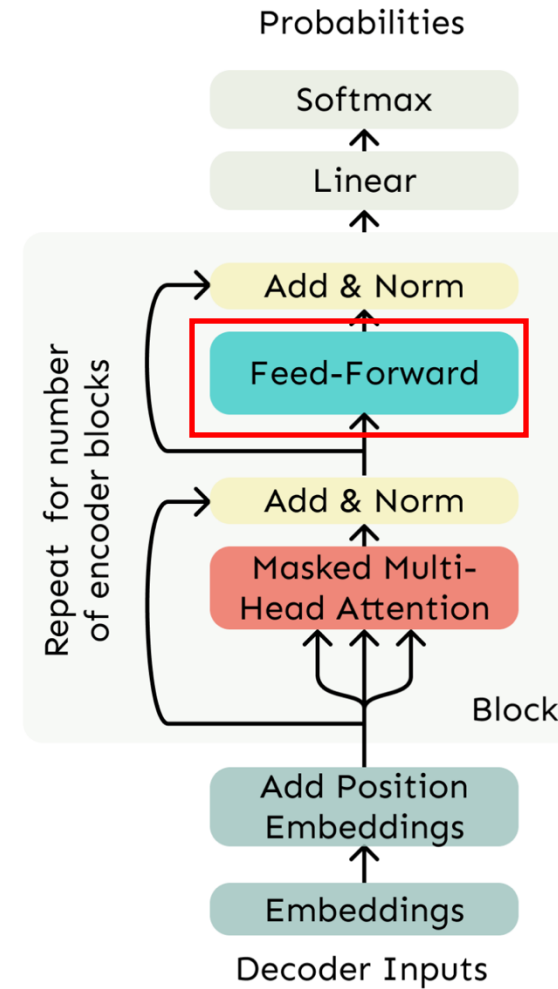
- Number of heads: a
- Dims of each W_{Q_i} , W_{K_i} and W_{V_i} : $h \times \frac{h}{a}$
- Dims of each W_O : $h \times h$



Feed-forward Layer

$$X' = \text{ReLU}(A \cdot W_1 + b_1) \cdot W_2 + b_2$$

- Dims of W_1 : $h \times 4h$
- Dims of each W_2 : $4h \times h$

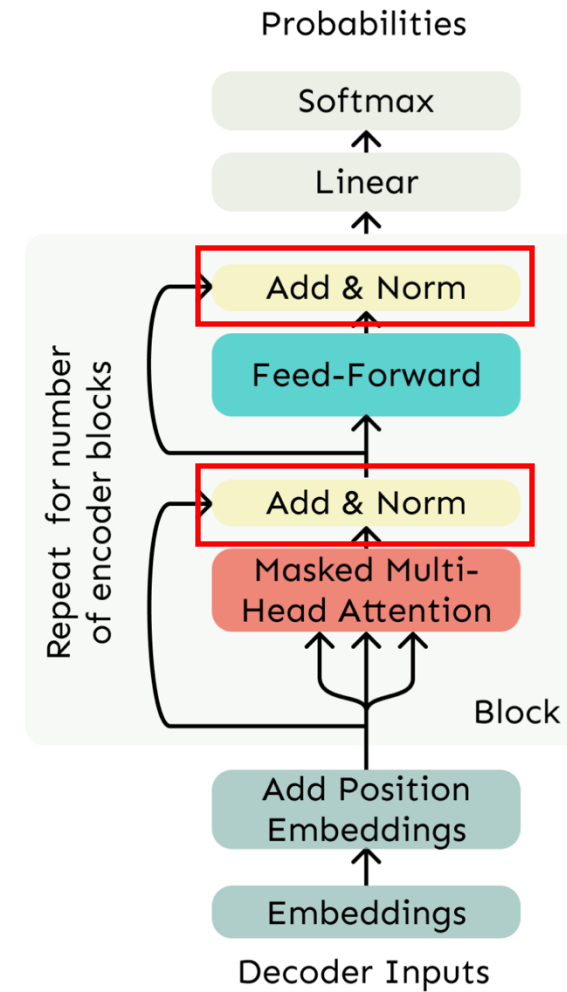


Layer normalization

- Then layer normalization computes:

$$\text{output} = \frac{x - \mu}{\sqrt{\sigma} + \epsilon} * \gamma + \beta$$

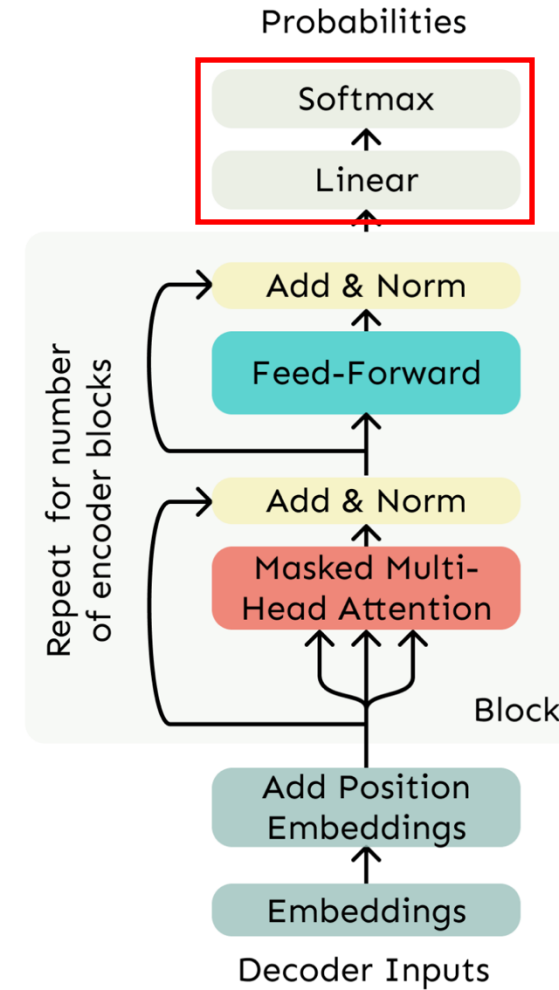
- Dims of γ and β : h



Probability prediction

$$p = \text{Softmax}(X \cdot W_v + b_v)$$

- Dims of W_v : $h \times v$




PART 02

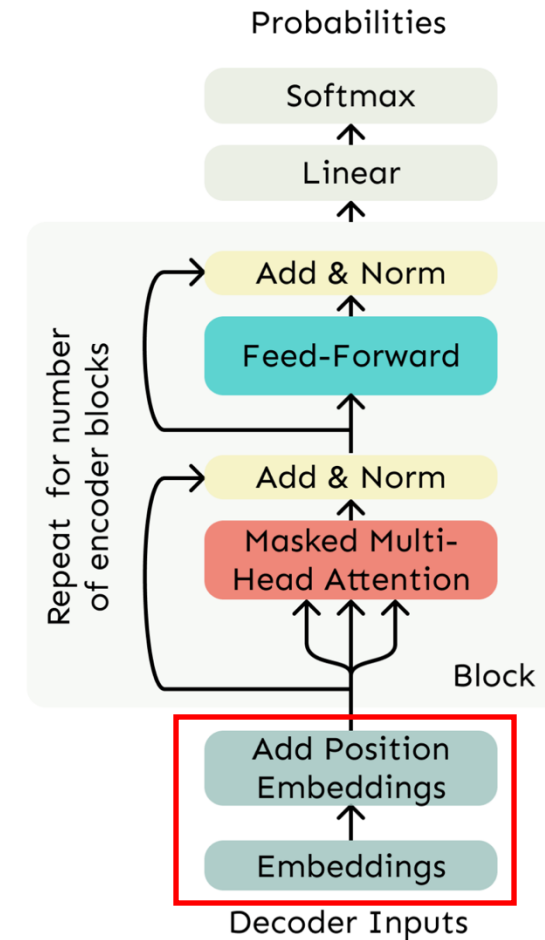
Parameters analysis



Embedding



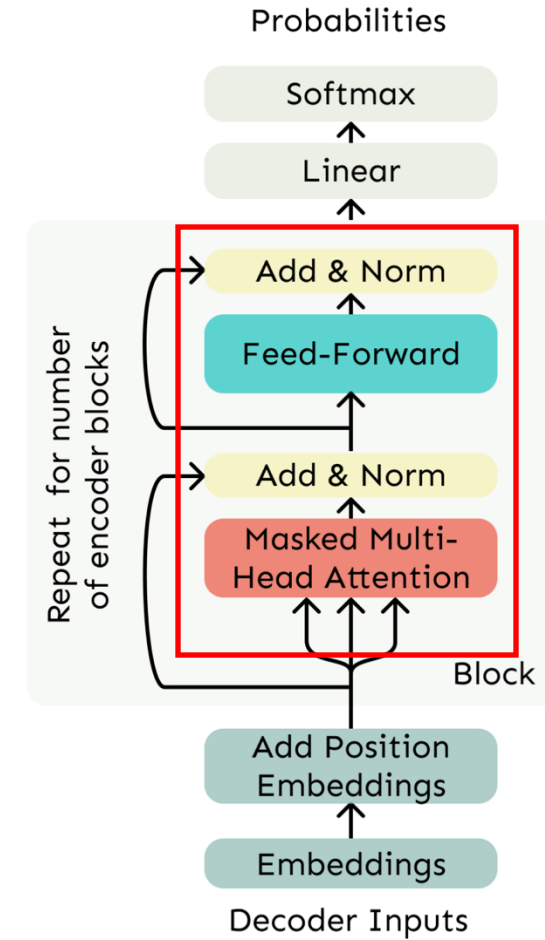

- We need to store the embedding with parameters vh
- Position embedding can be ignored when using RoPE and ALiBi




Multi-head attentions

- We need to store W_Q , W_K and W_V

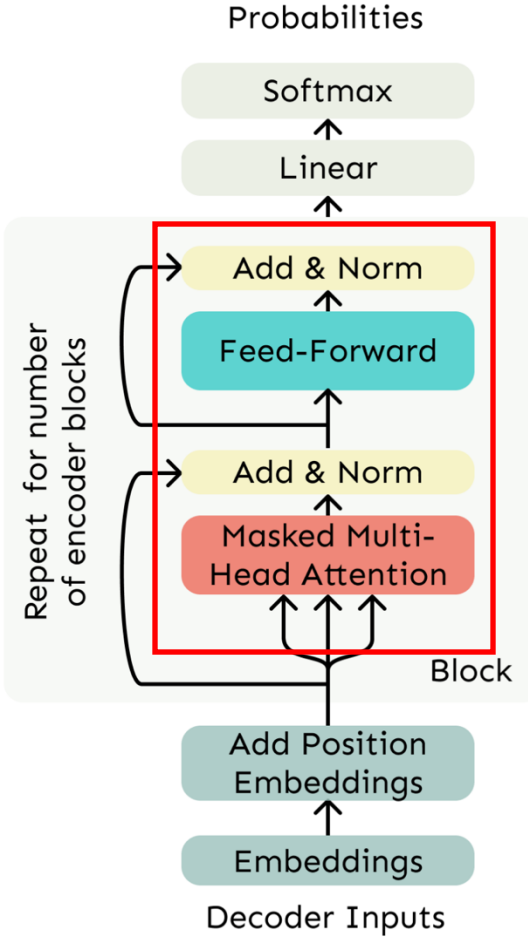
$$3(h^2/a) \times a = 3h^2$$



Multi-head attentions



- We need to store W_O : h^2

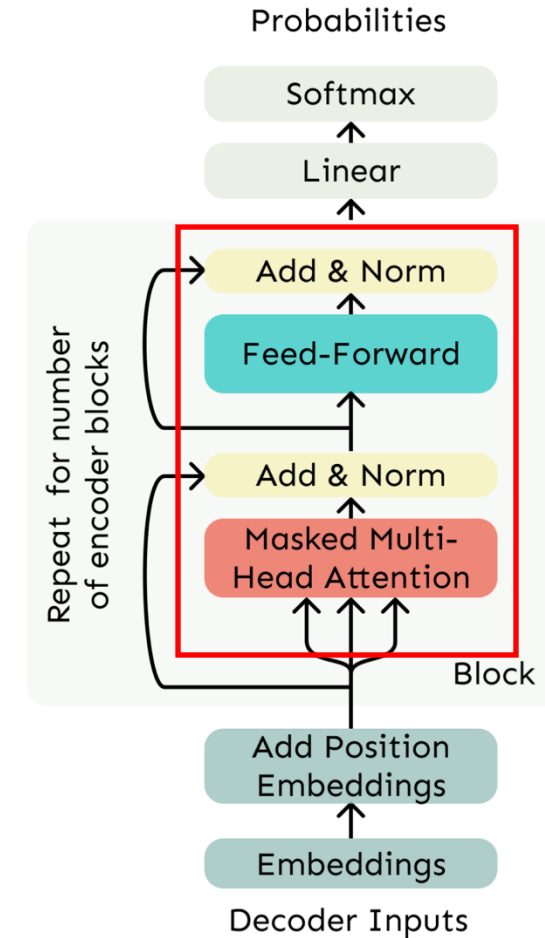


Layer normalization

- Then layer normalization computes:

$$\text{output} = \frac{x - \mu}{\sqrt{\sigma + \epsilon}} * \gamma + \beta$$


- Dims of γ and β : h
- We should store γ and β . Since their parameters are much smaller than h^2 and vh , we can ignore them



Feed-forward layers


$$X' = \text{ReLU}(A \cdot W_1 + b_1) \cdot W_2 + b_2$$

- Dims of W_1 : $h \times 4h$
- Dims of each W_2 : $4h \times h$
- We need to store W_1 and W_2 : $8h^2$
- The storage of b_1 and b_2 can be ignored



Transformer block

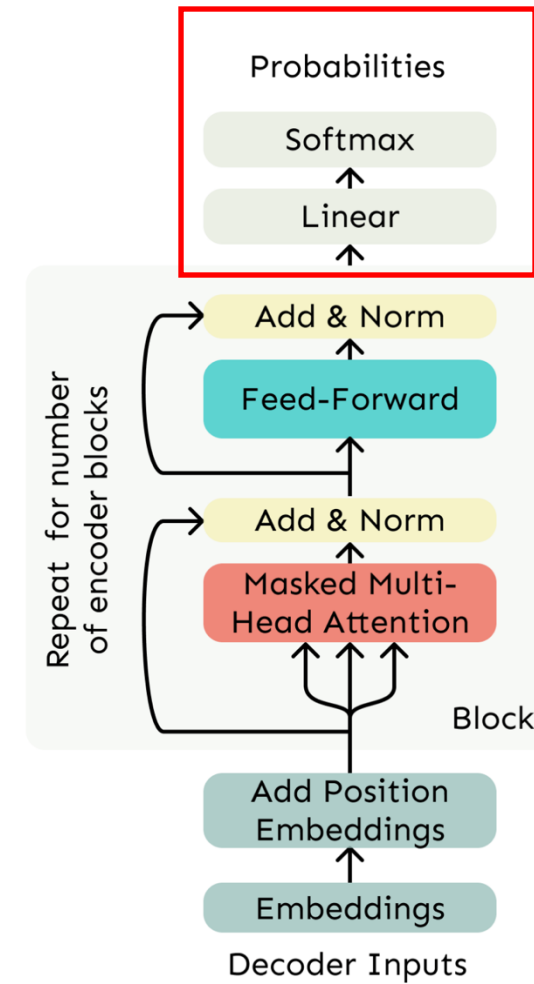
- Multi-head attentions: $4h^2$
- Feed-forward layers : $8h^2$
- l layers of attentions : $(4h^2 + 8h^2) \times l = 12lh^2$



Probability predictions

$$p = \text{Softmax}(X \cdot W_v + b_v)$$

- Dims of W_v : $h \times v$
- We need to store W_v : hv parameters
- b_v can be ignored

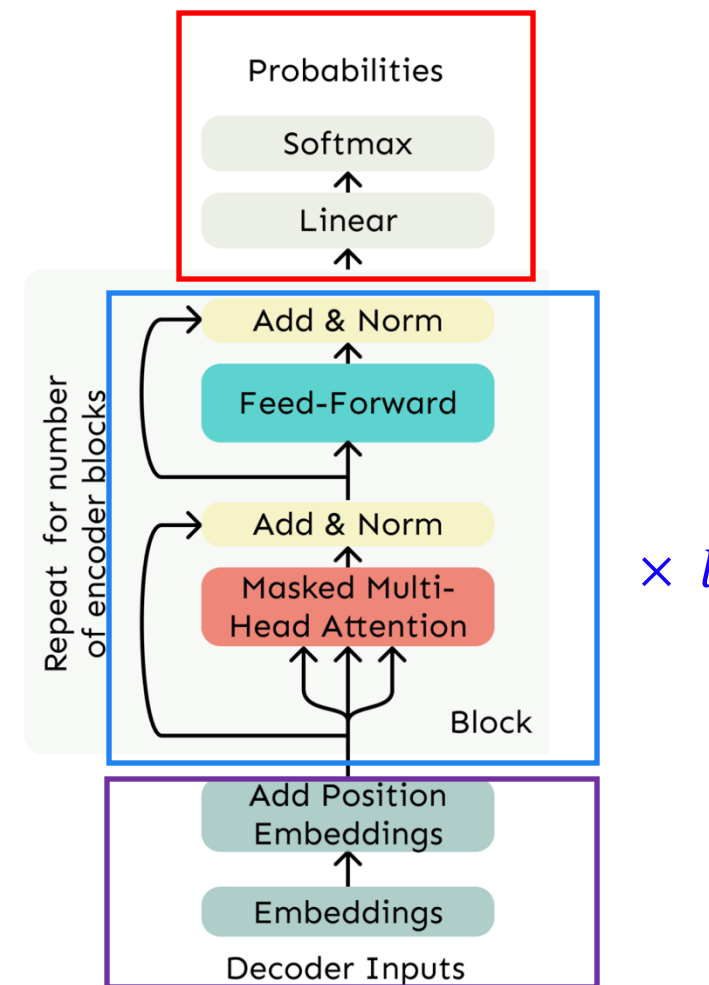


Total parameters

- Embeddings: vh
- Attention blocks: $12lh^2$
- Probability predictions: vh

Total parameters:

$$12\ell h^2 + 2vh$$



Example: LLaMA parameters

- Now we compare our theoretical evaluations with LLaMA model
- $12\ell h^2 + 2vh$ is a very accurate estimation

实际参数量	Embedding h	Attention层数l	Vocab大小v	预估参数量
6.7B	4096	32	32000	6,704,594,944
13.0B	5120	40	32000	12,910,592,000
32.5B	6656	60	32000	32,323,665,920
65.2B	8192	80	32000	64,948,797,440

PART 03


Computations analysis



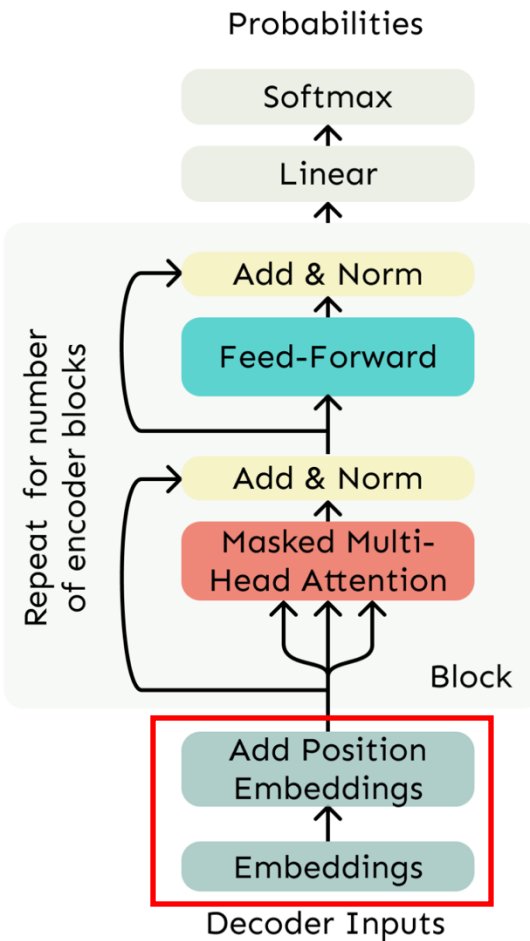
- FLOPs: Floating point operations; gauges the total amount of computations
- Given matrices $A \in \mathbb{R}^{m \times n}$, $B \in \mathbb{R}^{n \times p}$, to compute AB , we need

$$\left. \begin{array}{l} mnp \text{ additions} \\ mnp \text{ multiplications} \end{array} \right\} 2mnp \text{ FLOPs}$$

- In transformers, we only count computations raised by matrix operations and ignore vector operations since the later is trivial



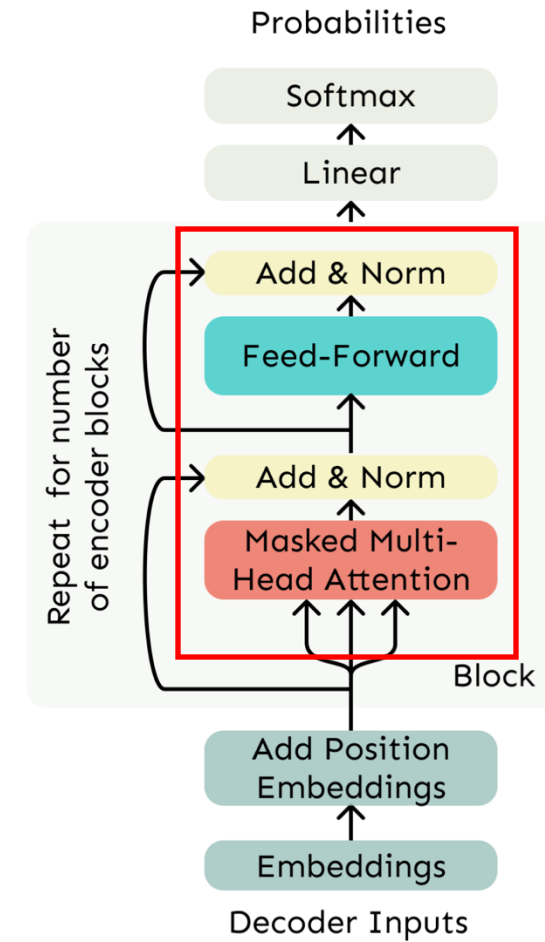

- Word embedding: $2svh$



Multi-head attentions

- Multi-head attentions

$$6(sh^2/a) \times a = 6sh^2$$




Multi-head attentions

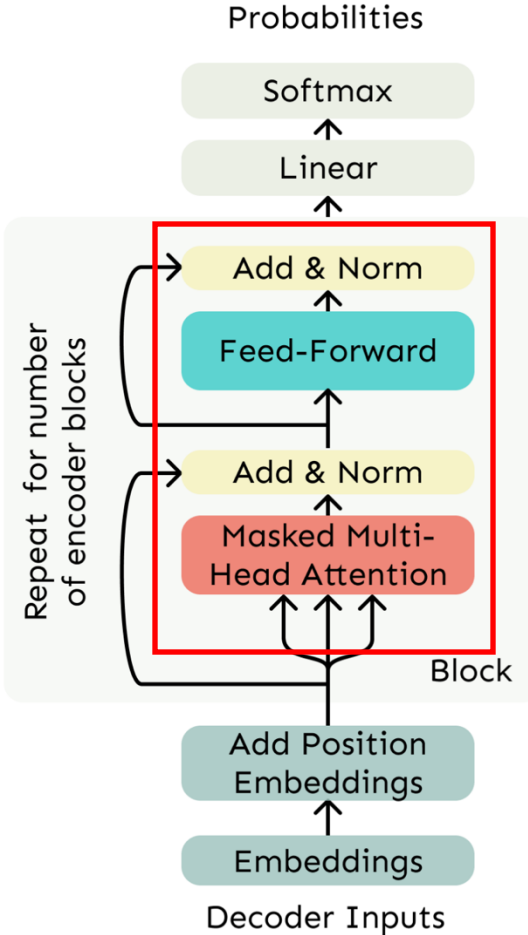
$$\text{softmax}\left(\frac{Q_i K_i^T}{\sqrt{h/a}}\right) V_i = \text{softmax} \left[\begin{array}{c} h/a \\ \begin{array}{|c|c|c|c|} \hline \text{blue grid} \\ \hline \end{array} \times \begin{array}{c} s \\ \begin{array}{|c|c|c|c|c|} \hline \text{yellow grid} \\ \hline \end{array} \end{array} \right] \times \begin{array}{c} \begin{array}{|c|c|c|c|} \hline \text{orange grid} \\ \hline \end{array} \end{array} = \begin{array}{c} h/a \\ \begin{array}{|c|c|c|c|} \hline \text{light blue grid} \\ \hline \end{array} s \\ Z_i \end{array}$$

$$(2s^2 h/a + 2s^2 h/a) \times a = 4s^2 h$$

Multi-head attentions



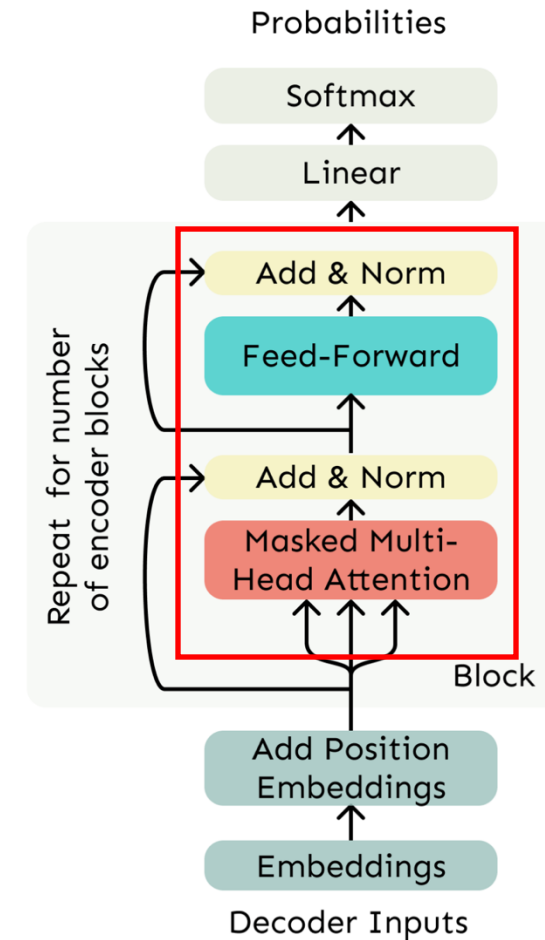
$2sh^2$ Flops



Feed-forward layers

$$X' = \text{ReLU}(A \cdot W_1 + b_1) \cdot W_2 + b_2$$

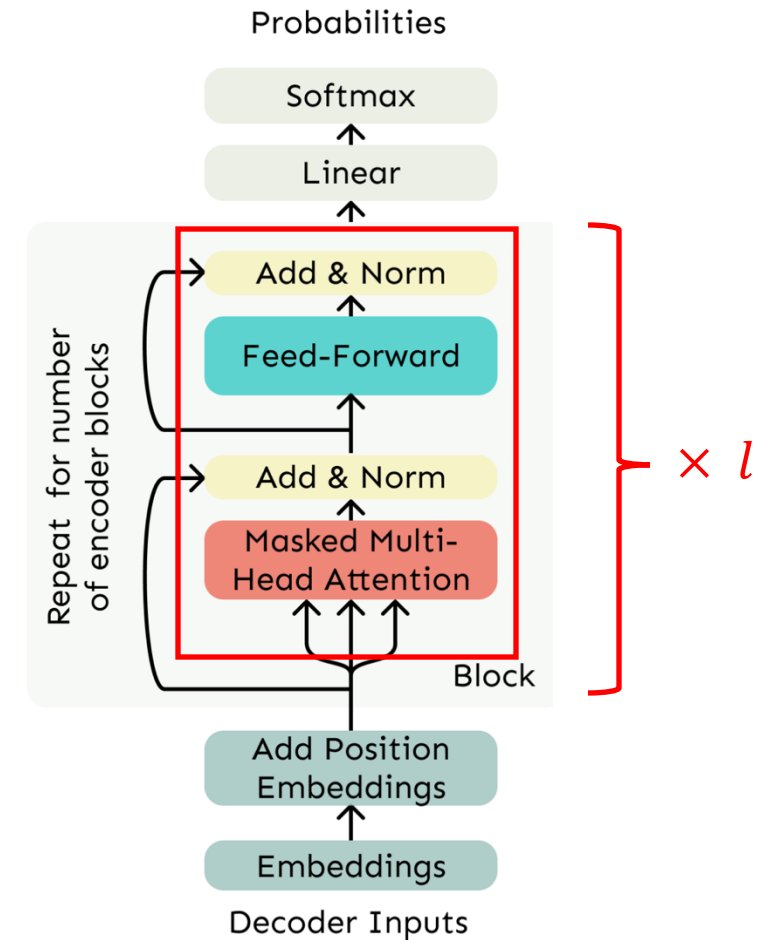
- Dims of W_1 : $h \times 4h$
 - Dims of each W_2 : $4h \times h$
 - $AW_1 + b_1$ needs: $8sh^2$
 - $A'W_2 + b_2$ needs: $8sh^2$
- } $16sh^2$



Transformer block

- Multi-head attentions: $8sh^2 + 4s^2h$
- Feed-forward layers : $16sh^2$
- l layers of attentions :

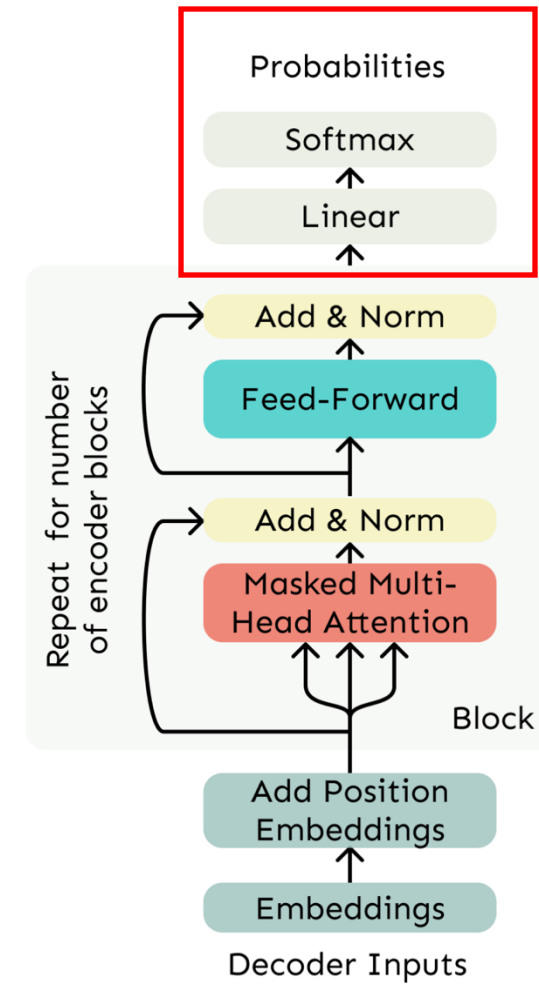
$$(8sh^2 + 16sh^2 + 4s^2h) \times l = 24slh^2 + 4s^2lh$$



Probability predictions

$$p = \text{Softmax}(X \cdot W_v + b_v)$$

- Dims of W_v : $h \times v$
- We need to : $2shv$ FLOPs



Total forward FLOPs

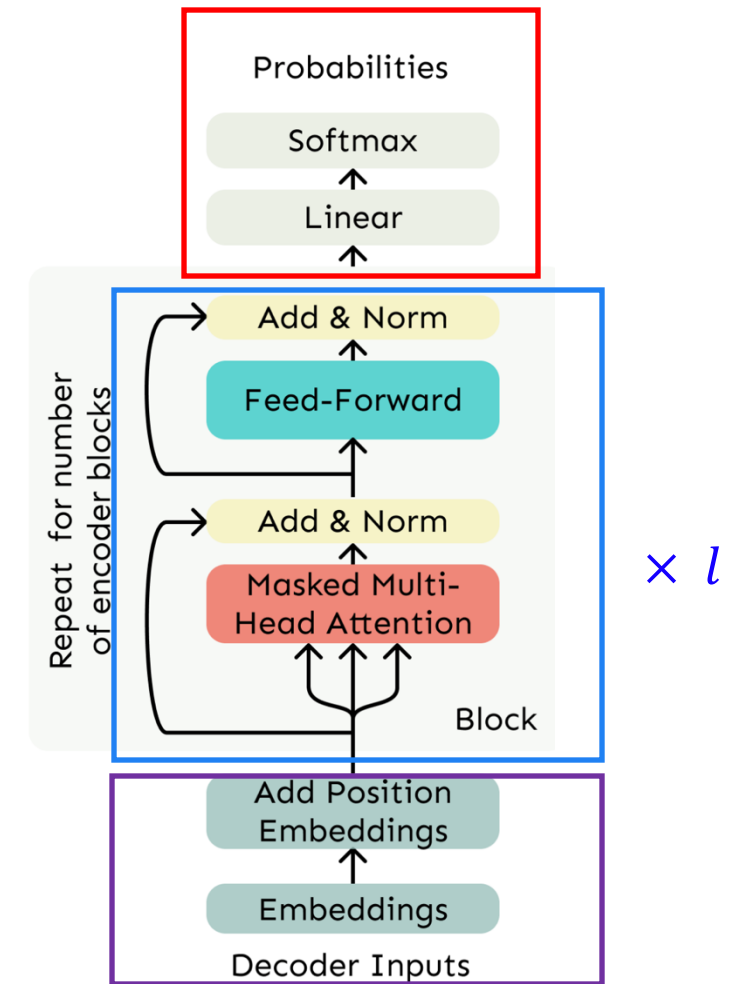
- Embeddings: $2svh$
- Attention blocks: $24lsh^2 + 4s^2lh$
- Probability predictions: $2svh$

Total forward FLOPs:

$$l(24sh^2 + 4s^2h) + 4svh$$

When using batch-size b , the total forward FLOPs:

$$bl(24sh^2 + 4s^2h) + 4bsvh$$



Total forward-backward FLOPs

The backward computations are **twice amount** of the forward computations

Total forward-backward FLOPs

$$\left(b\ell(24sh^2 + 4s^2h) + 4bsvh \right) \times 3 = 3 \left(b\ell(24sh^2 + 4s^2h) + 4bsvh \right)$$

When h^2 dominates, the above FLOPs can be simplified as $72bs\ell h^2$

When h^2 dominates, the parameters can be simplified as $P = 12\ell h^2$

Since $T = bs$ is the number of tokens, we thus have $\text{FLOPs} = 6TP$

Example: GPT3-175B

GPT-175B: 175B parameters, 300B tokens

$$6 \times 174600 \times 10^6 \times 300 \times 10^9 = 3.1428 \times 10^{23} \text{ flops}$$

Model	Total train compute (PF-days)	Total train compute (flops)	Params (M)	Training tokens (billions)	Flops per param per token	Mult for bwd pass	flops per active param per token	Fraction of params active for each token
T5-Small	2.08E+00	1.80E+20	60	1,000	3	3	1	0.5
T5-Base	7.64E+00	6.60E+20	220	1,000	3	3	1	0.5
T5-Large	2.67E+01	2.31E+21	770	1,000	3	3	1	0.5
T5-3B	1.04E+02	9.00E+21	3,000	1,000	3	3	1	0.5
T5-11B	3.82E+02	3.30E+22	11,000	1,000	3	3	1	0.5
BERT-Base	1.89E+00	1.64E+20	109	250	6	3	2	1.0
BERT-Large	6.16E+00	5.33E+20	355	250	6	3	2	1.0
RoBERTa-Base	1.74E+01	1.50E+21	125	2,000	6	3	2	1.0
RoBERTa-Large	4.93E+01	4.26E+21	355	2,000	6	3	2	1.0
GPT-3 Small	2.60E+00	2.25E+20	125	300	6	3	2	1.0
GPT-3 Medium	7.42E+00	6.41E+20	356	300	6	3	2	1.0
GPT-3 Large	1.58E+01	1.37E+21	760	300	6	3	2	1.0
GPT-3 XL	2.75E+01	2.38E+21	1,320	300	6	3	2	1.0
GPT-3 2.7B	5.52E+01	4.77E+21	2,650	300	6	3	2	1.0
GPT-3 6.7B	1.39E+02	1.20E+22	6,660	300	6	3	2	1.0
GPT-3 13B	2.68E+02	2.31E+22	12,850	300	6	3	2	1.0
GPT-3 175B	3.64E+03	3.14E+23	174,600	300	6	3	2	1.0