


Memory Analysis in Transformers

Kun Yuan

Center for Machine Learning Research @ Peking University



Warmup: Linear neural network



$$h = W_1 x$$

$$z = \sigma(h)$$

$$\hat{y} = W_2 z$$

$$f = L(\hat{y})$$

Forward

Store h, z and \hat{y}

$$\frac{\partial f}{\partial W_1} = \frac{\partial f}{\partial h} x^T$$

$$\frac{\partial f}{\partial h} = \frac{\partial f}{\partial z} \odot \nabla \sigma(h)$$


$$\frac{\partial f}{\partial W_2} = \frac{\partial L}{\partial \hat{y}} z^T, \quad \frac{\partial f}{\partial z} = W_2^T \frac{\partial L}{\partial \hat{y}}$$

$$\frac{\partial f}{\partial \hat{y}} = \nabla L(\hat{y})$$

Backward

Store $\nabla_{W_1} f(W_1)$ and $\nabla_{W_2} f(W_2)$

Warmup: Linear neural network with batch size



dims: d (p,d) p (q,p) q

$$h_b = W_1 x_b$$

$$z_b = \sigma(h_b)$$

$$\hat{y}_b = W_2 z_b$$

$$f = \frac{1}{B} \sum_{b=1}^B L(\hat{y}_b)$$

Forward

$$\frac{\partial f}{\partial W_1} = \frac{1}{B} \sum_{b=1}^B \frac{\partial f}{\partial h_b} x_b^T,$$

$$\frac{\partial f}{\partial h_b} = \frac{\partial f}{\partial z_b} \odot \nabla \sigma(h_b)$$

$$\frac{\partial f}{\partial W_2} = \frac{1}{B} \sum_{b=1}^B \frac{\partial L}{\partial \hat{y}_b} z_b^T, \quad \frac{\partial f}{\partial z_b} = W_2 \frac{\partial f}{\partial \hat{y}_b}$$

$$\frac{\partial f}{\partial \hat{y}_b} = \frac{\partial L}{\partial \hat{y}_b}$$

Backward

Store $\{h_b, z_b, \hat{y}_b\}_{b=1}^B$

Store $\nabla_{W_1} f(W_1)$ and $\nabla_{W_2} f(W_2)$

Memory = Model + Gradient + Optimizer states + Activations

- Given a model with P parameters, gradient will consume P parameters, and Optimizer states will consume $2P$ parameters; **$4P$ parameters in total.**
- When using FP32 to store parameters, each parameter takes **4** Bytes
- When using FP16 or BF16 to store parameters, each parameter takes **2** Bytes

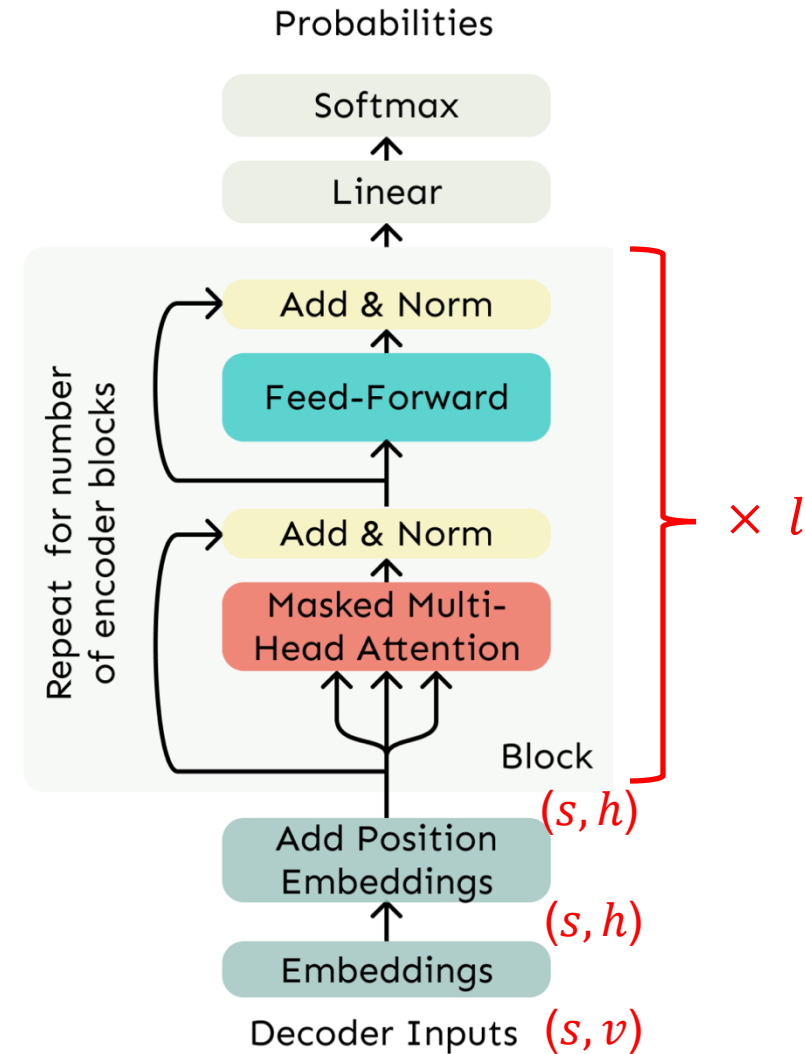
Decoder-only Transformers

- Number of the transformer layers: l
- Sequence length: s
- Vocabulary size: v
- Embedding representation dims: h


Parameters $P = 12\ell h^2 + 2vh$

Model + Gradient + Optimizer states = $4P * 4$ (Bytes)

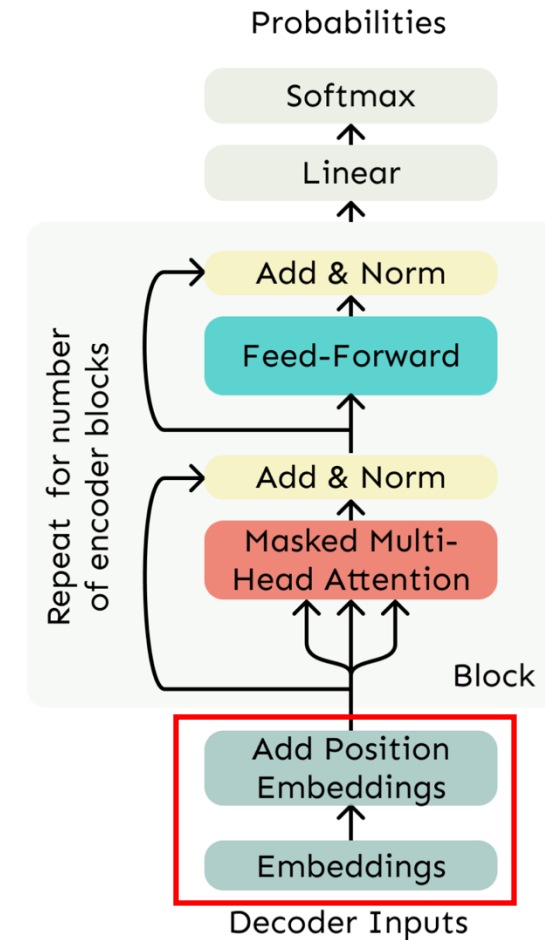
Next we estimate activations



Embedding





- We need to store the embedding activations with parameters sh
- Position embedding can be ignored when using RoPE and ALiBi



Multi-head attentions

- We need to store Q_i , K_i and V_i

$$3(sh/a) \times a = 3sh$$




$$\text{softmax}\left(\frac{Q_i K_i^T}{\sqrt{h/a}}\right) V_i = \text{softmax} \left[\begin{array}{|c|c|} \hline \text{Blue Grid} & \text{Yellow Grid} \\ \hline \end{array} \right] \times \begin{array}{|c|c|} \hline \text{Orange Grid} \\ \hline \end{array} = \begin{array}{|c|c|} \hline \text{Light Blue Grid} \\ \hline \end{array} \begin{matrix} h/a \\ s \\ Z_i \end{matrix}$$


One-head attention

- Store $Q_i K_i^T$ with s^2 parameters;
 - Store $\text{softmax}(Q_i K_i^T)$ with s^2 parameters
 - Store $\text{softmax}\left(\frac{Q_i K_i^T}{\sqrt{h/a}}\right) V_i$ with sh/a parameters
- } Store $2s^2 a + sh$ activations

Multi-head attentions



- We need to store A with sh parameters




Layer normalization

- Then layer normalization computes:

$$\text{output} = \frac{x - \mu}{\sqrt{\sigma + \epsilon}} * \gamma + \beta$$


- Dims of γ and β : h
- We can ignore the layer normalization



Feed-forward layers


$$X' = \text{ReLU}(A \cdot W_1 + b_1) \cdot W_2 + b_2$$

- Store AW_1 with 4sh parameters
 - Store $\text{ReLU}(AW_1)$ with 4sh parameters
 - Store $\text{ReLU}(AW_1) \cdot W_2$ with sh parameters
 - The activations of b_1 and b_2 can be ignored
- 9sh



Transformer block

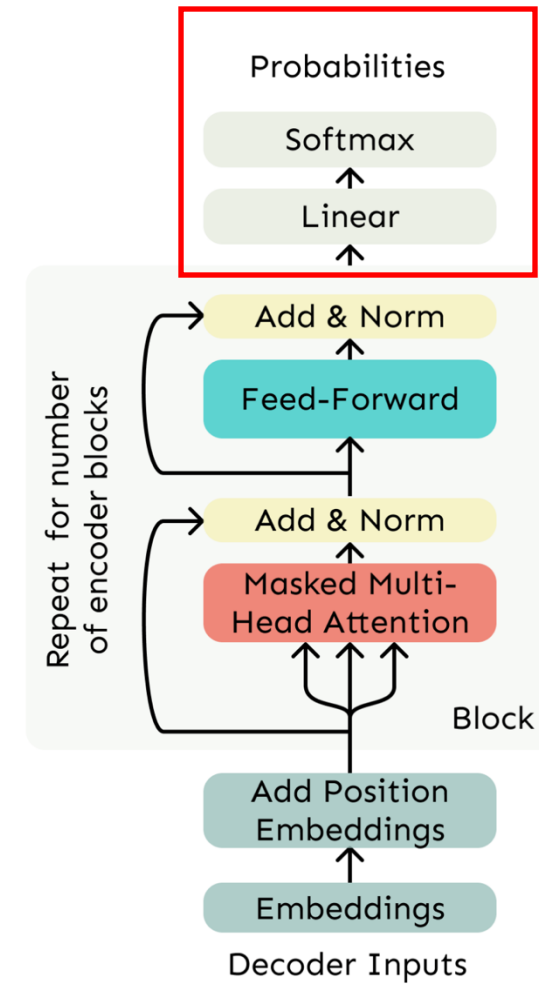
- Multi-head attentions: $5sh + 2s^2a$
- Feed-forward layers : $9sh$
- l layers of attentions : $(2s^2a + 14sh) \times l$



Probability predictions

$$p = \text{Softmax}(X \cdot W_v + b_v)$$

- Store XW_v with sv parameters
- Store $\text{Softmax}(XW_v)$ with sv parameters



Total activations

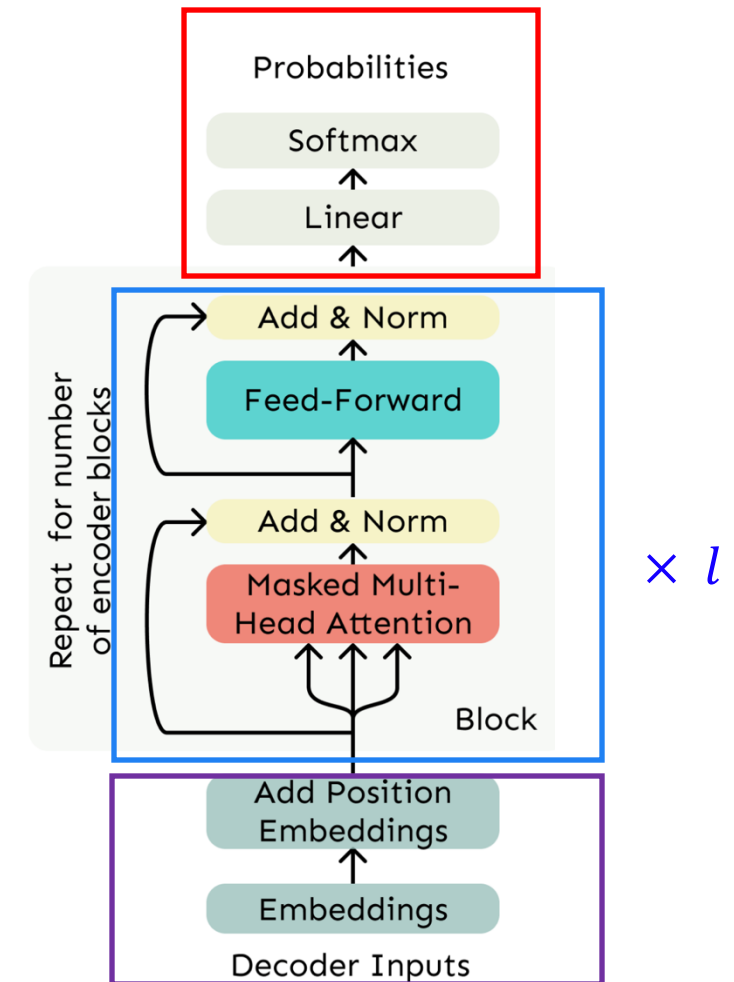
- Embedding activations: sh
- Self-attention activations: $(2s^2a + 14sh) \times l$
- Probability activations: $2sv$

Total activation parameters with batch-size 1:

$$(2s^2a + 14sh) \times l + 2sv + sh$$

Ignoring $2sv$ and sh and using batch-size b :

$$(2s^2a + 14sh) \times l \times b$$



Memory = **Model + Gradient + Optimizer states** + **Activations**

$$(48 l h^2 + bl(2s^2 a + 14sh)) \times 4 \text{ Bytes}$$

- When hidden state h is large, the model parameters dominate the memory
- When batch-size b or sequence length s is large, the activation dominates the memory

Examples

Model Name	n_{params}	n_{layers}	d_{model}	n_{heads}	d_{head}	Batch Size	Learning Rate
GPT-3 Small	125M	12	768	12	64	0.5M	6.0×10^{-4}
GPT-3 Medium	350M	24	1024	16	64	0.5M	3.0×10^{-4}
GPT-3 Large	760M	24	1536	16	96	0.5M	2.5×10^{-4}
GPT-3 XL	1.3B	24	2048	24	128	1M	2.0×10^{-4}
GPT-3 2.7B	2.7B	32	2560	32	80	1M	1.6×10^{-4}
GPT-3 6.7B	6.7B	32	4096	32	128	2M	1.2×10^{-4}
GPT-3 13B	13.0B	40	5140	40	128	2M	1.0×10^{-4}
GPT-3 175B or “GPT-3”	175.0B	96	12288	96	128	3.2M	0.6×10^{-4}

GPT3 has 175B parameters; its model consumes $4 \times 175 \times 10^9$ Bytes = 700 GB

Its gradient takes 700 GB parameters; Optimizer states take 1.4 TB

GPT has sequence length $s = 2048$. When $b=1$, its activation takes 444 GB, 63% of the model

When $b=128$, its activation is 81 times of the model size

Memory = **Model + Gradient + Optimizer states** + **Activations**

$$(48 l h^2 + bl(2s^2 a + 14sh)) \times 4 \text{ Bytes}$$

- When hidden state h is large, the model parameters dominate the memory
- When batch-size b or sequence length s is large, the activation dominates the memory
- The activation-incurred memory **cannot be ignored**