

Subspace Training in LLMs

Kun Yuan (袁坤)

Center for Machine Learning Research @ Peking University




PART 01

Basics and Motivation



LLM pretraining is essentially solving stochastic optimization

- The model weights in neural networks are a set of matrices $\mathbf{X} = \{\mathbf{X}_\ell\}_{\ell=1}^L$



- Let $h(\mathbf{X}; \xi)$ be the language model; $\hat{y} = h(\mathbf{X}; \xi)$ is the predicted token

LLM cost function:

$$\mathbf{X}^* = \arg \min_{\mathbf{X}} \left\{ \mathbb{E}_{\xi \sim \mathcal{D}} \left[\overset{\text{cross entropy}}{\underset{\substack{\text{data distribution} \quad \text{pred. token} \quad \text{real token}}}{L(h(\mathbf{X}; \xi), y)}} \right] \right\}$$

LLM pretraining is essentially solving stochastic optimization

- If we define $\xi = (\xi, y)$ and $F(\mathbf{X}; \xi) = L(h(\mathbf{X}; \xi), y)$, the LLM problem becomes

$$\text{Stochastic optimization: } \mathbf{X}^* = \arg \min_{\mathbf{X}} \left\{ \mathbb{E}_{\xi \sim \mathcal{D}} [F(\mathbf{X}; \xi)] \right\}$$

- In other words, LLM pretraining is essentially solving a stochastic optimization problem
- Adam is the standard approach in LLM pretraining

Optimizer states	$\mathbf{G}_t = \nabla F(\mathbf{X}_t; \xi_t)$	(stochastic gradient)
	$\mathbf{M}_t = (1 - \beta_1)\mathbf{M}_{t-1} + \beta_1 \mathbf{G}_t$	(first-order momentum)
	$\mathbf{V}_t = (1 - \beta_2)\mathbf{V}_{t-1} + \beta_2 \mathbf{G}_t \odot \mathbf{G}_t$	(second-order momentum)
	$\mathbf{X}_{t+1} = \mathbf{X}_t - \frac{\gamma}{\sqrt{\mathbf{V}_t} + \epsilon} \odot \mathbf{M}_t$	(adaptive SGD)

Memory = **Model** + **Gradient** + **Optimizer states** + **Activations**

- Given a **model** with **P** parameters, **gradient** will consume **P** parameters, and **optimizer states** will consume **2P** parameters; **4P parameters in total**.

P $G_t = \nabla F(\mathbf{X}_t; \xi_t)$

2P $\left\{ \begin{array}{l} M_t = (1 - \beta_1)M_{t-1} + \beta_1 G_t \\ V_t = (1 - \beta_2)V_{t-1} + \beta_2 G_t \odot G_t \end{array} \right.$

P $\mathbf{X}_{t+1} = \mathbf{X}_t - \frac{\gamma}{\sqrt{V_t} + \epsilon} \odot M_t$

Optimizer states introduces significant memory cost

$$\text{Memory} = \text{Model} + \text{Gradient} + \text{Optimizer states} + \text{Activations}$$

- Activations are auxiliary variables to facilitate the gradient calculations

Consider a linear neural network

$$z_i = X_i z_{i-1}, \forall i = 1, \dots, L$$

$$f = \mathcal{L}(z_L; y)$$

The gradient is derived as follows

$$\frac{\partial f}{\partial X_i} = \frac{\partial f}{\partial z_i} z_{i-1}^\top$$


Need to store activations z_1, z_2, \dots, z_L

- The size of activations depends on sequence length and batch size

Minimum memory requirement

- Pretrain LLaMA 7B model (BF16) from scratch with **a single batch size** requires

- Parameters: 7B
- Model storage: $7\text{B} * 2 \text{ Bytes} = 14 \text{ GB}$
- Gradient storage: 14 GB
- Optimizer states: 28 GB (using Adam)
- Activation storage: 2 GB [Zhao et. al., 2024]
- In total: **58 GB**



Minimum memory requirement: LLaMA 7B

- Pretrain LLaMA 7B model (BF16) from scratch with **a single batch size** requires

- Parameters: 7B
- Model storage: $7B * 2 \text{ Bytes} = 14 \text{ GB}$
- Gradient storage: 14 GB
- Optimizer states: 28 GB (using Adam)
- Activation storage: 2 GB [Zhao et. al., 2024]
- In total: **58 GB**



RTX 4090: 24GB




A100 80G



The minimum requirement is A100 * 1

Minimum memory requirement: GPT-3

- Pretrain GPT-3 model (BF16) from scratch with **a single batch size** requires
 - Parameters: 175B
 - Model storage: $175\text{B} * 2 \text{ Bytes} = 350 \text{ GB}$
 - Gradient storage: 350 GB
 - Optimizer states: 700 GB (using Adam)
 - Activation storage: ~220 GB
 - In total: **1620 GB**



Minimum memory requirement: GPT-3

- Pretrain GPT-3 model (BF16) from scratch with **a single batch size** requires
 - Parameters: 175B
 - Model storage: $175\text{B} * 2 \text{ Bytes} = 350 \text{ GB}$
 - Gradient storage: 350 GB
 - Optimizer states: 700 GB (using Adam)
 - Activation storage: ~220 GB
 - In total: **1620 GB**



x 21


A100 80G

The minimum requirement is $A100 * 21$

Very expensive!

Memory-efficient algorithm is in urgent need

- With memory-efficient algorithms, we can
 - **Train larger models** on limited computing resources
 - Use a larger training batch size to **improve throughput**
- Activation-incurred memory is **relatively minor** when using a single batch size
- Gradient-incurred memory can be **removed** by layer-wise calculation and dropping
- How to save memory caused by optimizer states?



PART 02

Subspace training



- Consider the following optimization problem

$$\min_{\mathbf{x}} f(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_s)$$

where $\mathbf{x} \in \mathbb{R}^d$ is decomposed into s block variables $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_s$

- Block coordinate descent

$$\mathbf{x}_j^k = \arg \min_{\mathbf{x}_{i_k}} f(\mathbf{x}_{i_k}, \mathbf{x}_{\neq i_k}^{k-1}), \quad \mathbf{x}_j^k = \mathbf{x}_j^{k-1} \quad \text{if } j \neq i_k$$

where $\mathbf{E}_{i_k} = [\mathbf{0}; \mathbf{0}; \dots; \mathbf{I}_{i_k}; \dots; \mathbf{0}] \in \mathbb{R}^{d \times d_{i_k}}$. The above recursion can be rewritten as

$$\mathbf{b}^k = \arg \min_{\mathbf{b}} f(\mathbf{x}^{k-1} + \mathbf{E}_{i_k} \mathbf{b}), \quad \mathbf{x}^k = \mathbf{x}^{k-1} + \mathbf{E}_{i_k} \mathbf{b}^k$$

- We only minimize a small block variable \mathbf{b}^k each time, which saves memory

- Now we consider optimization with matrix variables

$$\min_{X \in \mathbb{R}^{m \times n}} f(X)$$

Minimize X directly would result in large memory cost

- Inspired by block coordinate descent, we consider the subspace optimization

$$B^k = \arg \min_{B \in \mathbb{R}^{r \times n}} f(X^{k-1} + P^k B), \quad X^k = X^{k-1} + P^k B^k$$

where $P^k \in \mathbb{R}^{m \times r}$ is a randomly chosen matrix.

- **Instead of solving X directly, we solve subproblems with a smaller matrix variable**

Subspace optimization: GD variant

$$B^k = \arg \min_{B \in \mathbb{R}^{r \times n}} f(X^{k-1} + P^k B), \quad X^k = X^{k-1} + P^k B^k$$

- Now we consider solving the subproblem with GD, the above problem becomes

$$\begin{aligned} B^{(k,t)} &= B^{(k,t-1)} - \gamma (P^k)^\top \nabla_X f(X^{k-1} + P^k B^{(k,t-1)}), \quad \forall t = 1, 2, \dots, \tau \\ X^k &= X^{k-1} + P^k B^{(k,\tau)} \end{aligned}$$

- We let $X^{(k-1,t)} = X^{k-1} + P^k B^{(k,t)}$, the above method reduces to

$$\begin{aligned} X^{(k-1,t)} &= X^{(k-1,t-1)} - \gamma P^k (P^k)^\top \nabla_X f(X^{(k-1,t-1)}), \quad \forall t = 1, \dots, \tau \\ X^{(k,0)} &= X^{(k-1,\tau)} \end{aligned}$$

Subspace optimization: GD variant

- We let $X^{(k-1,t)} = X^{k-1} + P^k B^{(k,t)}$, the above method reduces to

$$\begin{aligned} X^{(k-1,t)} &= X^{(k-1,t-1)} - \gamma P^k (P^k)^\top \nabla_X f(X^{(k-1,t-1)}), \quad \forall t = 1, \dots, \tau \\ X^{(k,0)} &= X^{(k-1,\tau)} \end{aligned}$$

- The above algorithm can be rewritten as follows

$$\begin{aligned} X^t &= X^{t-1} - \gamma P^t (P^t)^\top \nabla f(X^{t-1}), \quad \forall t = 1, \dots, T \\ P^t &= \begin{cases} \text{Sample new } P & \text{if } \text{mod}(t, \tau) = 0 \\ P^{t-1} & \text{otherwise} \end{cases} \end{aligned}$$

Subspace optimization: advanced variant

$$B^k = \arg \min_{B \in \mathbb{R}^{r \times n}} f(X^{k-1} + P^k B), \quad X^k = X^{k-1} + P^k B^k$$

- Now we consider solving the subproblem with advanced approach $\rho(\cdot)$

$$B^{(k,t)} = B^{(k,t-1)} - \gamma \rho \left((P^k)^\top \nabla_X f(X^{k-1} + P^k B^{(k,t-1)}) \right), \quad \forall t = 1, 2, \dots, \tau$$

$$X^k = X^{k-1} + P^k B^{(k,\tau)}$$

- The operator $\rho(\cdot)$ can be either Momentum GD or Adam

(momentum)

$$m^t = (1 - \beta)m^{t-1} + \beta g^t$$
$$\rho(g^t) = m^t$$

$$m^t = (1 - \beta_1)m^{t-1} + \beta_1 g^t$$
$$v^t = (1 - \beta_2)v^{t-1} + \beta_2 g^t \odot g^t \quad (\text{Adam})$$
$$\rho(g^t) = \frac{m^t}{\sqrt{v^t} + \epsilon}$$

Subspace optimization: advanced variant

$$B^{(k,t)} = B^{(k,t-1)} - \gamma \boldsymbol{\rho} \left((P^k)^\top \nabla_X f(X^{k-1} + P^k B^{(k,t-1)}) \right), \quad \forall t = 1, 2, \dots, \tau$$

$$X^k = X^{k-1} + P^k B^{(k,\tau)}$$

- The above algorithm can be rewritten as follows

$$X^t = X^{t-1} - \gamma P^t \boldsymbol{\rho} \left((P^t)^T \nabla f(X^{t-1}) \right), \quad \forall t = 1, \dots, T$$

$$P^t = \begin{cases} \text{Sample new } P & \text{if } \text{mod}(t, \tau) = 0 \\ P^{t-1} & \text{otherwise} \end{cases}$$

PART 03

GaLore Algorithm



- In stochastic scenario, the algorithm in Page 18 recovers GaLore

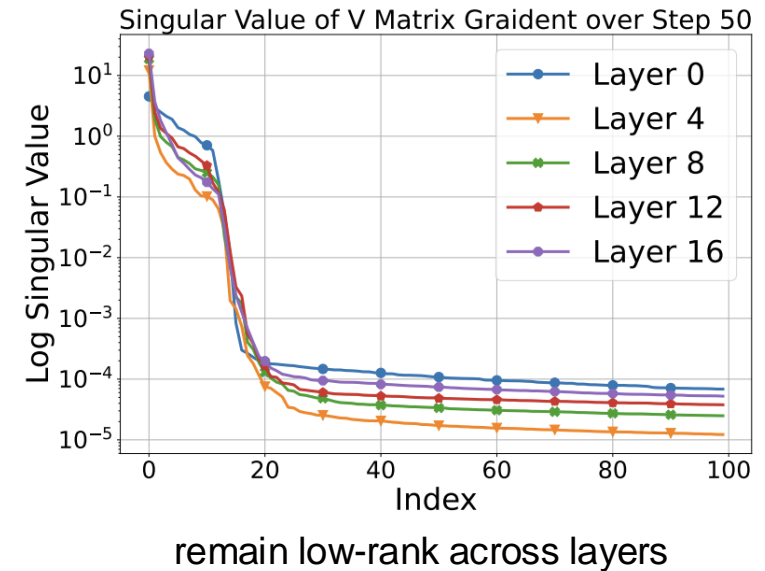
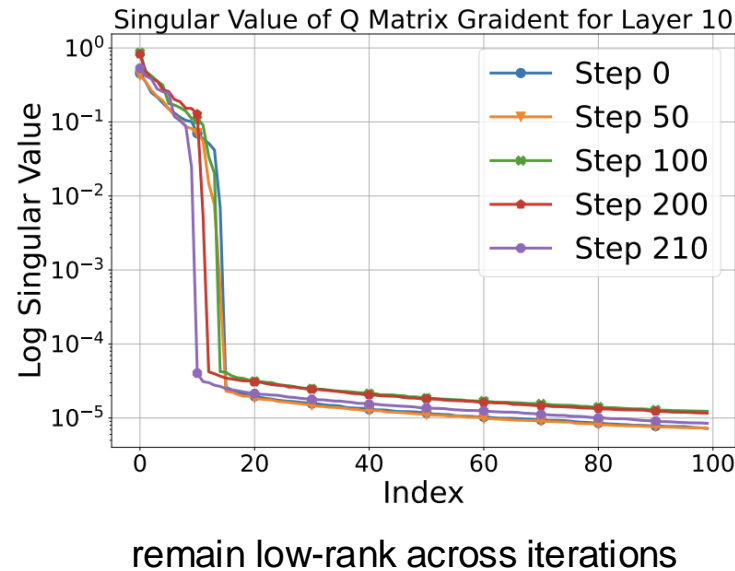

GaLore: Memory-Efficient LLM Training by Gradient Low-Rank Projection

Jiawei Zhao¹ Zhenyu Zhang³ Beidi Chen^{2,4} Zhangyang Wang³ Anima Anandkumar^{*1} Yuandong Tian^{*2}

-
- GaLore is a novel approach for reducing memory consumption from optimizer states
 - The first algorithm that enables LLaMA-7B pre-training on a single 4090 GPU (24GB)
 - Memory-efficient without severe performance degradation

GaLore: Gradient Low-Rank Projection



- Stochastic optimization: $\mathbf{X}^* = \arg \min_{\mathbf{X}} \left\{ \mathbb{E}_{\xi \sim \mathcal{D}} [F(\mathbf{X}; \xi)] \right\}$



- Given a gradient matrix with dimensions 2048 by 2048, around **top 10** eigenvalues dominate
- How to utilize the **low-rank** structure in gradients?

GaLore: Gradient Low-Rank Projection

- Main idea: Projecting gradient onto the low-rank subspace
- Given a gradient $G_t \in \mathbb{R}^{m \times n}$ and a projection $P_t \in \mathbb{R}^{m \times r}$, we project Gradient onto low-rank subspace



- Since $r \ll m$, the low-rank gradient g_t has much smaller parameters than G_t

GaLore: Gradient Low-Rank Projection

- Low-rank optimizer states:

$$\mathbf{g}_t = \mathbf{P}_t^\top \mathbf{G}_t$$

▷ dims $r \times n$

$$\mathbf{m}_t = (1 - \beta_1)\mathbf{m}_{t-1} + \beta_1\mathbf{g}_t$$

▷ dims $r \times n$

$$\mathbf{v}_t = (1 - \beta_2)\mathbf{v}_{t-1} + \beta_2\mathbf{g}_t \odot \mathbf{g}_t$$

▷ dims $r \times n$

$$\delta_t = \frac{\gamma}{\sqrt{\mathbf{v}_t} + \epsilon} \odot \mathbf{m}_t$$

▷ dims $r \times n$

Simplified as

$$\mathbf{X}_{t+1} = \mathbf{X}_t + \mathbf{P}_t \rho(\mathbf{P}_t^\top \mathbf{G}_t)$$

- Parameter updates:

$$\mathbf{X}_{t+1} = \mathbf{X}_t - \mathbf{P}_t \delta_t$$


▷ dims $m \times n$

- Memory cost: Model \mathbf{X} , Gradient \mathbf{G} , Projection \mathbf{P} , OptStates \mathbf{m}, \mathbf{v} and activations


trivial memory cost

GaLore: Gradient Low-Rank Projection


LLaMA 7B



Adam



Adam + GaLore




Adam + GaLore + Layerwise
gradient dropping (LWGD)



RTX 4090 affordable !!


GaLore: Gradient Low-Rank Projection

GPT3 175B



Adam

A100 * 21



Adam + GaLore + LWGD


A100 * 8

GaLore: Gradient Low-Rank Projection

- Recall the GaLore update: $\mathbf{X}_{t+1} = \mathbf{X}_t + \mathbf{P}_t \rho(\mathbf{P}_t^\top \mathbf{G}_t)$
- How to find the projection matrix? **SVD decomposition!**

$$\mathbf{G}_t = \mathbf{U} \Sigma \mathbf{V}^\top \longrightarrow \mathbf{P}_t = \boxed{\mathbf{U}[:, : r]} \in \mathbb{R}^{m \times r}$$

|
Select the first r columns



[Subspace Optimization for Large Language Models with Convergence Guarantees]


GaLore: Gradient Low-Rank Projection

- It is computationally expensive to perform SVD in each iteration
- **Lazy SVD**: perform SVD every τ iterations

(The complete GaLore algorithm)

$$\begin{cases} P_t \leftarrow \text{SVD}(G_t) & \text{if } t \bmod \tau = 0 \\ P_t \leftarrow P_{t-1} & \text{otherwise} \end{cases}$$

$$X_{t+1} = X_t + P_t \rho(P_t^\top G_t)$$



[GaLore: Memory-Efficient LLM Training by Gradient Low-Rank Projection]

GaLore: Gradient Low-Rank Projection

Pretraining LLaMA on C4 dataset

	60M	130M	350M	1B
Full-Rank	34.06 (0.36G)	25.08 (0.76G)	18.80 (2.06G)	15.56 (7.80G)
GaLore	34.88 (0.24G)	25.36 (0.52G)	18.95 (1.22G)	15.64 (4.38G)
Low-Rank	78.18 (0.26G)	45.51 (0.54G)	37.41 (1.08G)	142.53 (3.57G)
LoRA	34.99 (0.36G)	33.92 (0.80G)	25.58 (1.76G)	19.21 (6.17G)
ReLoRA	37.04 (0.36G)	29.37 (0.80G)	29.08 (1.76G)	18.33 (6.17G)
r/d_{model}	128 / 256	256 / 768	256 / 1024	512 / 2048
Training Tokens	1.1B	2.2B	6.4B	13.1B

GaLore: Gradient Low-Rank Projection

Fine-tuning RoBERTa-Base on GLUE

	Memory	CoLA	STS-B	MRPC	RTE	SST2	MNLI	QNLI	QQP	Avg
Full Fine-Tuning	747M	62.24	90.92	91.30	79.42	94.57	87.18	92.33	92.28	86.28
GaLore (rank=4)	253M	60.35	90.73	92.25	79.42	94.04	87.00	92.24	91.06	85.89
LoRA (rank=4)	257M	61.38	90.57	91.07	78.70	92.89	86.82	92.18	91.29	85.61
GaLore (rank=8)	257M	60.06	90.82	92.01	79.78	94.38	87.17	92.20	91.11	85.94
LoRA (rank=8)	264M	61.83	90.80	91.90	79.06	93.46	86.94	92.25	91.22	85.93

GaLore looks great

But does GaLore provably converge to the desired solution?

Not Necessarily True!

Y. He, P. Li, Y. Hu, C. Chen, and K. Yuan, *Subspace Optimization for Large Language Models with Convergence Guarantees*, 2024




PART 04

Non-convergence and convergence in GaLore



Intuition behind GaLore's non-convergence




When gradient noise dominates the stochastic gradient, SVD captures **noise-dominated** subspace!

All gradient information is lost !

Is non-convergence common? Yes!

Gradient dominates
during the initial stages

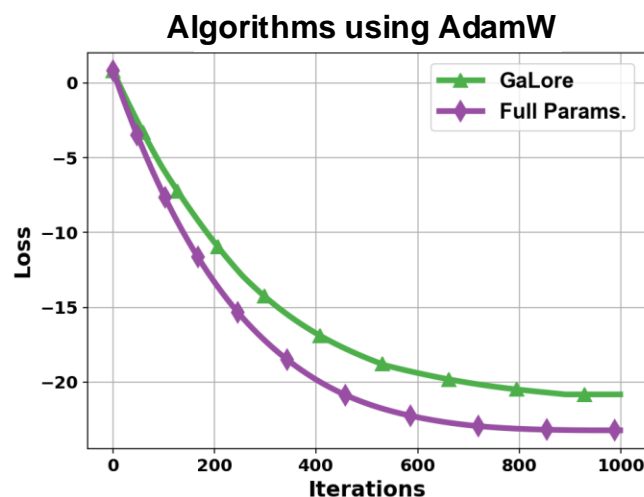
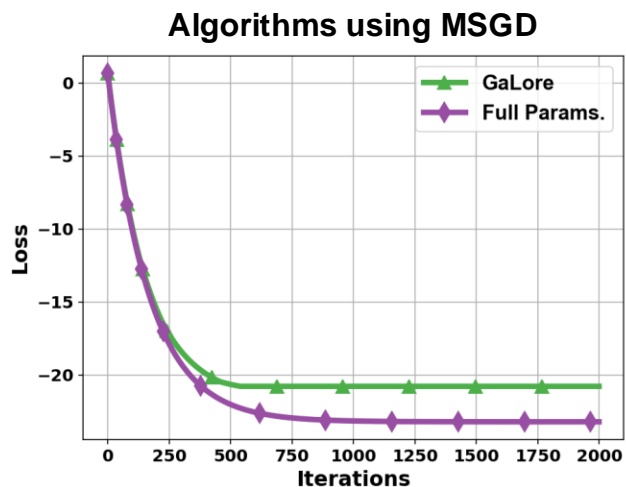


Noise dominates when
approaching the local minimum

Counter-Example. We consider the following quadratic problem with gradient noise:

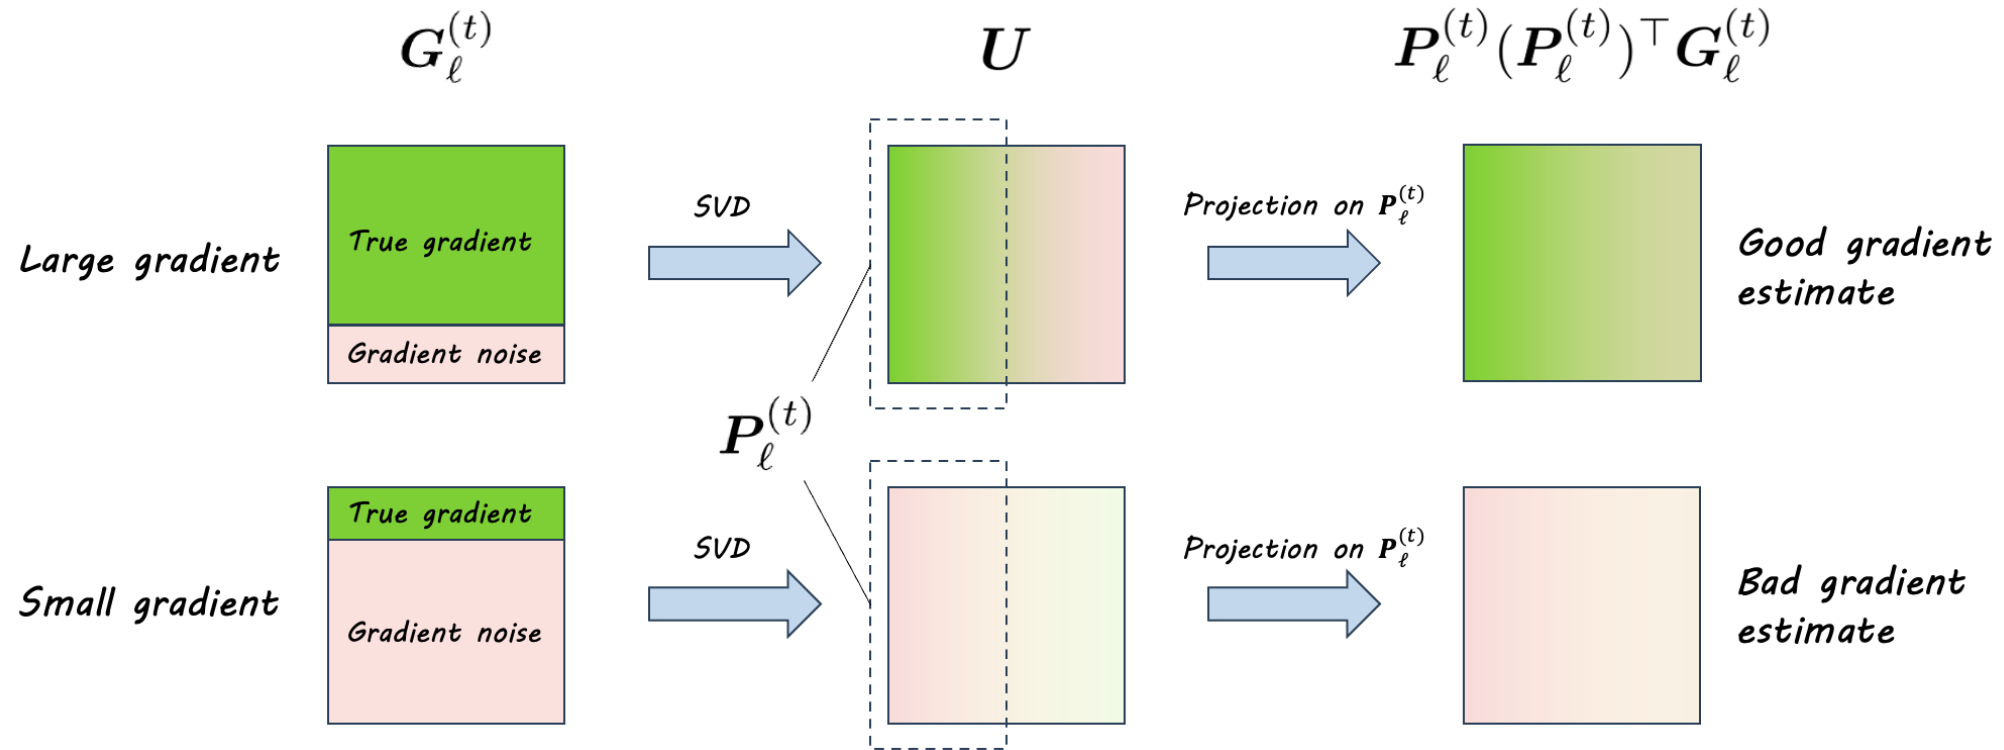
$$f(\mathbf{X}) = \frac{1}{2} \|\mathbf{A}\mathbf{X}\|_F^2 + \langle \mathbf{B}, \mathbf{X} \rangle_F, \quad \nabla F(\mathbf{X}; \xi) = \nabla f(\mathbf{X}) + \xi \sigma \mathbf{C}, \quad (1)$$

where $\mathbf{A} = (\mathbf{I}_{n-r} \ 0) \in \mathbb{R}^{(n-r) \times n}$, $\mathbf{B} = \begin{pmatrix} \mathbf{D} & 0 \\ 0 & 0 \end{pmatrix} \in \mathbb{R}^{n \times n}$ with $\mathbf{D} \in \mathbb{R}^{(n-r) \times (n-r)}$ generated randomly, $\mathbf{C} = \begin{pmatrix} 0 & 0 \\ 0 & \mathbf{I}_r \end{pmatrix} \in \mathbb{R}^{n \times n}$, ξ is a random variable uniformly sampled from $\{1, -1\}$ per iteration, and σ is used to control the gradient noise.



GaLore does **NOT** converge to desired solutions

Under what conditions can GaLore converge?



GaLore can converge if we can avoid the noise-dominant scenarios

- Consider GaLore with deterministic gradient: $G_\ell^{(t)} = \nabla_\ell f(\mathbf{x}^{(t)})$

Theorem 2 (Convergence rate of deterministic GaLore). *Under Assumptions 1-2, if the number of iterations $T \geq 64/(3\underline{\delta})$ and we choose*

$$\beta_1 = 1, \quad \tau = \left\lceil \frac{64}{3\underline{\delta}\beta_1} \right\rceil, \quad \text{and} \quad \eta = \left(4L + \sqrt{\frac{80L^2}{3\underline{\delta}\beta_1^2}} + \sqrt{\frac{80\tau^2 L^2}{3\underline{\delta}}} + \sqrt{\frac{16\tau L^2}{3\beta_1}} \right)^{-1},$$

GaLore using deterministic gradients and MSGD with MP converges as

$$\frac{1}{T} \sum_{t=0}^{T-1} \mathbb{E}[\|\nabla f(\mathbf{x}^{(t)})\|_2^2] = \mathcal{O}\left(\frac{L\Delta}{\underline{\delta}^{5/2}T}\right),$$

where $\Delta = f(\mathbf{x}^{(0)}) - \inf_{\mathbf{x}} f(\mathbf{x})$ and $\underline{\delta} := \min_\ell \frac{r_\ell}{\min\{m_\ell, n_\ell\}}$.

- Noise-free GaLore **converges** at rate $\mathcal{O}(1/T)$.

Condition II: Large batch-size

- Consider GaLore with large-batch stochastic gradient: $G_\ell^{(t)} = \frac{1}{\mathcal{B}} \sum_{b=1}^{\mathcal{B}} \nabla_\ell F(\mathbf{x}^{(t)}; \xi^{(t,b)})$
- Batch-size \mathcal{B} increases with iteration T , e.g., $\mathcal{B} = \mathcal{O}(\sqrt{T})$

Theorem 3 (Convergence rate of large-batch GaLore). *Under Assumptions 1-3, if $T \geq 2 + 128/(3\underline{\delta}) + (128\sigma)^2/(9\sqrt{\underline{\delta}}L\Delta)$ and we choose $\tau = \lceil 64/(3\underline{\delta}\beta_1) \rceil$, $\mathcal{B} = \lceil 1/(\underline{\delta}\beta_1) \rceil$,*

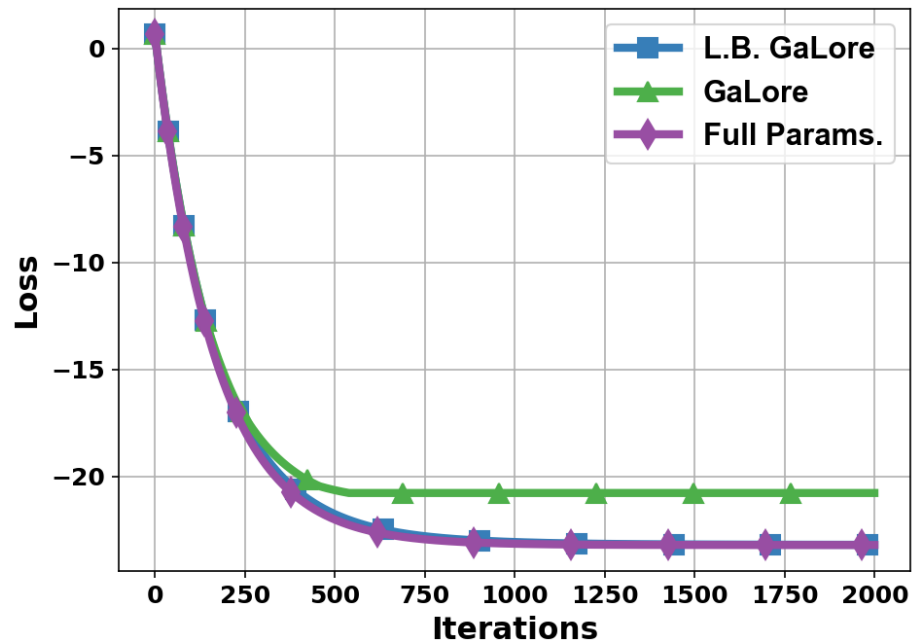
$$\beta_1 = \left(1 + \sqrt{\frac{\underline{\delta}^{3/2}\sigma^2 T}{L\Delta}} \right)^{-1}, \quad \text{and} \quad \eta = \left(4L + \sqrt{\frac{80L^2}{3\underline{\delta}\beta_1^2}} + \sqrt{\frac{40\tau^2 L^2}{\underline{\delta}}} + \sqrt{\frac{32\tau L^2}{3\beta_1}} \right)^{-1},$$

GaLore using large-batch gradients and MSGD with MP converges as

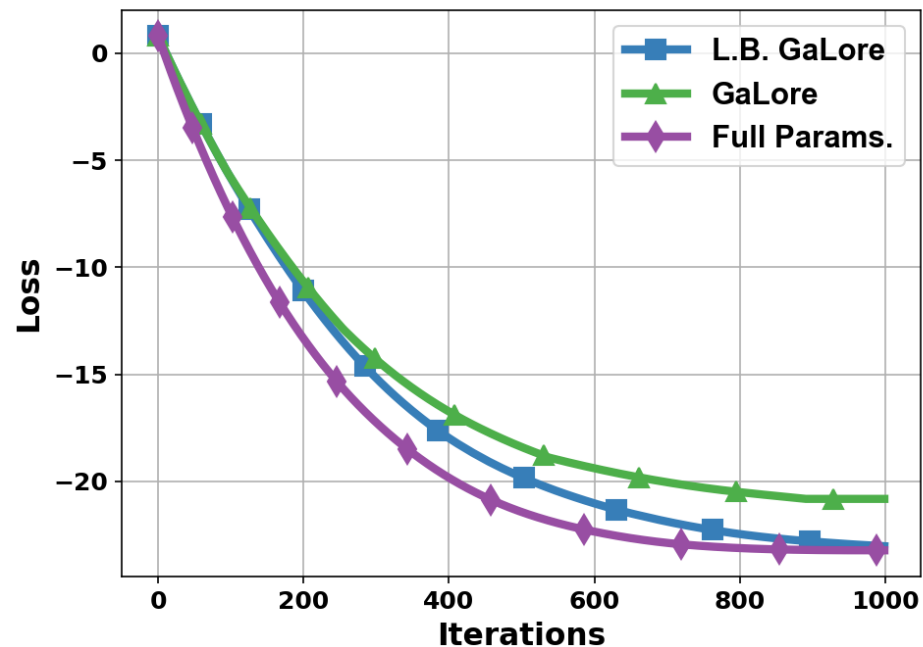
$$\frac{1}{T} \sum_{t=0}^{T-1} \mathbb{E}[\|\nabla f(\mathbf{x}^{(t)})\|_2^2] = \mathcal{O} \left(\frac{L\Delta}{\underline{\delta}^{5/2}T} + \sqrt{\frac{L\Delta\sigma^2}{\underline{\delta}^{7/2}T}} \right),$$

where $\Delta = f(\mathbf{x}^{(0)}) - \inf_{\mathbf{x}} f(\mathbf{x})$ and $\underline{\delta} := \min_\ell \frac{r_\ell}{\min\{m_\ell, n_\ell\}}$.

Algorithms using MSGD



Algorithms using AdamW



However, neither noise-free nor large-batch is practical for LLMs settings

PART 05

GoLore: Gradient random Low-rank projection



- In LLM settings, gradient noise exists and batch-size does not increase with iterations
- The root reason that GaLore has convergence issues is the SVD-incurred subspace
- Random projection can possibly capture gradient information when noise dominates

(Stiefel manifold) $\text{St}_{m,r} = \{\mathbf{P} \in \mathbb{R}^{m \times r} \mid \mathbf{P}^\top \mathbf{P} = \mathbf{I}_r\}.$

Proposition 1 (Chikuse (2012), Theorem 2.2.1). *A random matrix \mathbf{X} uniformly distributed on $\text{St}_{m,r}$ is expressed as $\mathbf{X} = \mathbf{Z}(\mathbf{Z}^\top \mathbf{Z})^{-1/2}$, where the elements of an $m \times r$ random matrix \mathbf{Z} are independent and identically distributed as normal $\mathcal{N}(0, 1)$.*

Following Prop. 1, we can sample random projections from Stiefel manifold

Stiefel random projection injects contractive error

- Instead of SVD, GoLore samples the projection matrix uniformly on the Stiefel manifold:

$$P_t \sim \mathcal{U}(\text{St}_{m,r})$$

- The following Lemma illustrates the projection error in GoLore:

Lemma 5 (Error of GoLore's projection). *Let $P \sim \mathcal{U}(\text{St}_{m,r})$, $Q \sim \mathcal{U}(\text{St}_{n,r})$, it holds for all $G \in \mathbb{R}^{m \times n}$ that*

$$\mathbb{E}[PP^\top] = \frac{r}{m} \cdot I, \quad \mathbb{E}[QQ^\top] = \frac{r}{n} \cdot I,$$

and

$$\mathbb{E}[\|PP^\top G - G\|_F^2] = \left(1 - \frac{r}{m}\right) \|G\|_F^2, \quad \mathbb{E}[\|GQQ^\top - G\|_F^2] = \left(1 - \frac{r}{n}\right) \|G\|_F^2.$$

Theorem 4 (Convergence rate of GoLore). *Under Assumptions 1-3, for any $T \geq 2 + 128/(3\underline{\delta}) + (128\sigma)^2/(9\sqrt{\underline{\delta}}L\Delta)$, if we choose $\tau = \lceil 64/(3\underline{\delta}\beta_1) \rceil$,*

$$\beta_1 = \left(1 + \sqrt{\frac{\underline{\delta}^{3/2}\sigma^2 T}{L\Delta}}\right)^{-1}, \quad \text{and} \quad \eta = \left(4L + \sqrt{\frac{80L^2}{3\underline{\delta}\beta_1^2}} + \sqrt{\frac{80\tau^2 L^2}{3\underline{\delta}}} + \sqrt{\frac{16\tau L^2}{3\beta_1}}\right)^{-1},$$

GoLore using small-batch stochastic gradients and MSGD with MP converges as

$$\frac{1}{T} \sum_{t=0}^{T-1} \mathbb{E}[\|\nabla f(\mathbf{x}^{(t)})\|_2^2] = \mathcal{O}\left(\frac{L\Delta}{\underline{\delta}^{5/2}T} + \sqrt{\frac{L\Delta\sigma^2}{\underline{\delta}^{7/2}T}}\right),$$

where $\Delta = f(\mathbf{x}^{(0)}) - \inf_{\mathbf{x}} f(\mathbf{x})$ and $\underline{\delta} := \min_{\ell} \frac{r_{\ell}}{\min\{m_{\ell}, n_{\ell}\}}$.

- Theoretically, GoLore **converges** at rate $\mathcal{O}(1/\sqrt{T})$.

Convergence analysis for subspace GD

- The minimization problem $\min_{X \in \mathbb{R}^{m \times n}} f(X)$
- Subspace gradient descent: $X^t = X^{t-1} - \gamma P^t (P^t)^\top \nabla f(X^{t-1}), \quad \forall t = 1, \dots, T$
- For simplicity, we assume there is no lazy update on the projection matrix P
- Assumptions on projection matrix P : $(P^t)^\top P^t = I_r$ and $\mathbb{E}[P^t (P^t)^\top] = \frac{r}{m} I_m$

$$\begin{aligned} f(X^t) &\leq f(X^{t-1}) + \langle \nabla f(X^{t-1}), X^t - X^{t-1} \rangle + \frac{L}{2} \|X^t - X^{t-1}\|^2 \\ &\leq f(X^{t-1}) - \gamma \langle \nabla f(X^{t-1}), P^t (P^t)^\top \nabla f(X^{t-1}) \rangle + \frac{\gamma^2 L}{2} \|P^t (P^t)^\top \nabla f(X^{t-1})\|^2 \\ &= f(X^{t-1}) - \gamma \|(P^t)^\top \nabla f(X^{t-1})\|^2 + \frac{\gamma^2 L}{2} \|(P^t)^\top \nabla f(X^{t-1})\|^2 \\ &= f(X^{t-1}) - \gamma \left(1 - \frac{\gamma L}{2}\right) \|(P^t)^\top \nabla f(X^{t-1})\|^2 \end{aligned}$$

Convergence analysis for subspace GD

- Taking expectation on the random projection matrix P , we have

$$\begin{aligned}\mathbb{E}f(X^t) &\leq \mathbb{E}f(X^{t-1}) - \gamma(1 - \frac{\gamma L}{2})\frac{r}{m}\mathbb{E}\|\nabla f(X^{t-1})\|^2 \\ &\leq \mathbb{E}f(X^{t-1}) - \frac{\gamma r}{2m}\mathbb{E}\|\nabla f(X^{t-1})\|^2\end{aligned}$$

- From the above inequality, we achieve

$$\|\nabla f(X^{t-1})\|^2 \leq \frac{2m(\mathbb{E}f(X^{t-1}) - \mathbb{E}f(X^t))}{\gamma r}$$

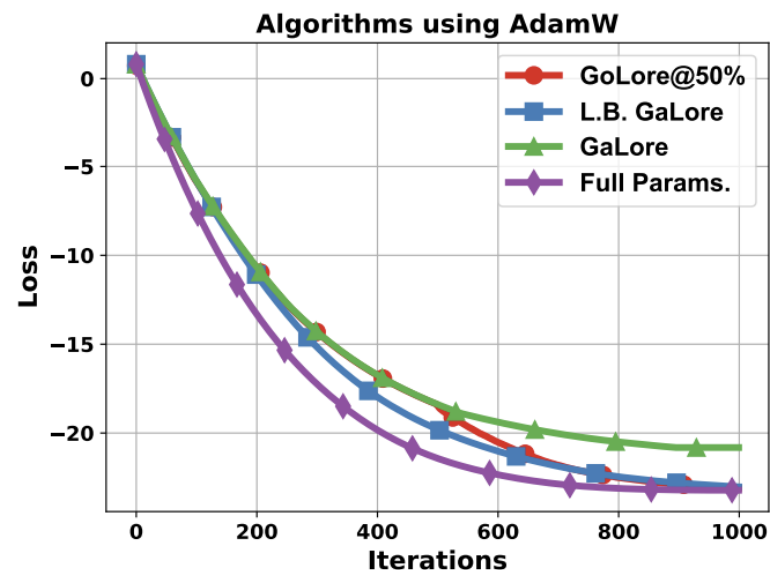
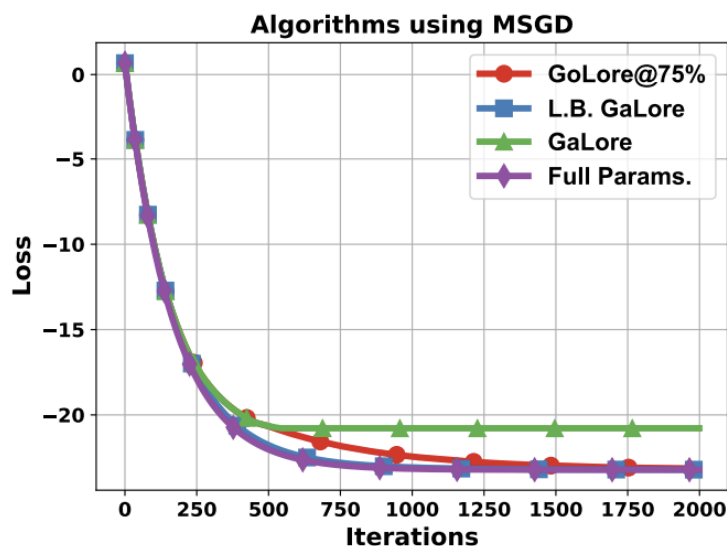
and therefore

$$\frac{1}{T+1} \sum_{t=0}^T \|\nabla f(X^t)\|^2 \leq \frac{2mL(f(X^0) - f^*)}{r(T+1)}$$

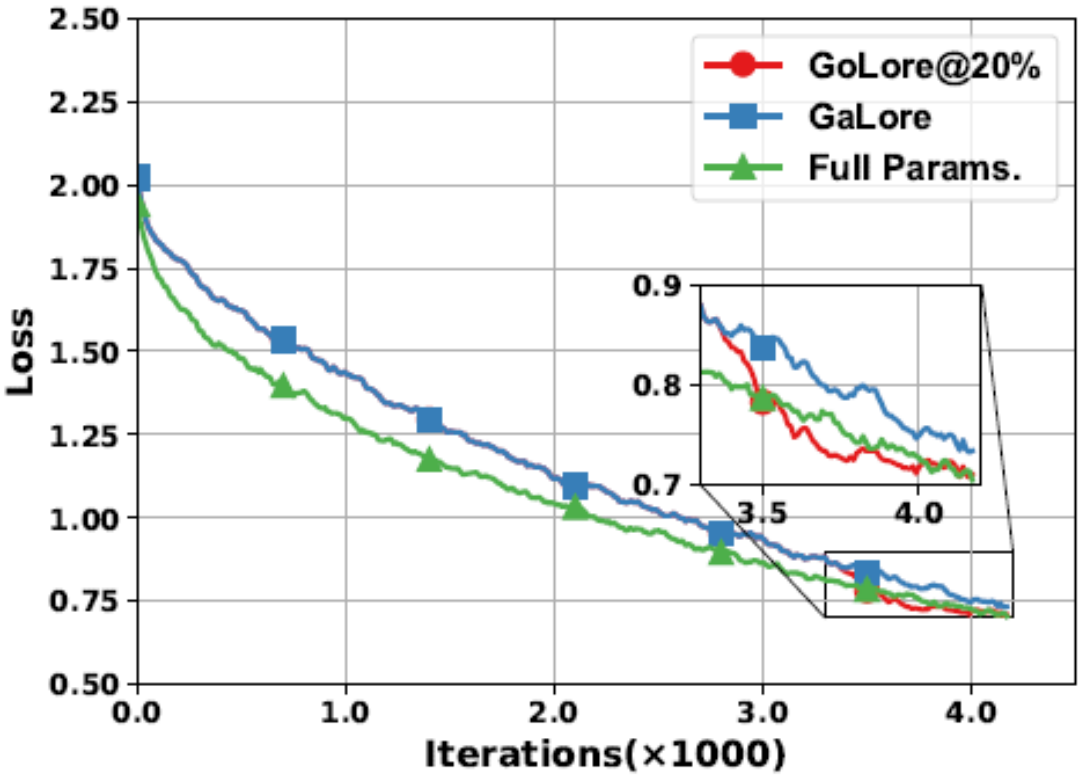
A hybrid strategy: GaLore + GoLore

- SVD projection is preferred in initial stages: effectively capture gradient information
- Random projection is preferred when approaching solutions: avoid losing gradient information

GoLore@x% = GaLore (first (100-x)% iters) + GoLore (last x% iters)



Fine-tuning LLaMA2-7B on WinoGrande:



Experimental results

- Fine-tuning RoBERTa-BASE on GLUE benchmark:

Algorithm	CoLA	STS-B	MRPC	RTE	SST2	MNLI	QNLI	QQP	Avg
Full Params.	62.07	90.18	92.25	78.34	94.38	87.59	92.46	91.90	86.15
GaLore	61.32	90.24	92.55	77.62	94.61	86.92	92.06	90.84	85.77
GoLore@20%	61.66	90.55	92.93	78.34	94.61	87.02	92.20	90.91	86.03

- GoLore shows **superior** performance than GaLore in the above experiments.

Improving the computational efficiency

- GaLore/GoLore always computes the full gradient before compressing them into subspaces.
- Can we compute the compressed gradient directly, without computing the full gradient?

Original Implementation

$$y = Wx$$

$$\nabla_W \mathcal{L} = (\nabla_y \mathcal{L})x^\top \quad \text{Backpropagated gradient}$$

$$W \leftarrow W + B\rho(B^\top(\nabla_W \mathcal{L}))$$

$$W = W_0 + BA$$
$$\longrightarrow$$

New Implementation

$$y = W_0x + BAx$$

$$\nabla_A \mathcal{L} = B^\top(\nabla_y \mathcal{L})x^\top$$

$$A \leftarrow A + \rho(\nabla_A \mathcal{L})$$



Improving the computational efficiency

- Comparing the computational complexities:

GaLore Implementation	Memory	Computation
(Zhao et al., 2024)	$mn + rm + rn + bm$	$6bmn + 4rmn + 2mn + 3rn$
Our ReLoRA-like version	$mn + rm + 2rn + bm + br$	$4bmn + 4brm + 6brn + 5rn$

- When $r \ll \min\{m, n\}$, this new version reduces the computation complexity from $(6b + 4r + 2)mn$ to $4bmn$, with minimal memory overhead.

Summary

- Gradient low-rank projection can effectively save optimizer states
- GaLore cannot converge to desired solutions due to SVD projections
- Random projections enable GaLore to converge

