



# Advanced LLM Optimizers

**Yihan Jin, Ming Sun, Jiahe Geng,  
Shuchen Zhu, Kun Yuan**

**Peking University**

## Recall: Memory Consumption in LLM Training

➤ Memory = **Model + Gradient + Optimizer States + Activation**

➤ Adam's Cost:

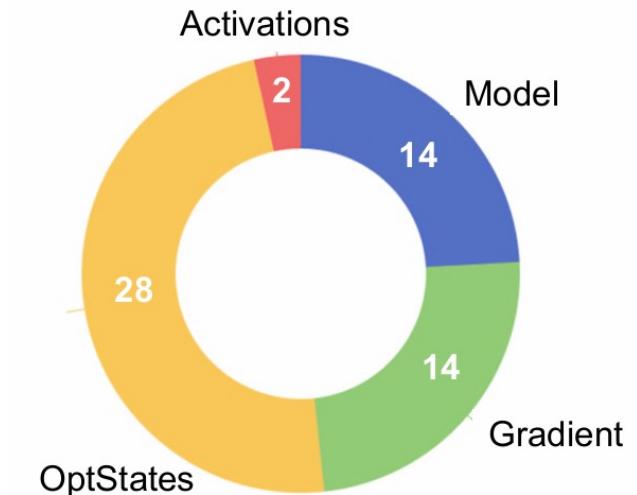
Model Parameter –  $\Phi$

Gradient -  $\Phi$

**Optimizer States** (First & Second moment) -  $2\Phi$

➤ Consequence:

The 7B-pretrained model(BF16) requires 28GB of Adam optimizer states.



**Goal: Decrease optimizer states while maintaining performance**

## Start with Adam: Destructing Adam's Redundancy

---

- Adam Update Rule:

$$W_{t+1} = W_t - \eta \frac{M_t}{\sqrt{V_t} + \epsilon}$$

where  $W_t \in R^{m \times n}$  is weight matrix,  $M_t \in R^{m \times n}$  is **momentum**,

$V_t \in R^{m \times n}$  is second moment as **adaptive learning rate (Preconditioning)**

- Question:

1. Is it necessary to maintain an **adaptive learning rate** for each components in  $W_t$  ?
2. If not, how should the **adaptability**( $V_t$ ) be arranged?

---

# Adafactor: Adaptive Learning Rates with Sublinear Memory Cost

---

Noam Shazeer<sup>1</sup> Mitchell Stern<sup>1,2</sup>

<sup>1</sup>Google Brain, Mountain View, California, USA <sup>2</sup>University of California, Berkeley, California, USA. Correspondence to: Noam Shazeer <noam@google.com>.

*Proceedings of the 35<sup>th</sup> International Conference on Machine Learning*, Stockholm, Sweden, PMLR 80, 2018. Copyright 2018 by the author(s).

## Adafactor: Factorize $V_t$

---

➤ Assumption:

$V_t \in \mathbb{R}^{m \times n}$  is **low-rank**, i.e.  $V_t \approx R_t \cdot C_t$ ,

where  $R_t \in \mathbb{R}^{m \times 1}$ ,  $C_t \in \mathbb{R}^{1 \times n}$

➤ Motivation:

Decrease the memory of optimizer states by **storing  $R_t$  &  $C_t$  instead of  $V_t$**

➤ Memory Saving Results:

$O(mn) \rightarrow O(m+n)$

# Adafactor: How to derive rank-1 factorization?

---



- Objective: Minimize **Generalized KL-Divergence (I-Divergence)**

$$D(V \parallel RC) = \sum_{i,j} \left( V_{ij} \log \frac{V_{ij}}{R_i C_j} - V_{ij} + R_i C_j \right)$$

- Constraints:

$$R_i \geq 0, C_j \geq 0$$

## Theorem

*The solution set of the optimization problem (minimizing  $I$ -divergence) when consists of all feasible pairs  $(R, S)$  satisfying:*

$$RC = \frac{V\mathbf{1}_m \mathbf{1}_n^\top V}{\mathbf{1}_n^\top V \mathbf{1}_m}$$

*where  $\mathbf{1}_\ell = (1, \dots, 1) \in \mathbb{R}^\ell$  denotes a column vector of  $\ell$  ones.*

➤ Solution (Closed-form):

$$R = V \cdot \mathbf{1}_m, \quad C = \frac{\mathbf{1}_n^\top \cdot V}{\mathbf{1}_n^\top \cdot V \cdot \mathbf{1}_m}$$

# Proof of Theorem

---

$$\begin{aligned}
 D(V||RC) &= \sum_{i=1}^n \sum_{j=1}^m \left( V_{ij} \log \frac{V_{ij}}{R_i C_j} - V_{ij} + R_i C_j \right) \\
 &= \sum_{i=1}^n \sum_{j=1}^m V_{ij} \log V_{ij} - \sum_{i=1}^n \sum_{j=1}^m V_{ij} \log R_i \\
 &\quad - \sum_{i=1}^n \sum_{j=1}^m V_{ij} \log C_j - \sum_{i=1}^n \sum_{j=1}^m V_{ij} + \sum_{i=1}^n \sum_{j=1}^m R_i C_j
 \end{aligned}$$

Setting the derivatives with respect to  $R_i$  and  $C_j$  to 0:

$$\frac{\partial \mathcal{L}}{\partial R_i} = - \sum_{j=1}^m \frac{V_{ij}}{R_i} + \sum_{j=1}^m C_j = 0 \quad \Rightarrow \quad R_i = \frac{\sum_{j=1}^m V_{ij}}{\sum_{j=1}^m C_j}$$

$$\frac{\partial \mathcal{L}}{\partial C_j} = - \sum_{i=1}^n \frac{V_{ij}}{C_j} + \sum_{i=1}^n R_i = 0 \quad \Rightarrow \quad C_j = \frac{\sum_{i=1}^n V_{ij}}{\sum_{i=1}^n R_i}$$

# Proof of Theorem

---

The solution has a scaling symmetry  $(\alpha R, C/\alpha)$ . We break this symmetry by enforcing the constraint  $\sum_i R_i = \sum_{i,j} V_{ij}$ .

This leads to the **canonical minimizer**:

$$R_i = \sum_{j=1}^m V_{ij}, \quad C_j = \frac{\sum_{i=1}^n V_{ij}}{\sum_{i=1}^n \sum_{j=1}^m V_{ij}}$$

In vector notation:

$$R = V\mathbf{1}_m, \quad C = \frac{\mathbf{1}_n^\top V}{\mathbf{1}_n^\top V\mathbf{1}_m}$$

Thus, the product  $RS$  is unique:

$$RC = V\mathbf{1}_m \left( \frac{\mathbf{1}_n^\top V}{\mathbf{1}_n^\top V\mathbf{1}_m} \right) = \frac{V\mathbf{1}_m \mathbf{1}_n^\top V}{\mathbf{1}_n^\top V\mathbf{1}_m} \quad \square$$

# Adafactor: Removing Momentum

---

- Motivation: Further save the memory of momentum –  $\Phi$
- Pseudo-code:

**Algorithm 2** Adam for a matrix parameter  $X$  with factored second moments and first moment decay parameter  $\beta_1 = 0$ .

```

1: Inputs: initial point  $X_0 \in \mathbb{R}^{n \times m}$ , step sizes  $\{\alpha_t\}_{t=1}^T$ ,
   second moment decay  $\beta_2$ , regularization constant  $\epsilon$ 
2: Initialize  $R_0 = 0$  and  $C_0 = 0$ 
3: for  $t = 1$  to  $T$  do
4:    $G_t = \nabla f_t(X_{t-1})$ 
5:    $R_t = \beta_2 R_{t-1} + (1 - \beta_2)(G_t^2)1_m$ 
6:    $C_t = \beta_2 C_{t-1} + (1 - \beta_2)1_n^\top(G_t^2)$ 
7:    $\hat{V}_t = (R_t C_t / 1_n^\top R_t) / (1 - \beta_2^t)$ 
8:    $X_t = X_{t-1} - \alpha_t G_t / (\sqrt{\hat{V}_t} + \epsilon)$ 
9: end for

```

Rank-1 Factorization of  $V_t$

Discard Momentum

## Summary of Adafactor

---

➤ Pros:

Less memory usage of optimizer state:  $2mn (M_t + V_t) \rightarrow m+n (R_t + C_t)$

➤ Cons:

**Approximation Error:**  $V_t$  is not always of rank 1  $\rightarrow$  Slow convergence

**Throughput Cost:** Factorize  $V_t$  and compute RMS increase **computation cost**

➤ Result:

In order to achieve better performance, momentum  $M_t$  is **re-introduced** to the actual use of Adafactor, sacrificing memory to gain faster convergence speed.

# ADAM-MINI: USE FEWER LEARNING RATES TO GAIN MORE

**Yushun Zhang<sup>\*13</sup>, Congliang Chen<sup>\*13</sup>, Ziniu Li<sup>13</sup>, Tian Ding<sup>23</sup>, Chenwei Wu<sup>4</sup>,  
Diederik P. Kingma<sup>5</sup>, Yinyu Ye<sup>16</sup>, Zhi-Quan Luo<sup>13</sup>, Ruoyu Sun<sup>†123</sup>**

<sup>1</sup> The Chinese University of Hong Kong, Shenzhen, China;

<sup>2</sup> Shenzhen International Center for Industrial and Applied Mathematics;

<sup>3</sup> Shenzhen Research Institute of Big Data; <sup>4</sup> Duke University;

<sup>5</sup> Anthropic; <sup>6</sup> Stanford University

## Adam-mini: Start from Pre-conditioning Perspective

### ➤ Critique of Adafactor:

It assumes  $V_t$  is rank-1. The assumption lacks **observation** to support.  
→ Need to investigate NN's **structure!**

### ➤ View Adam as Pre-conditioning method:

$$w_{t+1} = w_t - \eta_t D_t m_t$$

where  $D_t = \text{Diag}\left(\frac{1}{\sqrt{v_t}}\right)$  is the pre-conditioning matrix

The ideal pre-conditioning matrix is Hessian's inverse  $H^{-1}$

**Need to inspect NN's Hessian structure!**

# Adam-mini: Hessian is near-block-diagonal

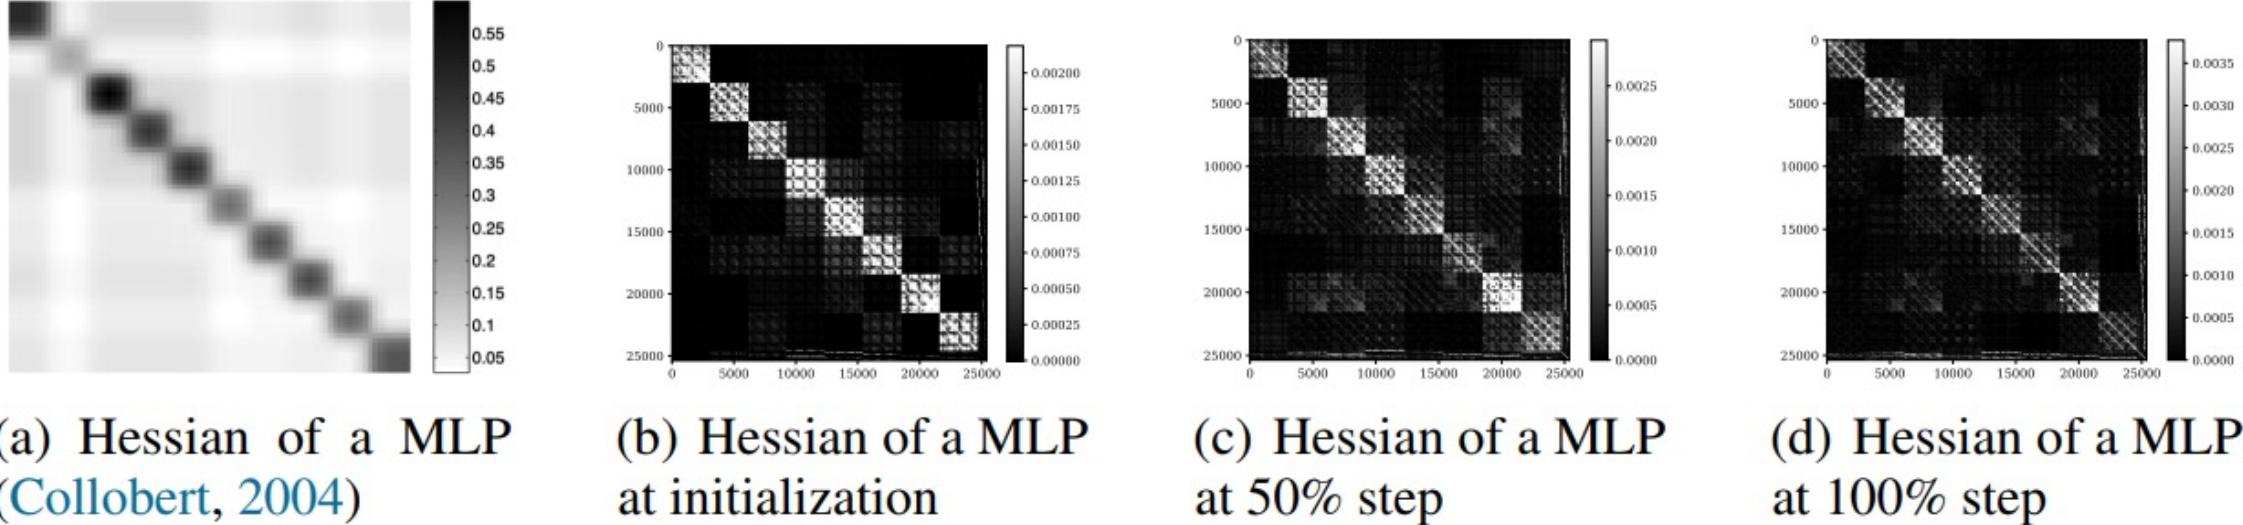
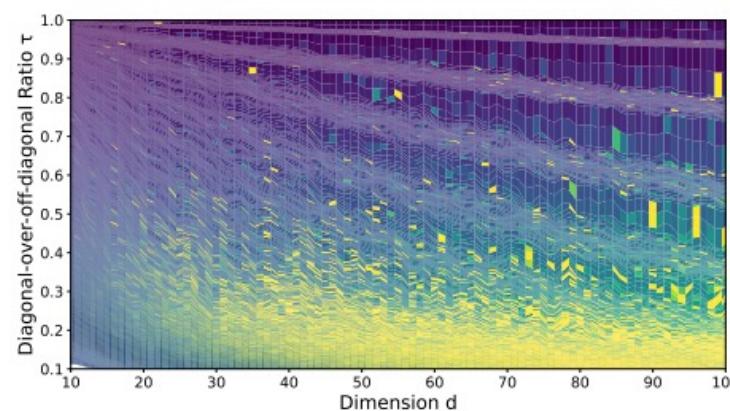


Figure 3: The near-block-diagonal Hessian structure of neural nets. (a) is the Hessian of an MLP after 1 training step reported in (Collobert, 2004). (b,c,d): the Hessians of a 1-hidden-layer MLP on CIFAR-100. The near-block-diagonal structure maintains throughout training, where each block corresponds to one neuron.

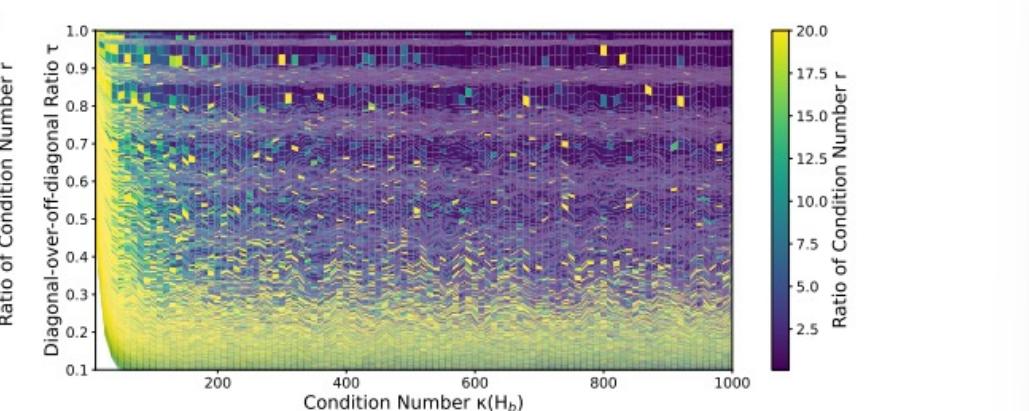
# Adam-mini: How is Adam's pre-conditioning effect?

$\tau = \frac{\sum_i |H_{b,i,i}|}{\sum_{i,j} |H_{b,i,j}|}$ : the degree of dominance of diagonal elements

$r = \frac{\kappa(D_{Adam} H_b)}{\kappa(H_b)}$ : the pre-conditioning effect of Adam



(a)  $r$  v.s. dimension  $d$



(b)  $r$  v.s. dimension  $\kappa(H_b)$

Figure 5: The effectiveness of Adam's preconditioner  $D_{Adam}$  on different matrix structures of  $H_b$ . (a): for most dimension  $d$ ,  $r$  is large when  $\tau$  is small ( $r$  and  $\tau$  are defined in Eq. (2)). This indicates that Adam might not be so effective when  $H_b$  is dense. We fix  $\kappa(H_b) = 500$  here. (b): We use the same setups as (a), except that we fix the dimension  $d = 50$  and change the  $x$ -axis to  $\kappa(H_b)$ .

# Adam-mini: Case study of Adam's pre-conditioning effect

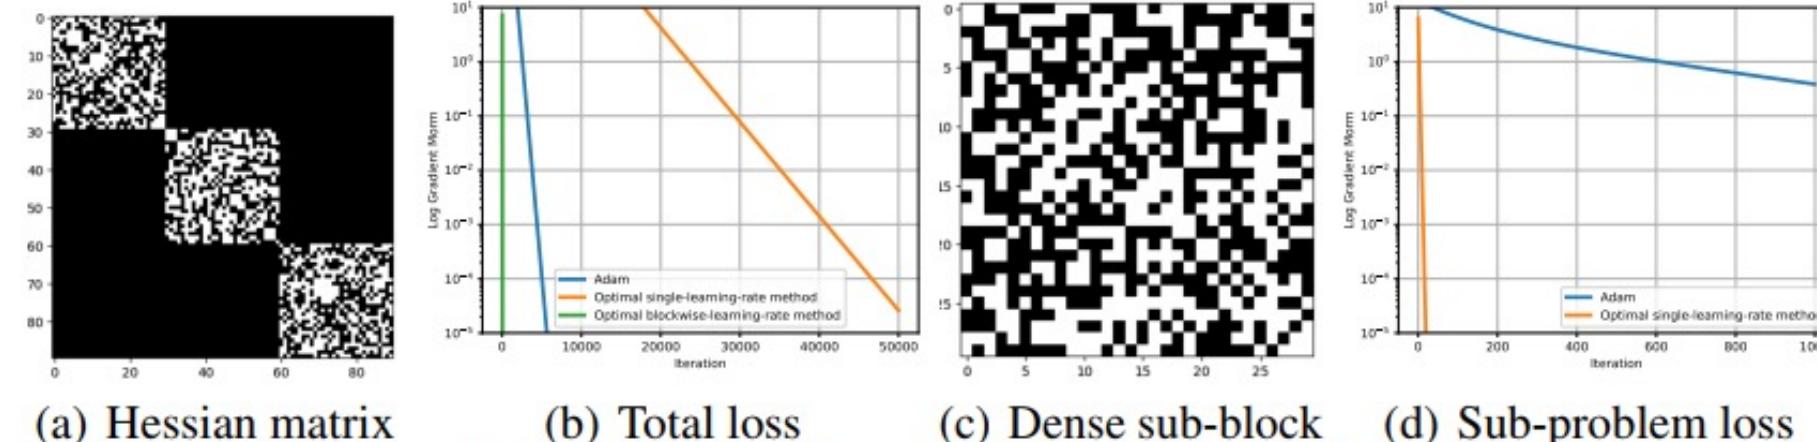


Figure 4: (a): The Hessian of a three-block random quadratic problem. (b): Training curves for the problem associated with the full Hessian in (a). The optimal single (blockwise) learning rate is chosen based on the full (blockwise) Hessian in (a). (c): The 1st dense Hessian sub-blocks in (a). (d): Training curves for the new problem associated with the Hessian in (c).

Conclusion:

- For **dense Hessian case**, Adam is far inferior to **optimal single-learning-rate**.
- For block-diagonal Hessian case, Adam surpasses optimal single-learning-rate

➤ (Recall) Question:

1. Is it necessary to maintain an adaptive learning rate for each components in  $W_t$ ?
2. If not, how should the adaptability( $V_t$ ) be arranged?

➤ For Q1:

Under **near-block-diagonal Hessian structure**, Adam's maintaining an adaptive learning rate( $V_t$ ) for each components in  $W_t$  involves **redundancy**.

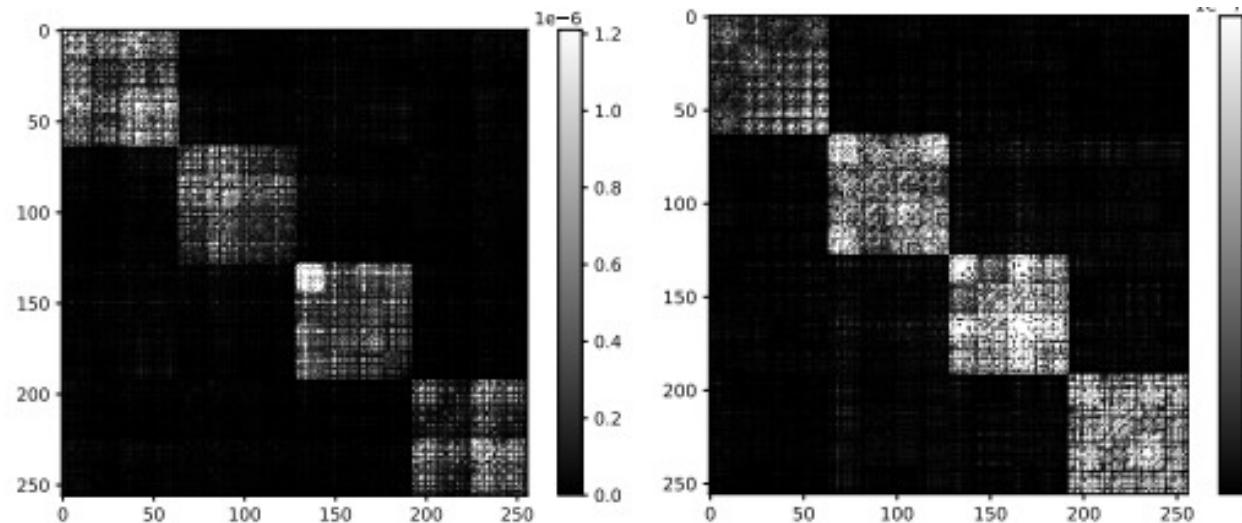
➤ For Q2:

For **each dense sub-block of Hessian**, carefully chose **single learning rate** is good enough.

Use Hessian information to partition variables into groups:

➤ **Query/Key: Head-wise**

Weight components in the same head as a block.



(a) query (4 heads)

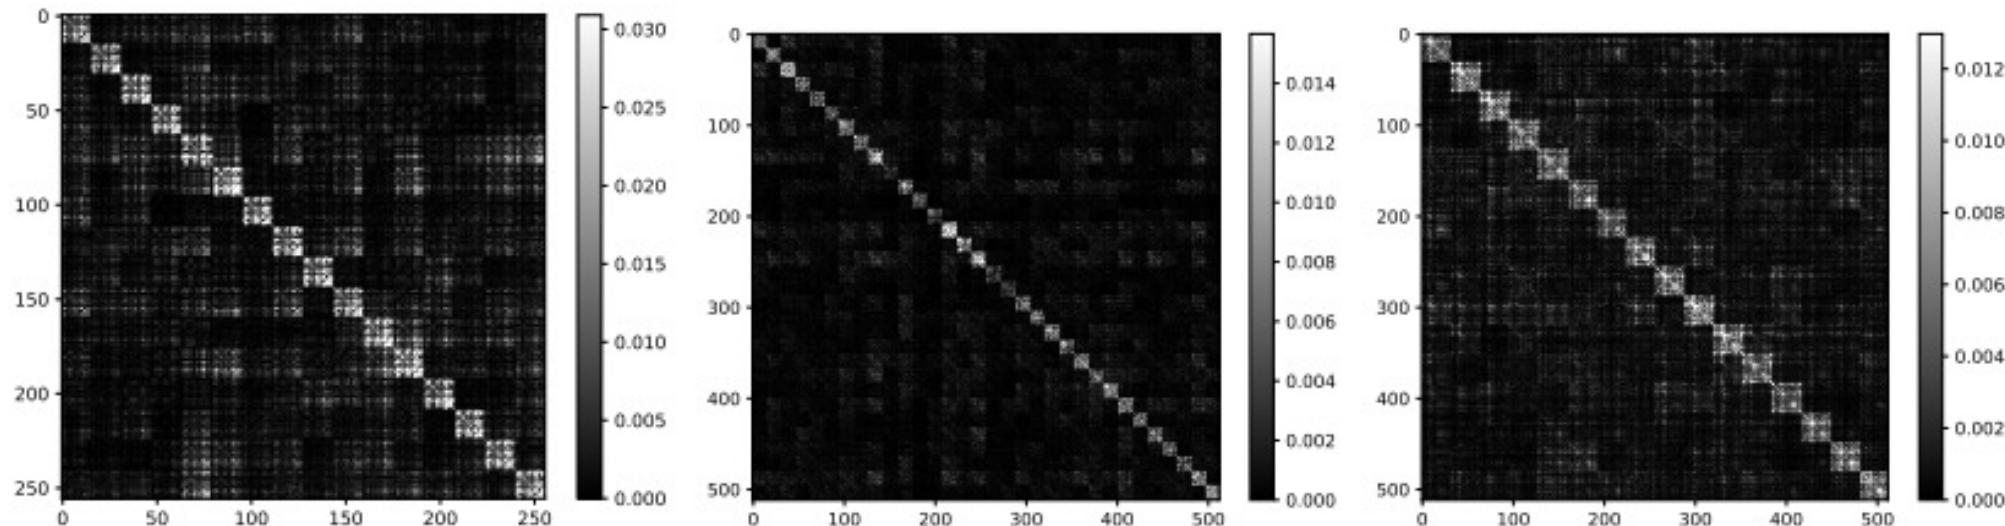
(b) key (4 heads)

# Adam-mini: Hessian based Transformer Partition Strategy

Use Hessian information to partition variables into groups:

- **attn.proj/MLP: Neuron-wise**

Weight components in the same row as a block



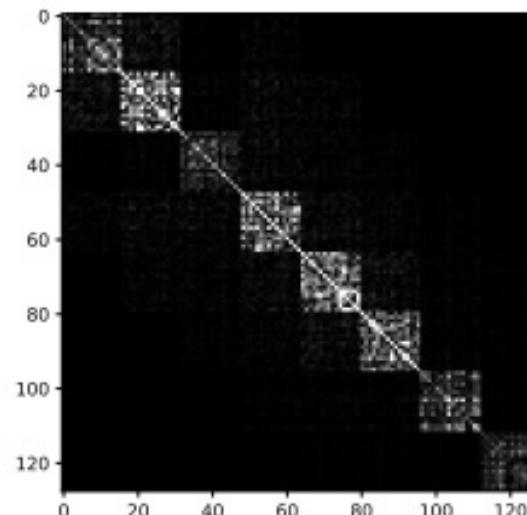
(d) attn.proj (16 neurons) (e) mlp.fc\_1 (32 neurons) (f) mlp.proj (16 neurons)

# Adam-mini: Hessian based Transformer Partition Strategy

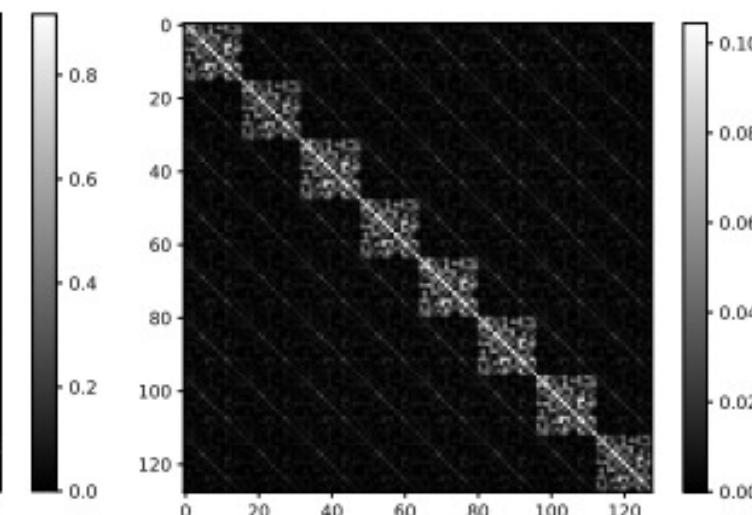
Use Hessian information to partition variables into groups:

- **Embedding/Output: Token-wise**

Embedding – Each word vector as a block



(g) embed (8 tokens)



(h) output (8 tokens)

# Adam-mini: How to set learning rate?

- For Adam:  $u_{\text{Adam}} = \left( \frac{\eta}{\sqrt{v_1}}, \frac{\eta}{\sqrt{v_2}}, \frac{\eta}{\sqrt{v_3}}, \frac{\eta}{\sqrt{v_4}}, \frac{\eta}{\sqrt{v_5}} \right)$ .
- For Adam-mini: suppose the partition is  $(1, 2, 3)$  and  $(4, 5)$  then

$$u_{\text{mini}} = \left( \frac{\eta}{\sqrt{(v_1+v_2+v_3)/3}}, \frac{\eta}{\sqrt{(v_1+v_2+v_3)/3}}, \frac{\eta}{\sqrt{(v_1+v_2+v_3)/3}}, \frac{\eta}{\sqrt{(v_4+v_5)/2}}, \frac{\eta}{\sqrt{(v_4+v_5)/2}} \right).$$

- Why  $\text{lr} = \text{mean}(g \odot g)$  ?
  1. Convenience
  2. Best among common statistics

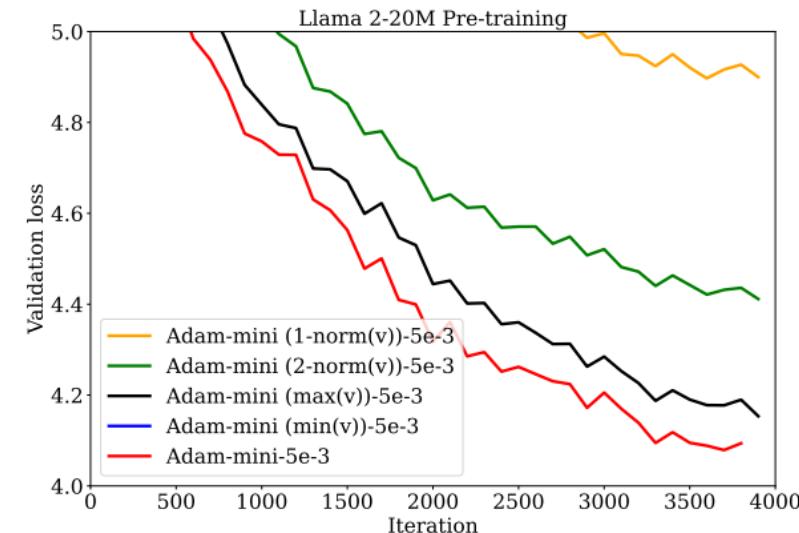


Figure 15: Ablation studies on the design of Adam-mini. We find that  $\text{mean}(v)$  performs better than other candidates. The blue curve does not show because the algorithm diverges and the curve is out of range.

## Algorithm 1 Adam-mini (General form)

```
1: Input weight-decay coefficient  $\lambda$  and  
   current step  $t$   
2: Partition params into param_blocks  
   by Principle 1 in Section 2.3  
3: for param in param_blocks do  
4:   g = param.grad  
5:   param = param -  $\eta_t * \lambda * \text{param}$   
6:   m =  $(1 - \beta_1) * g + \beta_1 * m$   
7:    $\hat{m} = \frac{m}{1 - \beta_1^t}$   
8:   v =  $(1 - \beta_2) * \text{mean}(g \odot g) + \beta_2 * v$  → Single lr for a sub-block  
9:    $\hat{v} = \frac{v}{1 - \beta_2^t}$   
10:  param = param -  $\eta_t * \frac{\hat{m}}{\sqrt{\hat{v}} + \epsilon}$   
11: end for
```

Partition blocks  
based on Hessian



Single lr for a sub-block

# Evaluation: Scaling Law

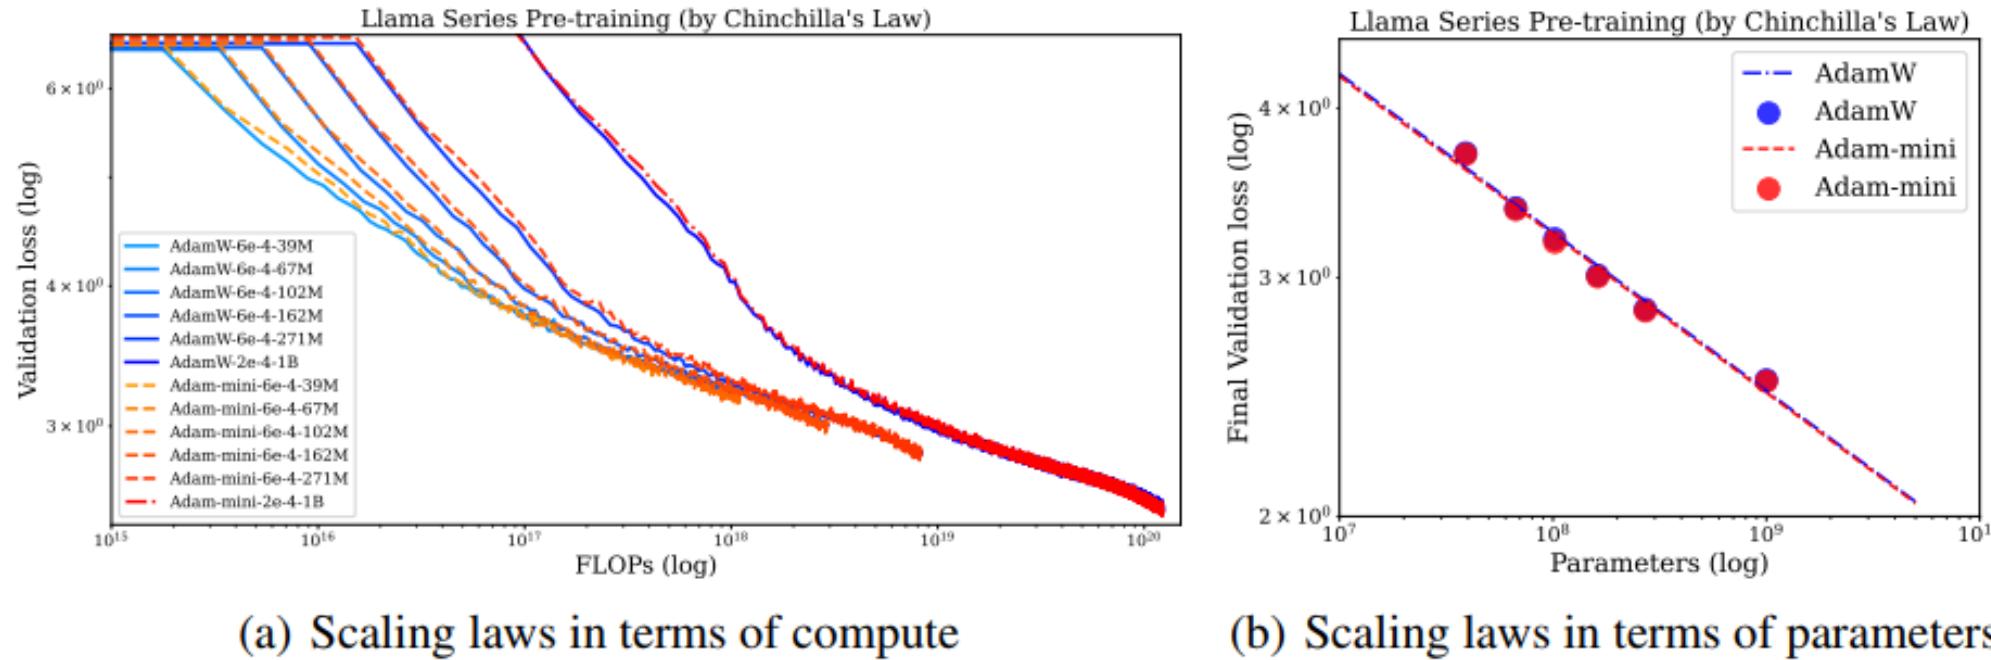


Figure 11: (a, b): Scaling laws of Adam-mini. We pre-train Llama 2 architectures by Chinchilla's law. For all models sized from 39M to 1B, Adam-mini's loss curves are consistently similar to AdamW, but Adam-mini uses 50% less memory. Further, as shown in (b), Adam-mini reaches a lower final loss than AdamW for all models. The fitted lines in (b) suggest that Adam-mini can be scaled up to larger models (if the scaling law holds).

**Adam-mini's loss curves closely resembles AdamW's**

**Adam-mini performs well using the same hyperparameter as AdamW**

# Evaluation: Memory & Throughput

Table 1: Memory cost of AdamW v.s. Adam-mini. Calculation is based on float32, which is a standard choice for optimizer states.

Model	Optimizer	Memory (GB)
GPT-2-1.5B	AdamW	12.48
GPT-2-1.5B	Adam-mini	6.24 (50% ↓)
Llama 2-1B	AdamW	8.80
Llama 2-1B	Adam-mini	4.40 (50% ↓)
Llama 2-7B	AdamW	53.92
Llama 2-7B	Adam-mini	26.96 (50% ↓)
Llama 3-8B	AdamW	64.24
Llama 3-8B	Adam-mini	32.12 (50% ↓)
Llama 2-13B	AdamW	104.16
Llama 2-13B	Adam-mini	52.08 (50% ↓)

Table 2: Throughput ( $\uparrow$ ) test on  $2 \times$  A800-80GB GPUs for Llama 2-7B pre-training.  $\times$  means out of memory. GPU hours ( $\downarrow$ ) to pre-train Llama 2-7B with the optimal token amount by Chinchila's law.

Optimizer	bs_per_GPU	total_bs	Throughput ( $\uparrow$ )
Adam-mini	4	256	5572.19 ( $\uparrow 49.6\%$ )
AdamW	2	256	$\times$
AdamW	1	256	3725.59
Optimizer	# Tokens (B)		GPU hours (h) ( $\downarrow$ )
AdamW	1		74.56
Adam-mini	1		49.85 ( $\downarrow 33.1\%$ )
AdamW	70		5219.16
Adam-mini	70		3489.55 ( $\downarrow 33.1\%$ )
AdamW	140		10438.32
Adam-mini	140		6979.10 ( $\downarrow 33.1\%$ )

Compared to AdamW, Adam-mini saves 50% memory, has 49.6% higher throughput.

Efficiency:

- **Memory:** Less memory usage of optimizer state:  $2mn (M_t + V_t) \rightarrow mn (M_t)$
- **Hyperparameter:** Performs well using the same hyperparameter as AdamW
- **Computation:** Substitute vector operations like *sqrt* & *div* by scalar operation

- Non-constrained Optimization Problem:

$$\min_{\mathbf{x} \in \mathbb{R}^n} f(\mathbf{x})$$

- Taylor's Formula:

$$f(\mathbf{x}_k + \mathbf{p}) \approx f(\mathbf{x}_k) + \nabla f(\mathbf{x}_k)^\top \mathbf{p} + \frac{1}{2} \mathbf{p}^\top \mathbf{H}(\mathbf{x}_k) \mathbf{p}$$

- Update Vector:

$$\mathbf{p}_k^{\text{Newton}} = -\mathbf{H}_k^{-1} \nabla f(\mathbf{x}_k)$$

## Second-order Methods

- Advantage of Second-order Methods: Curvature Calibration

Eigen Value Decomposition:

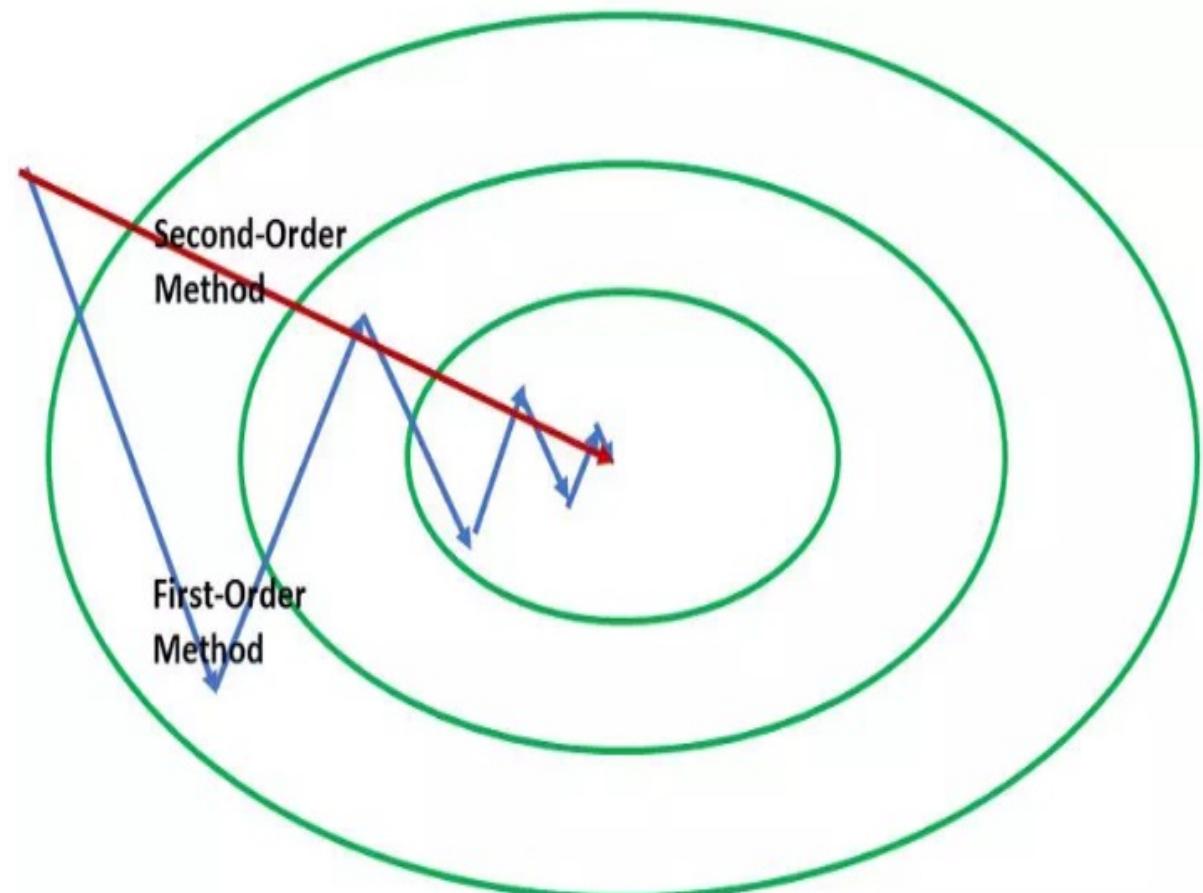
$$\mathbf{H}_k = \mathbf{Q} \boldsymbol{\Lambda} \mathbf{Q}^\top$$

Then we have:

$$\mathbf{p}_k^{\text{Newton}} = -\mathbf{Q} \boldsymbol{\Lambda}^{-1} \mathbf{Q}^\top \nabla f(\mathbf{x}_k)$$

In the Eigenspace:

$$\tilde{\mathbf{p}} = -\boldsymbol{\Lambda}^{-1} \tilde{\nabla} f$$



## Second-order Methods

---

- However, it is always hard to get Hessian Matrix
- For Weight matrix  $W \in \mathbb{R}^{m \times n}$ , the Hessian Matrix is:

$$H \in \mathbb{R}^{mn \times mn}$$

Computational complexity for  $H^{-1}$ :  $O(m^3n^3)$

- For LLM with 1B parameters, the memory cost for Hessian Matrix is:

$$10^{18} \text{ elements} \times 8 \text{ bytes/element} = 8 \times 10^{18} \text{ bytes} = 10^6 TB$$

The computational complexity for  $H^{-1}$  is:  $10^{27} FLOP$

- Can we decompose the Hessian Matrix?
- Kronecker Product: Given two matrices:  $A \in \mathbb{R}^{m \times n}, B \in \mathbb{R}^{p \times q}$ ,

The Kronecker product of  $A$  and  $B$ , denoted  $A \otimes B$ , is defined as the block matrix:

$$A \otimes B = \begin{bmatrix} a_{11}B & a_{12}B & \cdots & a_{1n}B \\ a_{21}B & a_{22}B & \cdots & a_{2n}B \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1}B & a_{m2}B & \cdots & a_{mn}B \end{bmatrix}$$

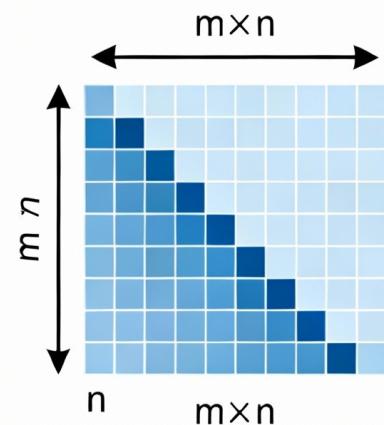
The resulting matrix  $A \otimes B$  has dimensions  $(mp) \times (nq)$ .

- Motivation: Exploits the structure of the parameter space



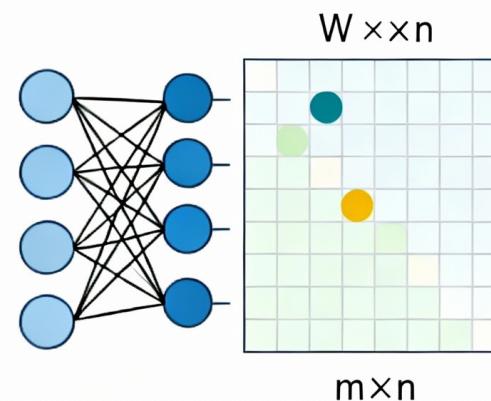
## Parameter Space Structures in Machine Learning

Multiclass Problem



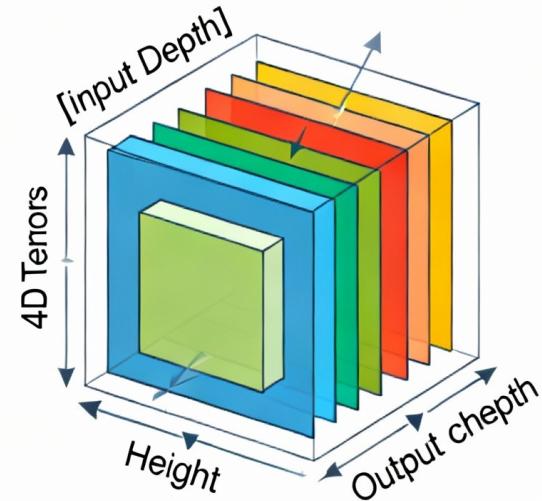
● features    ● classes

Fully-connected Neural Network Layer



● weight    ● parameter

Convolutional Neural Network (CNN)



# Shampoo

---



- Vector-form Second-order Method:

$$\mathbf{W}_t = \mathbf{W}_{t-1} - \eta_t \mathbf{H}_t^{-1} \mathbf{G}_t$$
 $O(m^2n^2)$   $O(m^3n^3)$

- Shampoo:

$$\mathbf{L}_t = \beta \mathbf{L}_{t-1} + \mathbf{G}_t \mathbf{G}_t^\top$$
 $O(m^2)$   $O(m^2n)$

$$\mathbf{R}_t = \beta \mathbf{R}_{t-1} + \mathbf{G}_t^\top \mathbf{G}_t$$
 $O(n^2)$   $O(mn^2)$

$$\mathbf{W}_t = \mathbf{W}_{t-1} - \eta_t \mathbf{L}_t^{-1/4} \mathbf{G}_t \mathbf{R}_t^{-1/4}$$
 $O(mn)$   $O(m^3 + n^3 + m^2n + mn^2)$

$$F \approx \mathbb{E}[gg^T]$$



$$g = \text{vec}(G) \in \mathbb{R}^{mn}$$



$$F = \mathbb{E}[gg^T] = \mathbb{E}[\text{vec}(G) \text{vec}(G)^T]$$

$$G = \delta x^T \in \mathbb{R}^{m \times n}$$



利用kronecker乘积的性质

$$F = \mathbb{E}[\text{vec}(G) \text{ vec}(G)^T] = \mathbb{E}[(xx^T) \otimes (\delta\delta^T)]$$



认为输入  $x$  和误差信号  $\delta$  的联合统计可以近似为“独立”

$$F \approx \mathbb{E}[xx^T] \otimes \mathbb{E}[\delta\delta^T]$$

# Shampoo

---



$$\mathbb{E}[GG^T] = \mathbb{E}[\|x\|^2 \delta \delta^T] \quad \mathbb{E}[G^T G] = \mathbb{E}[\|\delta\|^2 x x^T]$$



$\|x\|^2$ 、 $\|\delta\|^2$  与方向统计可以分离

$$L_t = \beta L_{t-1} + (1 - \beta) G_t G_t^T$$

$$R_t = \beta R_{t-1} + (1 - \beta) G_t^T G_t$$

# Shampoo

---

$$(L \otimes R)^{-1} = L^{-1} \otimes R^{-1}$$

$$\text{vec}(AXB^T) = (B \otimes A) \text{vec}(X)$$



$$H^{-1} \text{vec}(G) = (L \otimes R)^{-1} \text{vec}(G) = \text{vec}(L^{-1}GR^{-1})$$



$$\text{Update Matrix} = L^{-1} \cdot G \cdot R^{-1}$$

# Shampoo

- Algorithm in matrix case & tensor case

---

## Algorithm 1 Shampoo, matrix case.

Initialize  $W_1 = \mathbf{0}_{m \times n}$ ;  $L_0 = \epsilon I_m$ ;  $R_0 = \epsilon I_n$

**for**  $t = 1, \dots, T$  **do:**

    Receive loss function  $f_t : \mathbb{R}^{m \times n} \mapsto \mathbb{R}$

    Compute gradient  $G_t = \nabla f_t(W_t)$  //  $G_t \in \mathbb{R}^{m \times n}$

    Update preconditioners:

$$L_t = L_{t-1} + G_t G_t^\top$$

$$R_t = R_{t-1} + G_t^\top G_t$$

    Update parameters:

$$W_{t+1} = W_t - \eta L_t^{-1/4} G_t R_t^{-1/4}$$

---

## Algorithm 2 Shampoo, general tensor case.

Initialize:  $W_1 = \mathbf{0}_{n_1 \times \dots \times n_k}$ ;  $\forall i \in [k] : H_0^i = \epsilon I_{n_i}$

**for**  $t = 1, \dots, T$  **do:**

    Receive loss function  $f_t : \mathbb{R}^{n_1 \times \dots \times n_k} \mapsto \mathbb{R}$

    Compute  $G_t = \nabla f_t(W_t)$  //  $G_t \in \mathbb{R}^{n_1 \times \dots \times n_k}$

$\tilde{G}_t \leftarrow G_t$  //  $\tilde{G}_t$  is preconditioned gradient

**for**  $i = 1, \dots, k$  **do:**

$$H_t^i = H_{t-1}^i + G_t^{(i)}$$

$$\tilde{G}_t \leftarrow \tilde{G}_t \times_i (H_t^i)^{-1/2k}$$

    Update:  $W_{t+1} = W_t - \eta \tilde{G}_t$

# Shampoo

## ➤ Experiment Results

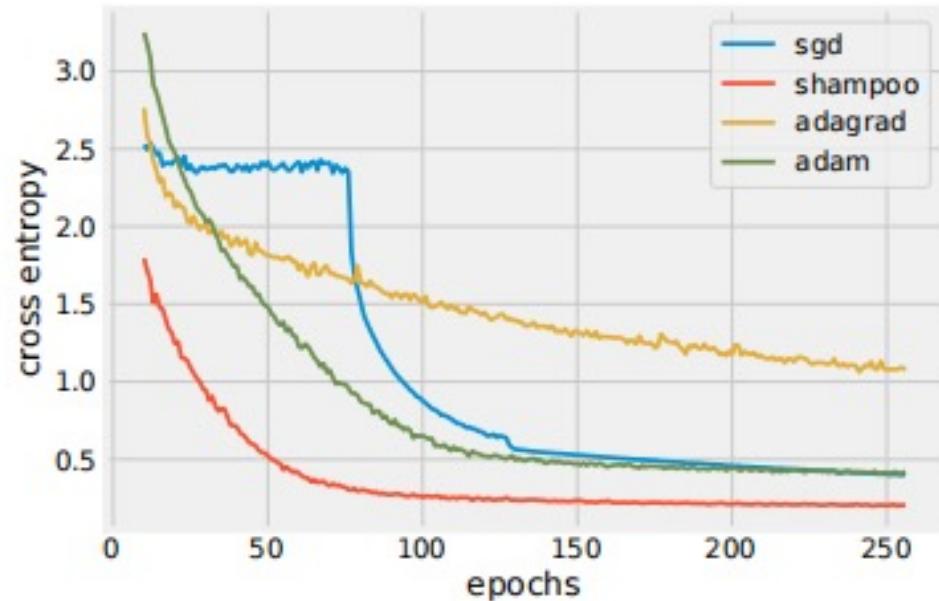


Figure 3. Convergence of training loss for a 55-layer ResNet on CIFAR-100.

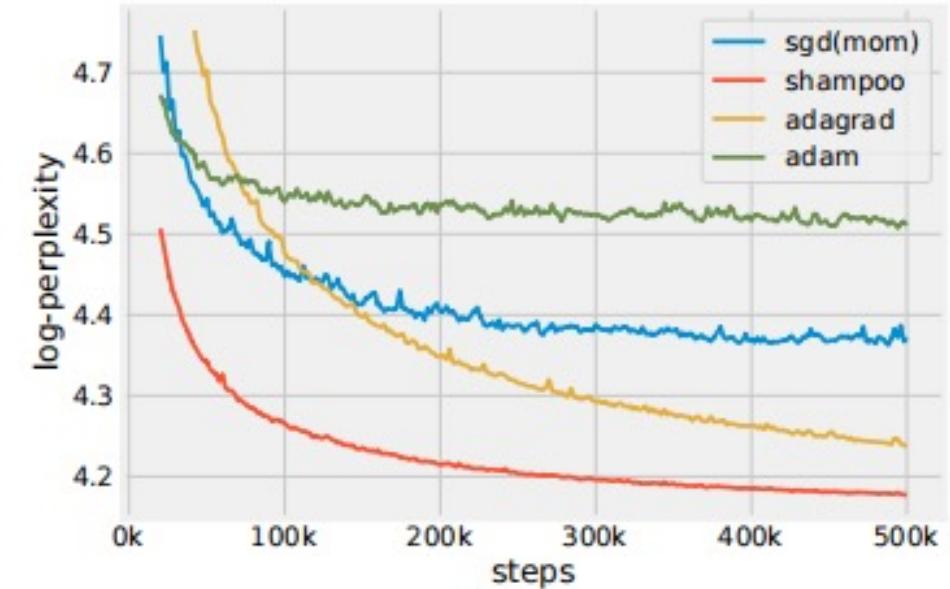
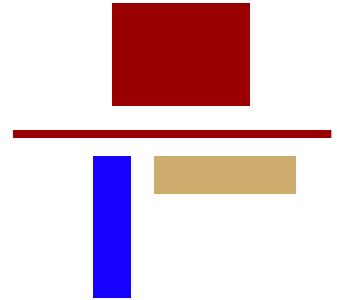


Figure 4. Convergence of test loss for the Transformer model for machine translation (Vaswani et al., 2017) on LM1B.

- Eigen Decomposition:

$$L = Q_L \Lambda_L Q_L^T$$



$$\text{Update} = Q_L (\Lambda_L^{-1/4} \underbrace{Q_L^T G Q_R}_{\text{旋转后的梯度 } G'} \Lambda_R^{-1/4}) Q_R^T$$



$$\Lambda_L^{-1/4} \cdot G' \cdot \Lambda_R^{-1/4} = \frac{G'}{\text{diag}\left(\Lambda_L^{\frac{1}{4}}\right) \text{diag}\left(\Lambda_R^{\frac{1}{4}}\right)^T}$$

- Observation: The variant of Shampoo is equivalent to running Adafactor in the eigenbasis provided by Shampoo's preconditioner

**Algorithm 1** Single step of idealized Shampoo with power 1/2.

- 1: Sample batch  $B_t$ .
- 2:  $G_t \in \mathbb{R}^{m \times n} \leftarrow -\nabla_W \phi_{B_t}(W_t)$
- 3:  $L \leftarrow \mathbb{E}_B[G_B G_B^T]$  {Where the expectation is over a random batch  $B$ .}
- 4:  $R \leftarrow \mathbb{E}_B[G_B^T G_B]$
- 5:  $\hat{H} \leftarrow L \otimes R / \text{Trace}(L)$
- 6:  $W_t \leftarrow W_{t-1} - \eta \hat{H}^{-1/2} G_t = W_{t-1} - \eta L^{-1/2} G_t R^{-1/2} / \text{Trace}(L)^{-1/2}$

**Algorithm 2** Single step of idealized Adafactor in Shampoo's eigenspace.

- 1: Sample batch  $B_t$ .
- 2:  $G_t \in \mathbb{R}^{m \times n} \leftarrow -\nabla_W \phi_{B_t}(W_t)$
- 3:  $L \leftarrow \mathbb{E}_B[G_B G_B^T]$
- 4:  $R \leftarrow \mathbb{E}_B[G_B^T G_B]$
- 5:  $Q_L \leftarrow \text{Eigenvectors}(L)$
- 6:  $Q_R \leftarrow \text{Eigenvectors}(R)$
- 7:  $G'_t \leftarrow Q_L^T G_t Q_R$
- 8: {Idealized version of code for Adafactor taking  $G'_t$  to be the gradient}
- 9:  $G'_B \leftarrow Q_L^T G_B Q_R$
- 10:  $A = \mathbb{E}_B[G'_B \odot G'_B] \mathbf{1}_m$  where  $G'_B = Q_L^T G_B Q_R$
- 11:  $C = \mathbf{1}_n^\top \mathbb{E}_B[G'_B \odot G'_B]$
- 12:  $\hat{V}_t = \frac{AC^T}{\mathbf{1}_n^\top A}$  {Elementwise division}
- 13:  $G''_t \leftarrow \frac{G'_t}{\sqrt{\hat{V}_t + \epsilon}}$  {Elementwise division and square root}
- 14:  $G'''_t \leftarrow Q_L G''_t Q_R^T$  {Projecting back to original space}
- 15:  $W_t \leftarrow W_{t-1} - \eta G'''_t$

- Inspiration: a broader design space for combining first and second order methods——running a first-order method in the eigenbasis provided by a second-order method

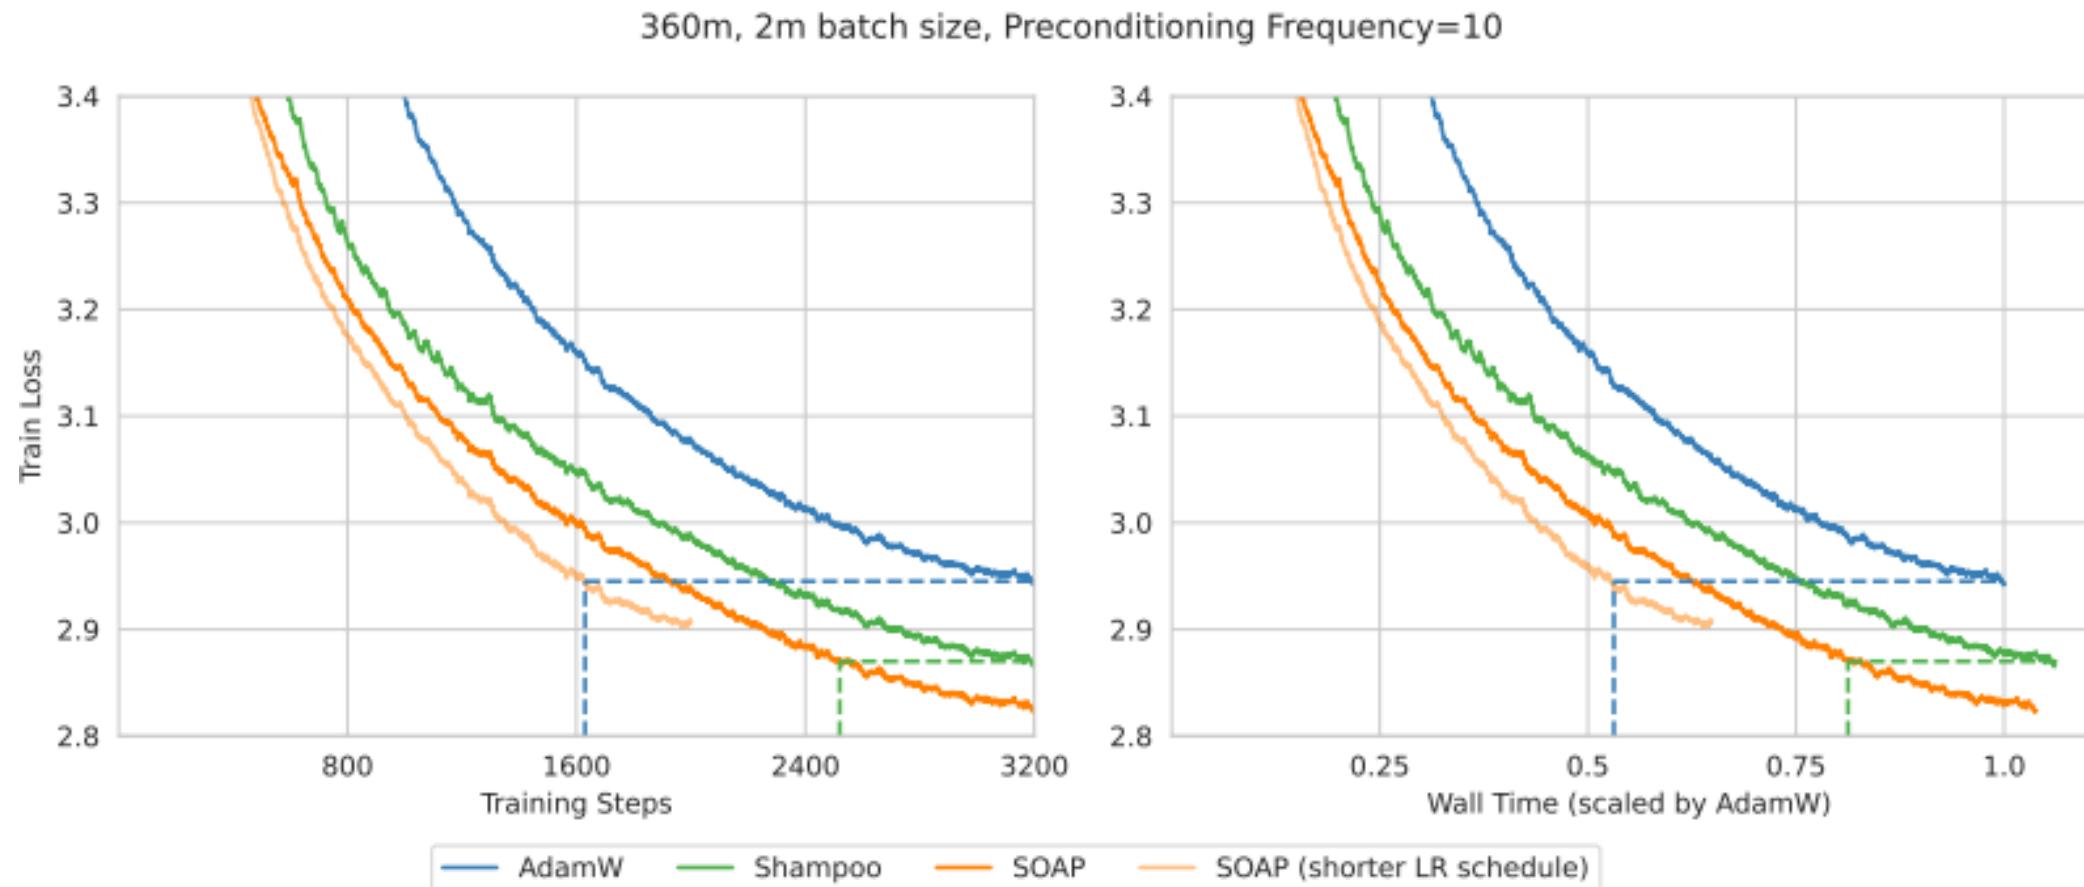
**Algorithm 3** Single step of SOAP for a  $m \times n$  layer. Per layer, we maintain four matrices:  $L \in \mathbb{R}^{m \times m}$ ,  $R \in \mathbb{R}^{n \times n}$  and  $V, M \in \mathbb{R}^{m \times n}$ . For simplicity we ignore the initialization and other boundary effects such as bias correction. Hyperparameters: Learning rate  $\eta$ , betas =  $(\beta_1, \beta_2)$ , epsilon  $\epsilon$ , and preconditioning frequency  $f$ .

An implementation of SOAP is available at <https://github.com/nikhilvyas/SOAP/tree/main>.

```

1: Sample batch  $B_t$ .
2:  $G \in \mathbb{R}^{m \times n} \leftarrow -\nabla_W \phi_{B_t}(W_t)$ 
3:  $G' \leftarrow Q_L^T G Q_R$ 
4:  $M \leftarrow \beta_1 M + (1 - \beta_1)G$ 
5:  $M' \leftarrow Q_L^T M Q_R$ 
6: {Now we “run” Adam on  $G'$ }
7:  $V \leftarrow \beta_2 V + (1 - \beta_2)(G' \odot G')$  {Elementwise multiplication}
8:  $N' \leftarrow \frac{M'}{\sqrt{V_t + \epsilon}}$  {Elementwise division and square root}
9: {Now that we have preconditioned by Adam in the rotated space, we go back to the original space.}
10:  $N \leftarrow Q_L N' Q_R^T$ 
11:  $W \leftarrow W - \eta N$ 
12: {End of gradient step, we now update  $L$  and  $R$  and possibly also  $Q_L$  and  $Q_R$ . }
13:  $L \leftarrow \beta_2 L + (1 - \beta_2)GG^T$ 
14:  $R \leftarrow \beta_2 R + (1 - \beta_2)G^T G$ 
15: if  $t \% f == 0$  then
16:    $Q_L \leftarrow \text{Eigenvectors}(L, Q_L)$ 
17:    $Q_R \leftarrow \text{Eigenvectors}(R, Q_R)$ 
18: end if
```

- Better robustness, Faster training





# Muon: An Orthogonalization-Based Optimizer for Deep Networks

Keller Jordan

# Motivation: Limitations of Standard Optimizers

---

- **Standard view:** parameters in deep learning are a **long vector**; we use SGD/Adam/AdamW on this vector. But hidden layers are actually **matrices**.
- **Empirical issue:**
  - Gradients/updates often have **highly skewed singular values** → poor conditioning
  - Many directions updated very weakly → slow learning of rare / subtle patterns
- **Question:** can we design an optimizer that respects **matrix structure** and **improves conditioning?**

# Muon Update Rules:

---

- Outline:

$$G_t \xrightarrow{\text{momentum}} M_t \xrightarrow{-\eta} U_t \xrightarrow{\text{NS-ortho}} Q_t \xrightarrow{+\alpha} W_{t+1}$$

- Details:

**Gradient:**  $G_t = \nabla_W L(W_t),$

**Momentum:**  $M_t = \beta M_{t-1} + (1 - \beta) G_t,$

**Raw update:**  $U_t = -\eta M_t,$

**Orthogonalization:**  $Q_t = \text{Ortho}_{\text{NS}}(U_t),$

**Parameter update:**  $W_{t+1} = W_t + \alpha Q_t.$

# Orthogonalization via K-step Newton–Schulz:

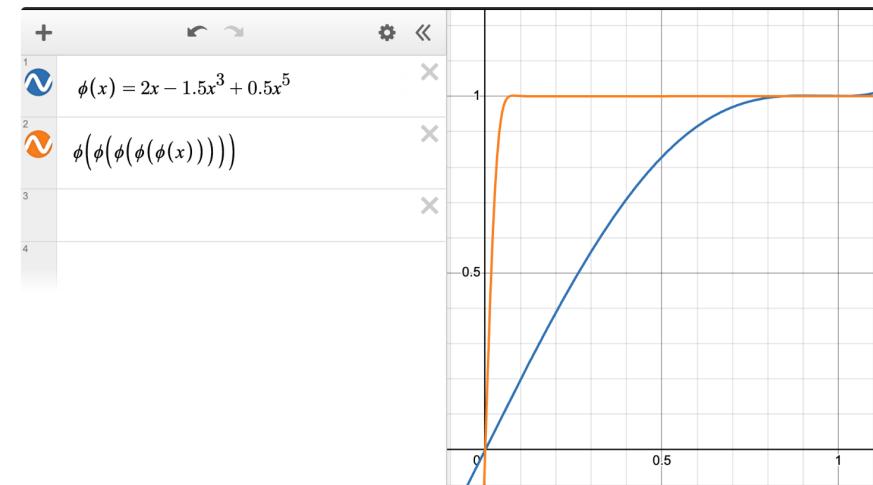
---

- Newton–Schulz iteration performs a **fast** and **low-cost approximate orthogonalization** of the target matrix. (SVD for parameter matrix is expensive.)
- How it works:

**Initialization:**  $U^{(0)} = \frac{U_t}{\|U_t\|_F}$ ,

**Iterations:**  $U^{(k+1)} = a U^{(k)} + b U^{(k)} U^{(k)\top} U^{(k)} + c (U^{(k)} U^{(k)\top})^2 U^{(k)},$   
 $k = 0, \dots, K - 1,$

**Output:**  $Q_t = U^{(K)} \approx \text{Ortho}_{\text{NS}}(U_t)$ .



- (a,b,c)=(3.1415,4.7750,2.0315) ,N=5 for final Muon design.

# Why is it good to orthogonalize the update?

---

## ➤ What does orthogonalization in Muon do:

Let  $W_t \in \mathbb{R}^{n \times m}$  be the weight matrix at time step  $t$ , and  $G_t$  be the gradient of the loss function with respect to  $W_t$ :

$$G_t = \nabla_W L(W_t)$$

The key idea is to apply an orthogonalization operator to  $U_t$ . The polar decomposition of a matrix  $G$  is given by:

$$G = QP$$

where  $Q$  is a semi-orthogonal matrix and  $P$  is a symmetric positive semidefinite matrix. The matrix  $Q$  is the nearest orthogonal matrix to  $G$ .

$$Q = \arg \min_{O: O^\top O = I} \|O - G\|_F$$

where  $O$  is any semi-orthogonal matrix. This is equivalent to:

$$Q = G(G^\top G)^{-\frac{1}{2}}$$

# Why is it good to orthogonalize the update?

## ➤ Properties of orthogonalization:

- Orthogonalization forces the **singular values** of  $\tilde{U}_t$  to be equal to 1, which improves the **conditioning** of the update.

$$\text{Singular values of } \tilde{U}_t : \quad \sigma_i(\tilde{U}_t) = 1 \quad \forall i$$

## ➤ Why is it good to orthogonalize the update?

- Updates produced by both SGD-momentum and Adam for the 2D parameters in transformer-based neural networks typically are almost **low-rank matrices**, with the updates for all neurons being dominated by just a few directions.
- Orthogonalization effectively increases the scale of other “**rare directions**” which have small magnitude in the update but are nevertheless **important for learning**.

## ➤ NS Iteration and Extra FLOPs:

- Before the NS iteration is applied, Muon is just a standard SGD-momentum optimizer, so it has the same memory requirement.
- For each  $n \times m$  matrix parameter in the network, each step of the NS iteration requires  $2(2nm^2 + m^3)$  matmul FLOPs.
- Therefore, the extra FLOPs required by Muon compared to SGD is at most  $2T(2nm^2 + m^3)$ , where  $T$  is the number of NS iterations (typically  $T = 5$ ).

## ➤ Extra computation rate:

- If the parameter parametrizes a linear layer, then the baseline amount of FLOPs used to perform a training step (i.e., a forward and backward pass) is  $6nmB$ , where  $B$  is the batch size in tokens.
- Therefore, the FLOP overhead of Muon is at most  $Tm/B$ , where  $m$  is the model dimension,  $B$  is the batch size in tokens, and  $T$  is the number of NS iteration steps

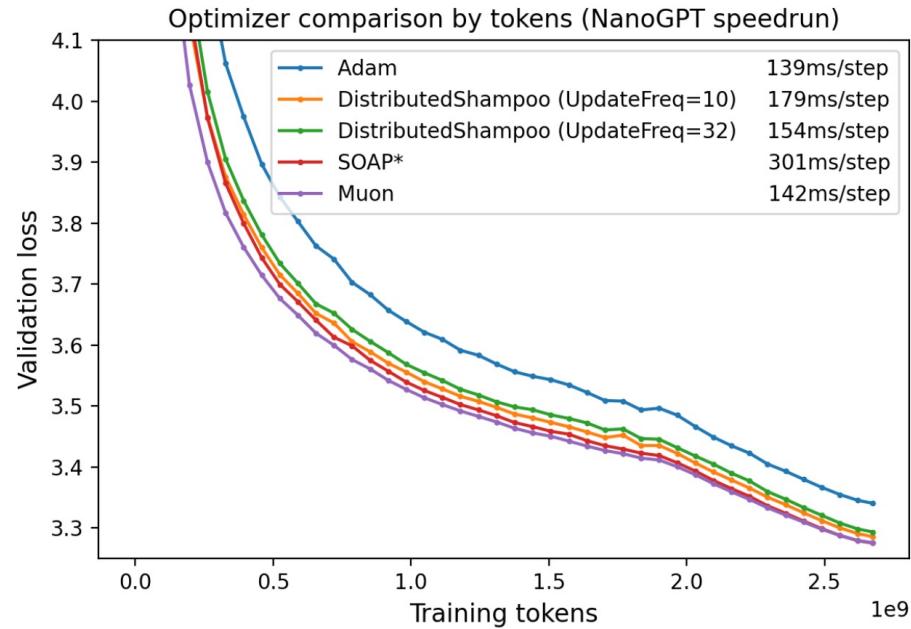
## ➤ NanoGPT overhead of using Muon:

- For the current NanoGPT speedrunning record, the model dimension is  $m = 768$  and the number of tokens per batch is  $B = 524288$ .
- Therefore, the overhead of using Muon is  $\frac{5 \times 768}{524288} = 0.7\%$ .

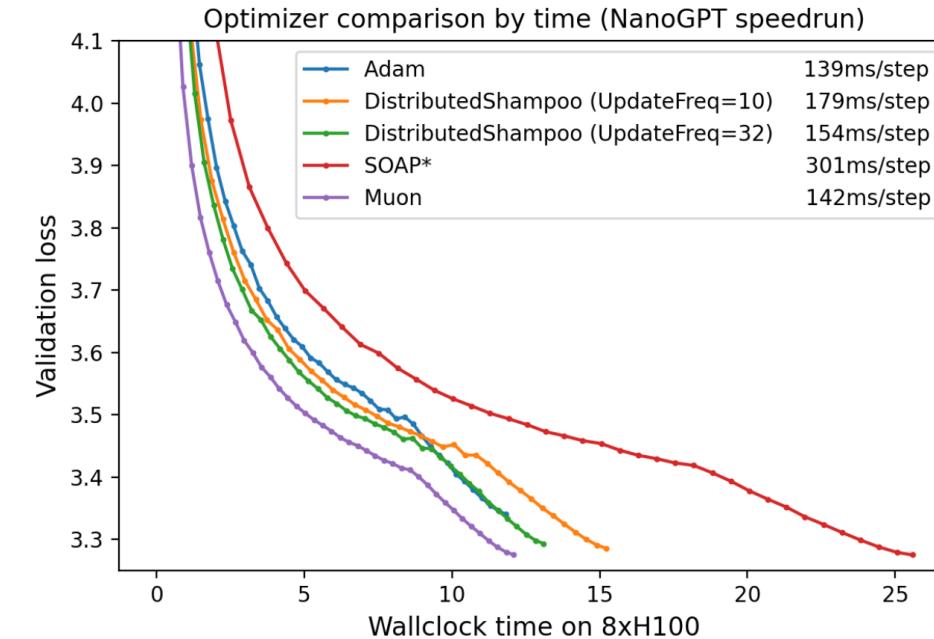
## ➤ Llama 405B overhead of using Muon:

- For Llama 405B training, the model dimension is  $m = 16384$  and the number of tokens per batch is reported to be  $B = 16000000$ .
- Therefore, the overhead of using Muon for this training would be  $\frac{5 \times 16384}{16000000} = 0.5\%$ .

# Muon Empirical Results (Jordan)



\*SOAP is under active development. Future versions will significantly improve the wallclock overhead.

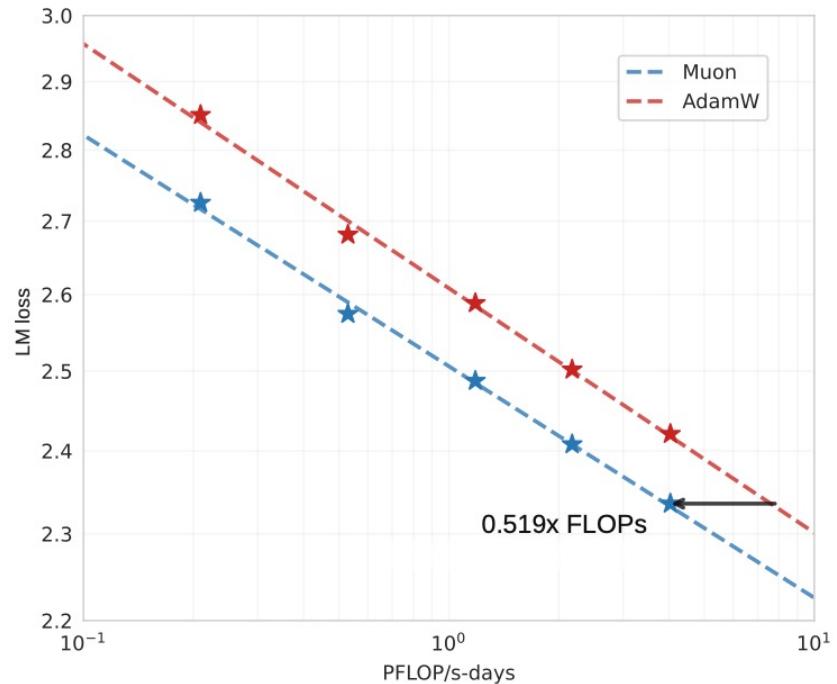


\*SOAP is under active development. Future versions will significantly improve the wallclock overhead.

**Fig1:** Improved the speed record for training to 3.28 val loss on FineWeb (a competitive task known as NanoGPT speedrunning) by a factor of **1.35x**.

**Fig2:** Optimizer comparison by wallclock time.

# Muon Empirical Results (Kimi)



**Fig1:** Muon uses ~ 52% less computational cost (FLOPs) during training compared to Adam optimizer in **Llama architechture**.

Table 4: Comparison of different models at around 1.2T tokens.

Benchmark (Metric)	DSV3-Small	Moonlight-A@1.2T	Moonlight@1.2T
Activated Params <sup>†</sup>	2.24B	2.24B	2.24B
Total Params <sup>†</sup>	15.29B	15.29B	15.29B
Training Tokens	1.33T	1.2T	1.2T
Optimizer	AdamW	AdamW	Muon
English	MMLU	53.3	60.2
	MMLU-pro	-	26.8
	BBH	41.4	45.3
	TriviaQA	-	57.4
Code	HumanEval	26.8	29.3
	MBPP	36.8	49.2
Math	GSM8K	31.4	43.8
	MATH	10.7	16.1
	CMath	-	57.8
Chinese	C-Eval	-	57.2
	CMMLU	-	58.2

<sup>†</sup> The reported parameter counts exclude the embedding parameters.

**Table4:** Moonlight(trained by Muon) performs significantly better than Moonlight-A(trained by AdamW), proving the scaling effectiveness of Muon. We observed that Muon especially excels on **Math and Code** related tasks.



# Sophia: A Scalable Stochastic Second-order Optimizer for Language Model Pre-training

Hong Liu   Zhiyuan Li   David Hall   Percy Liang   Tengyu Ma

Stanford University

{hliu99, zhiyuanli, dlwh, pliang, tengyuma}@cs.stanford.edu

# Motivations: Challenges in Large-Scale LLM Training

- AdamW uses only **first-order gradients**, while Loss landscape is **highly anisotropic**.
  - Some directions: high curvature
  - Others: flat
- Uniform step sizes cause:
  - Instability in steep directions
  - Slow progress in flat directions

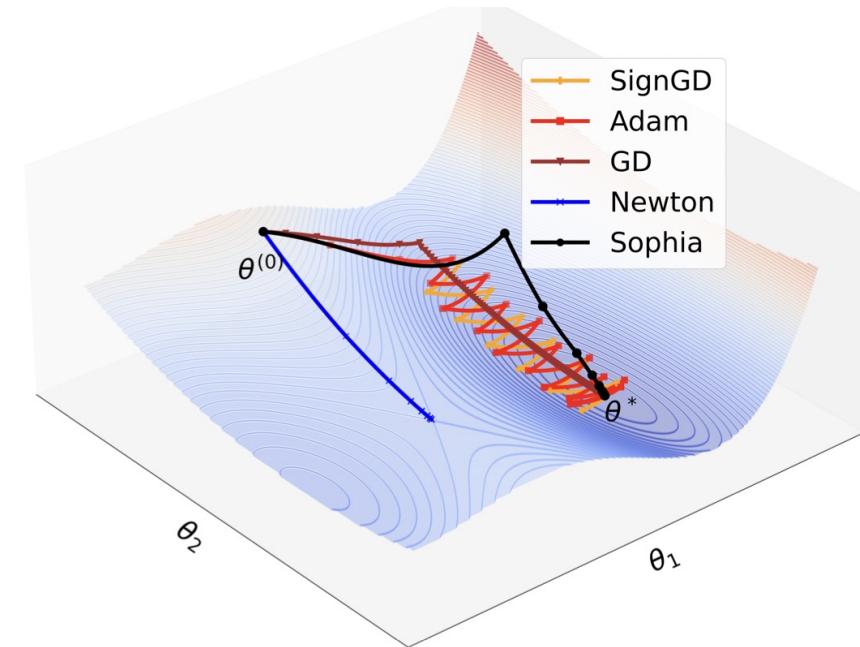


Fig: The motivating toy example.  $\theta[1]$  is the sharp dimension and  $\theta[2]$  is the flat dimension. GD's learning rate is limited by the sharpness in  $\theta[1]$ , and makes slow progress along  $\theta[2]$ . Adam and SignGD bounce along  $\theta[1]$  while making slow progress along  $\theta[2]$ . Vanilla Newton's method converges to a saddle point. Sophia makes fast progress in both dimensions and converges to the minimum with a few steps.

# Motivations: Need for Second-Order Information

---

- Second-order curvature helps scale updates per coordinate, but Full Hessian is impossible to compute or store.
  - Size =  $O(n^2)$  for  $n$  parameters.
  - Inverting or factorizing Hessian is infeasible.
- Goal: Use curvature approximation that is:
  - Informative
  - Cheap
  - Stable and simple like AdamW

# Preconditioner Estimation

---

- Idea: Use **Hessian Diagonal** as Preconditioner
- Approximates curvature per parameter dimension
- Scales updates as:

$$\Delta\theta_i \propto \frac{1}{H_{ii}}$$

- Large curvature → smaller update
- Small curvature → larger update
- **However, computing the Hessian directly is too expensive, so we use an approximation.**

# Preconditioner Estimation: Efficient Estimators for Hessian Diagonal

---

- **Option 1: Hutchinson Estimator**
  - Uses random vectors  $v$  to approximate  $\text{diag}(H)$
- **Option 2 : Gauss–Newton–Bartlett Estimator**
  - Uses gradients/Jacobians instead of Hessian
- Update **every k steps**, not every step.
- Maintain **EMA** smoothing for stability.

# How to estimate Hessian Diagonal?

➤ Option 1: Hutchinson Estimator

$$\hat{h} = u \odot (\nabla^2 \ell(\theta) u)$$

$$\mathbb{E}[\hat{h}] = \text{diag}(\nabla^2 \ell(\theta))$$

---

### Algorithm 1 Hutchinson( $\theta$ )

---

- 1: **Input:** parameter  $\theta$ .
  - 2: Compute mini-batch loss  $L(\theta)$ .
  - 3: Draw  $u$  from  $\mathcal{N}(0, I_d)$ .
  - 4: **return**  $u \odot \nabla(\langle \nabla L(\theta), u \rangle)$ .
- 

- Remark: The H-method only needs to compute the Hessian-vector product, which makes it efficient.

# How to estimate Hessian Diagonal?

---

## ➤ Option 2: Gauss–Newton–Bartlett Estimator

- The core idea of GNB is to estimate the Hessian diagonal (or, more precisely, the diagonal of its Fisher information) using the **element-wise squared gradient**.

---

### Algorithm 2 Gauss-Newton-Bartlett( $\theta$ )

---

- 1: **Input:** parameter  $\theta$ .
  - 2: Draw a mini-batch of input  $\{x_b\}_{b=1}^B$ .
  - 3: Compute logits on the mini-batch:  
 $\{f(\theta, x_b)\}_{b=1}^B$ .
  - 4: Sample  $\hat{y}_b \sim \text{softmax}(f(\theta, x_b)), \forall b \in [B]$ .
  - 5: Calculate  $\hat{g} = \nabla(1/B \sum \ell(f(\theta, x_b), \hat{y}_b))$ .
  - 6: **return**  $B \cdot \hat{g} \odot \hat{g}$ .
-

# GNB: Why Can the Squared Gradient Represent 2nd-Order Curvature?



- We consider the log-likelihood:

$$\log p_\theta(y \mid x)$$

- Negative log-likelihood (NLL):

$$\ell(\theta; x, y) = -\log p_\theta(y \mid x)$$

- Define the score (the gradient of the log-likelihood) as:

$$s(\theta) = \nabla_\theta \log p_\theta(y \mid x).$$

- Then the gradient of the NLL is:

$$\nabla_\theta \ell(\theta; x, y) = -s(\theta).$$

# GNB: Why Can the Squared Gradient Represent 2nd-Order Curvature?



## ➤ Bartlett Identity:

- Under mild regularity conditions, the Bartlett identity states that, for data drawn from the model's own distribution (i.e.,  $y \sim p_\theta(\cdot | x)$ ),

$$\mathbb{E}_{y \sim p_\theta}[s(\theta)] = 0, \quad \mathbb{E}_{y \sim p_\theta}[s(\theta)s(\theta)^\top] = -\mathbb{E}_{y \sim p_\theta}[\nabla_\theta^2 \log p_\theta(y | x)].$$

## ➤ Rewriting in Terms of the NLL:

$$\mathbb{E}_{y \sim p_\theta}[\nabla_\theta \ell \nabla_\theta \ell^\top] = \mathbb{E}_{y \sim p_\theta}[\nabla_\theta^2 \ell].$$

- **Interpretation:** This means when labels are sampled from the model distribution, the expected Hessian equals the expected gradient outer product.

# GNB: Why Can the Squared Gradient Represent 2nd-Order Curvature?



- Therefore, a very natural estimator for the Hessian diagonal is:

$$\text{diag}(\nabla_{\theta}^2 \ell) \approx \mathbb{E}[(\nabla_{\theta} \ell) \odot (\nabla_{\theta} \ell)],$$

where  $\odot$  denotes element-wise multiplication.

# Bartlett Identity: Mathematical Derivation

---



We consider a conditional probabilistic model  $p_\theta(y | x)$ , where  $x$  is treated as fixed and  $\theta \in R^d$  denotes the model parameters. Define the *score function* as

$$s(\theta; y, x) := \nabla_\theta \log p_\theta(y | x).$$

All expectations below are taken with respect to  $y \sim p_\theta(\cdot | x)$ .

**Regularity assumptions.** We assume that:

- $p_\theta(y | x)$  is sufficiently smooth in  $\theta$ ;
- differentiation and integration (or summation) can be interchanged;
- the support of  $p_\theta(y | x)$  does not depend on  $\theta$ .

# Bartlett Identity: Mathematical Derivation

---

**First identity:**  $E[s(\theta)] = 0$ . Since  $p_\theta(y | x)$  is a conditional probability distribution, it satisfies

$$\int p_\theta(y | x) dy = 1.$$

Taking the gradient with respect to  $\theta$  yields

$$\nabla_\theta \int p_\theta(y | x) dy = \nabla_\theta 1 = 0.$$

Under the regularity assumptions, we may interchange differentiation and integration:

$$\int \nabla_\theta p_\theta(y | x) dy = 0.$$

# Bartlett Identity: Mathematical Derivation

---



Using the identity

$$\nabla_{\theta} p_{\theta}(y \mid x) = p_{\theta}(y \mid x) \nabla_{\theta} \log p_{\theta}(y \mid x),$$

we obtain

$$\int p_{\theta}(y \mid x) s(\theta; y, x) dy = 0,$$

which implies

$$E_{y \sim p_{\theta}} [s(\theta; y, x)] = 0.$$

# Bartlett Identity: Mathematical Derivation

---

**Second identity:**  $E[s(\theta)s(\theta)^\top] = -E[\nabla_\theta^2 \log p_\theta]$ . We start from the second derivative of the log-likelihood:

$$\nabla_\theta \log p_\theta(y | x) = \frac{\nabla_\theta p_\theta(y | x)}{p_\theta(y | x)}.$$

Taking another derivative with respect to  $\theta$  gives

$$\nabla_\theta^2 \log p_\theta(y | x) = \frac{\nabla_\theta^2 p_\theta(y | x)}{p_\theta(y | x)} - \frac{\nabla_\theta p_\theta(y | x) \nabla_\theta p_\theta(y | x)^\top}{p_\theta(y | x)^2}.$$

Noting that

$$\frac{\nabla_\theta p_\theta(y | x)}{p_\theta(y | x)} = \nabla_\theta \log p_\theta(y | x) = s(\theta; y, x),$$

we can rewrite the above as

$$\nabla_\theta^2 \log p_\theta(y | x) = \frac{\nabla_\theta^2 p_\theta(y | x)}{p_\theta(y | x)} - s(\theta; y, x)s(\theta; y, x)^\top.$$

# Bartlett Identity: Mathematical Derivation

---



Multiplying both sides by  $p_\theta(y | x)$  and integrating over  $y$ , we obtain

$$\int p_\theta(y | x) \nabla_\theta^2 \log p_\theta(y | x) dy = \int \nabla_\theta^2 p_\theta(y | x) dy - \int p_\theta(y | x) s(\theta; y, x) s(\theta; y, x)^\top dy.$$

The first term on the right-hand side satisfies

$$\int \nabla_\theta^2 p_\theta(y | x) dy = \nabla_\theta^2 \int p_\theta(y | x) dy = \nabla_\theta^2 1 = 0.$$

Therefore,

$$E_{y \sim p_\theta} [\nabla_\theta^2 \log p_\theta(y | x)] = -E_{y \sim p_\theta} [s(\theta; y, x) s(\theta; y, x)^\top],$$

or equivalently,

$$E_{y \sim p_\theta} [s(\theta; y, x) s(\theta; y, x)^\top] = -E_{y \sim p_\theta} [\nabla_\theta^2 \log p_\theta(y | x)].$$

# Bartlett Identity: Mathematical Derivation

---



**Negative log-likelihood form.** Let  $\ell(\theta; y, x) = -\log p_\theta(y | x)$ . Then

$$\nabla_\theta \ell = -s(\theta), \quad \nabla_\theta^2 \ell = -\nabla_\theta^2 \log p_\theta.$$

Hence, the Bartlett identity can be written as

$$E[\nabla_\theta \ell \nabla_\theta \ell^\top] = E[\nabla_\theta^2 \ell],$$

which underlies the Gauss–Newton–Bartlett estimator used in second-order optimization methods.

# Sophia Update Rule

---

## Algorithm 3 Sophia

```

1: Input:  $\theta_1$ , learning rate  $\{\eta_t\}_{t=1}^T$ , hyperparameters  $\lambda, \gamma, \beta_1, \beta_2, \epsilon$ , and estimator choice
   Estimator  $\in \{\text{Hutchinson, Gauss-Newton-Bartlett}\}$ 
2: Set  $m_0 = 0, v_0 = 0, h_{1-k} = 0$ 
3: for  $t = 1$  to  $T$  do
4:   Compute minibatch loss  $L_t(\theta_t)$ .
5:   Compute  $g_t = \nabla L_t(\theta_t)$ .
6:    $m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$ 
7:   if  $t \bmod k = 1$  then
8:     Compute  $\hat{h}_t = \text{Estimator}(\theta_t)$ .
9:      $h_t = \beta_2 h_{t-k} + (1 - \beta_2) \hat{h}_t$ 
10:  else
11:     $h_t = h_{t-1}$ 
12:     $\theta_t = \theta_t - \eta_t \lambda \theta_t$  (weight decay)
13:     $\theta_{t+1} = \theta_t - \eta_t \cdot \text{clip}(m_t / \max\{\gamma \cdot h_t, \epsilon\}, 1)$ 

```

Adam-like Momentum Term

# Sophia Update Rule

## ➤ Core Update Formula:

$$\theta_{t+1} = \theta_t - \eta_t \cdot \text{clip}\left(\frac{m_t}{\max(\gamma h_t, \epsilon)}, 1\right)$$

## ➤ Notations:

$m_t$ : EMA of gradients

$h_t$ : EMA of Hessian diagonal estimate

$\gamma$ : curvature scaling factor

$\epsilon$ : numerical stability

## ➤ Interpretation:

- Hessian diagonal acts as adaptive step size
- Each coordinate has its own curvature-aware scaling
- Update stabilizer:  $\max(\gamma h_t, \epsilon)$

# Per-coordinate Clipping

---

## ➤ Why Needed?

- Hessian estimates are noisy
- High variance may produce excessively large updates

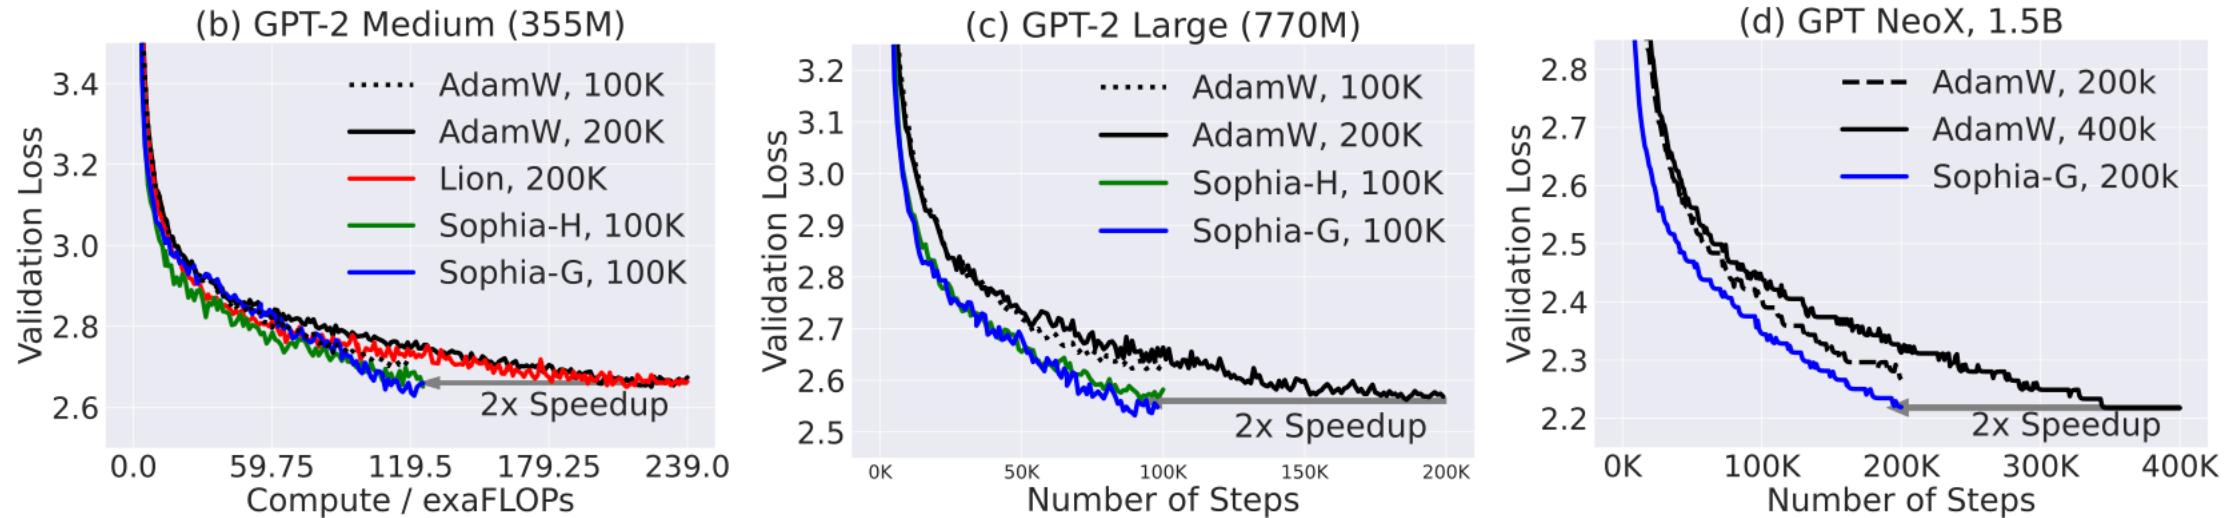
## ➤ Solution: Per-coordinate Clip

- Bound each coordinate's update:

$$|u_i| \leq 1$$

- In practice: ensures stable training even with imperfect curvature

# Experiments



(b) GPT-2 Medium (355M). (c) GPT-2 Large (770M). (d) GPT NeoX 1.5B. Across all model sizes, Sophia achieves a 2x speedup.

# Summary

---

- Sophia is designed to:
  - Use informative **second-order curvature** (diagonal Hessian)
  - Maintain very **low computation** cost
  - Remain **stable** via per-coordinate clipping
  - Achieve faster convergence than AdamW in large-scale models
- Key innovation:
  - A practical second-order optimizer with LLM-scale compatibility.