

Introduction to Large Language Models

Adaptive SGD

Kun Yuan

Peking University

Main contents in this lecture

- Preconditioned SGD
- AdaGrad
- RMSProp
- Adam

Preconditioned GD

- Consider an ill-conditioned quadratic problem

$$\min_x x^T Q x + c^T x$$

where Q is an ill-conditioned matrix. GD is slow when solving the problem




Figure: An ill-conditioned QP problem. (From Prof. Chris De Sa's lecture notes)

Preconditioned GD

- We now let $x = P^{\frac{1}{2}}w$ for some positive definite matrix P . Since P is positive definite, x and w is an 1 – 1 mapping
- If we choose $P = Q^{-1}$, we have $x^T Q x = w^T Q^{-\frac{1}{2}} Q Q^{\frac{1}{2}} w = \|w\|^2$
- With $x = P^{\frac{1}{2}}w$ and $P = Q^{-1}$, the ill-conditioned problem becomes

$$\min_w \quad \frac{1}{2} \|w\|^2 + c^T Q^{-\frac{1}{2}} w$$

which is a benign problem. GD is fast to achieve w^* .

- Once w^* is determined, we have $x^* = P^{\frac{1}{2}} w^*$.

Preconditioned GD




Figure: Left: an ill-conditioned QP problem. Right: a benign QP problem after transformation.(From Prof. Chris De Sa's lecture notes)

Preconditioned GD: derivation

- Consider a general ill-conditioned optimization problem

$$\min_{x \in \mathbb{R}^d} f(x)$$

- We let $x = P^{\frac{1}{2}}w$ so that $g(w) = f(P^{\frac{1}{2}}w)$ is a nice function.
- Use gradient descent to minimize $g(w)$, i.e.,

$$w_{k+1} = w_k - \gamma \nabla g(w_k) = w_k - \gamma P^{\frac{1}{2}} \nabla f(P^{\frac{1}{2}}w_k)$$

- Left-multiplying $P^{\frac{1}{2}}$ to both sides, we achieve

$$\begin{aligned} P^{\frac{1}{2}}w_{k+1} &= P^{\frac{1}{2}}w_k - \gamma P \nabla f(P^{\frac{1}{2}}w_k) \\ \iff x_{k+1} &= x_k - \gamma P \nabla f(x_k) \end{aligned}$$

where P is called the preconditioning matrix.

Preconditioned GD

- The preconditioned GD algorithm

$$x_{k+1} = x_k - \gamma P_k \nabla f(x_k)$$

where P_k varies with iteration k .

- It is critical to choose the preconditioning matrix P_k
- If $P_k = [\nabla^2 f(x_k)]^{-1}$, then preconditioned GD reduces to Newton's method
- It is critical to construct an efficient and effective P matrix

Stochastic optimization

- Consider the stochastic optimization problem:

$$\min_{x \in \mathbb{R}^d} f(x) = \mathbb{E}_{\xi \sim \mathcal{D}}[F(x; \xi)]$$

- ξ is a random variable indicating data samples
 - \mathcal{D} is the data distribution; unknown in advance
 - $F(x; \xi)$ is differentiable in terms of x
- Similar to preconditioned GD, **preconditioned SGD** iterates as follows

$$x_{k+1} = x_k - \gamma P_k \nabla F(x_k; \xi_k)$$

Adaptive gradient method (AdaGrad)

- Adaptive gradient method

$$g_k = \nabla F(x_k; \xi_k)$$

$$s_k = s_{k-1} + g_k \odot g_k$$

$$x_{k+1} = x_k - \frac{\gamma}{\sqrt{s_k} + \epsilon} \odot g_k$$

where $1/\sqrt{s_k} = \text{col}\{1/\sqrt{s_{k,1}}, \dots, 1/\sqrt{s_{k,d}}\} \in \mathbb{R}^d$ is an element-wise operation, s_0 is initialized as 0, and a small ϵ is added for safe-guard.

Adaptive gradient method (AdaGrad)

- AdaGrad falls into preconditioned SGD
- If we let $P_k = \text{diag}\{\frac{1}{\sqrt{s_{k,1}}+\epsilon}, \dots, \frac{1}{\sqrt{s_{k,d}}+\epsilon}\} \in \mathbb{R}^{d \times d}$, AdaGrad becomes

$$x_{k+1} = x_k - \gamma P_k g_k$$

where P_k is a time-varying preconditioning matrix.

- AdaGrad imposes smaller learning rates for notable gradient directions
- AdaGrad imposes larger learning rates for insignificant gradient directions

Adaptive gradient method (AdaGrad)

AdaGrad alleviates the “Zig-Zag” phenomenon




Figure: GD converges slow for ill-conditioned problem

Adaptive gradient method (AdaGrad)

AdaGrad alleviates the “Zig-Zag” phenomenon





Figure: AdaGrad has alleviated “Zig-Zag” phenomenon


Adaptive gradient method (AdaGrad)

The learning rate in AdaGrad is adaptive; no need to tune.



Subgradient g_k stays constant

$$\gamma_k = \frac{\gamma}{\sqrt{\sum_{t=1}^k g_t^2}} = O\left(\frac{1}{\sqrt{T}}\right)$$



Gradient decays at $g_k = O(\rho^k)$

$$\gamma_k = \frac{\gamma}{\sqrt{\sum_{t=1}^k g_t^2}} = O(1)$$

Figure: AdaGrad automatically adapts to problem structure¹.

¹These examples are from <https://conferences.mpi-inf.mpg.de/adfocs/material/alina/adaptive-L1.pdf>

RMSProp

- Since s_k keeps increasing, the rate γ_k in AdaGrad keeps decreasing
- AdaGrad may suffer from slow convergence
- RMSProp proposes a different way to construct s_k

$$s_k = \beta s_{k-1} + (1 - \beta) g_k \odot g_k$$

where $\beta \in (0, 1)$. A typical value for β is 0.9.

- In RMSProp, only the most recent g_k influences the convergence rate

RMSProp

- Suppose $g_k = 1/k$, we can visualize s_k from AdaGrad and RMSProp

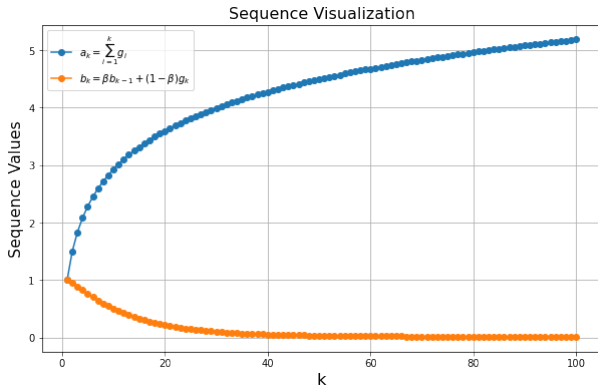


Figure: AdaGrad increases very fast while RMSProp decays slowly with $\beta = 0.9$

RMSPProp

- We also visualize s_k from RMSPProp with different β .

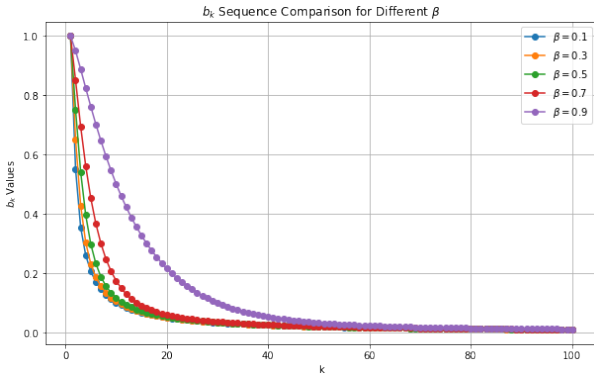


Figure: Gradient accumulation in RMSPProp with different β .

- RMSProp has the following update

$$g_k = \nabla F(x_k; \xi_k)$$

$$s_k = \beta s_{k-1} + (1 - \beta) g_k \odot g_k$$

$$x_{k+1} = x_k - \frac{\gamma}{\sqrt{s_k} + \epsilon} \odot g_k$$

where s_0 is initialized as 0, and a small ϵ is added for safe-guard.

- Adam applies both momentum and adaptive rate to alleviate “Zig-Zag”.

$$g_k = \nabla F(x_k; \xi_k)$$

$$m_k = \beta_1 m_{k-1} + (1 - \beta_1) g_k$$

$$s_k = \beta_2 s_{k-1} + (1 - \beta_2) g_k \odot g_k$$

$$x_{k+1} = x_k - \frac{\gamma}{\sqrt{s_k} + \epsilon} \odot m_k$$

where m_0 and s_0 are initialized as 0, and a small ϵ is added for safe-guard.

- It is good to set $\beta_1 = 0.9$ and $\beta_2 = 0.999$.

Animation of different adaptive SGD

<https://imgur.com/a/Hqolp>

Numerical performance




Figure: This figure is from the Adam paper (Kingma and Ba, 2014)

References I

- D. P. Kingma and J. Ba, "Adam: A method for stochastic optimization," *arXiv preprint arXiv:1412.6980*, 2014.
- Z. Guo, Y. Xu, W. Yin, R. Jin, and T. Yang, "A novel convergence analysis for algorithms of the adam family," 2021.
- I. Loshchilov and F. Hutter, "Decoupled weight decay regularization," in *International Conference on Learning Representations*, 2018.